

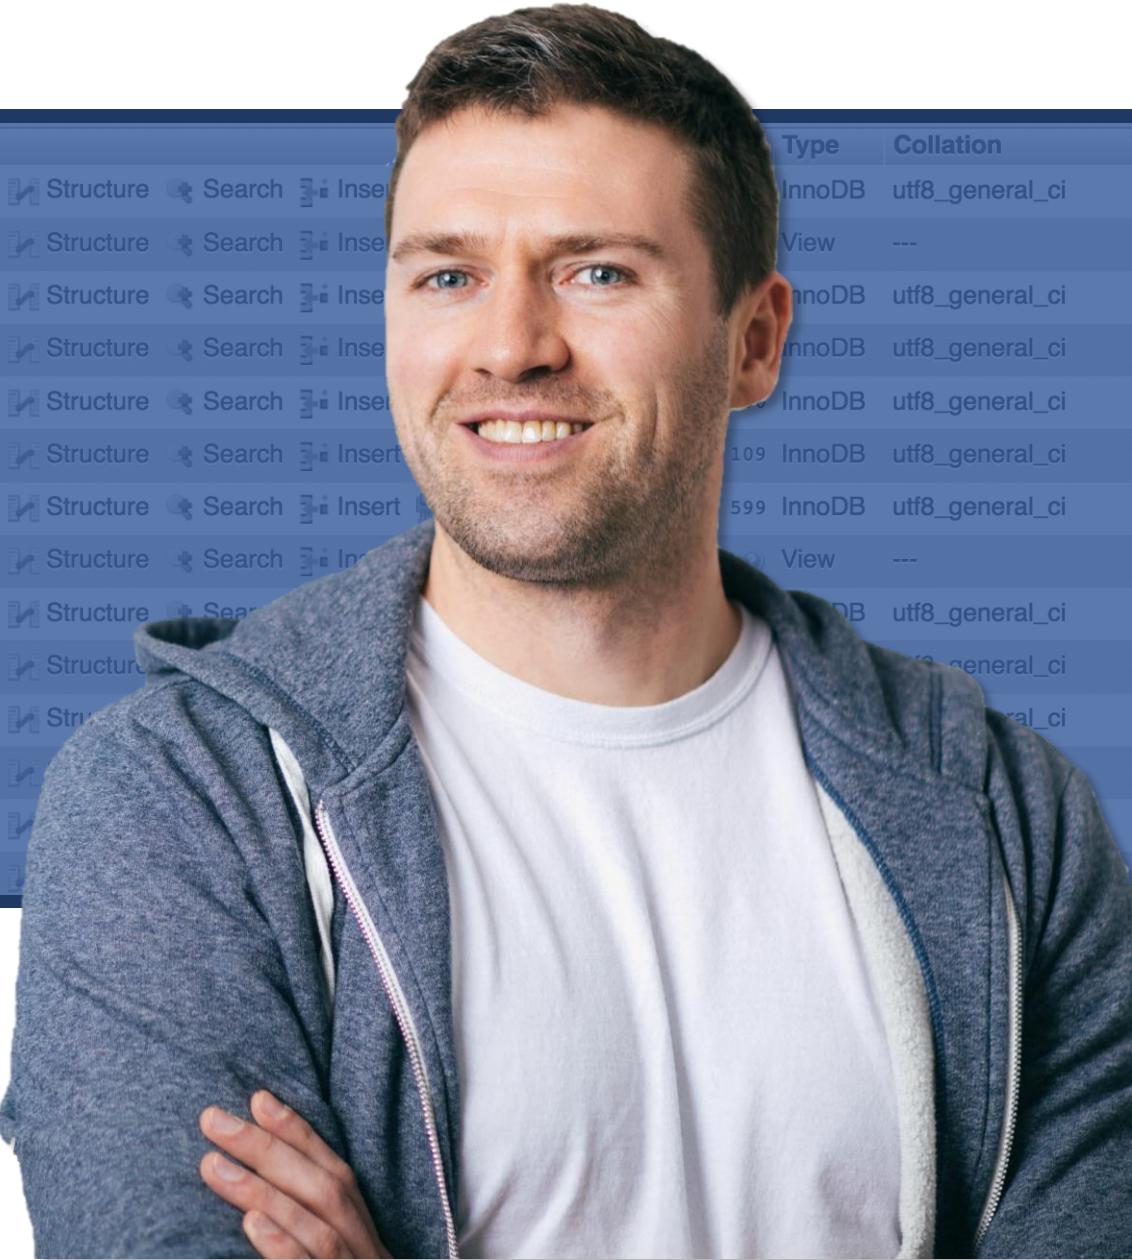
# ADVANCED SQL DATABASE ADMINISTRATION

★★★★★ With Expert SQL Instructor John Pauler

The screenshot shows a MySQL Workbench interface. On the left, there's a 'Customer Data' schema diagram with tables like city, country, customer, and film. In the center, a query editor displays the following SQL code:

```
1 SELECT
2     customer.first_name,
3     customer.last_name,
4     rental.rental_date,
5     rental.return_date,
6     rental.rental_id,
7     rental.inventory_id,
8     rental.store_id,
9     rental.rental_date,
10    rental.rental_id,
11    rental.inventory_id,
12    rental.store_id
13 FROM
14     customer
15     JOIN rental ON customer.customer_id = rental.customer_id
16 GROUP BY
17     customer.first_name,
18     customer.last_name
19 ORDER BY
20     rental.rental_date
```

Below the code, a message says: "Showing rows 0 to 24 (599 total, Query took 0.0212 seconds.)". On the right, a list of actions for various tables is shown, such as Browse, Structure, Search, Insert, and View.



# COURSE STRUCTURE

---



This is a **project-based course**, for students looking for a *practical, hands-on, and highly engaging* approach to building & maintaining databases with MySQL

*Additional resources include:*

- ★ **Downloadable Ebook** to serve as a helpful reference when you're offline or on the go
- ★ **Quizzes & Homework Exercises** to test and reinforce key concepts, with step-by-step solutions
- ★ **Bonus Projects** to test your abilities and apply the skills developed throughout the course

# COURSE OUTLINE

1	<b>Introduction &amp; Setup</b>	<i>Discuss the basic prerequisites, download Community Server and Workbench, and review the Workbench Interface</i>
2	<b>File Import, Alter, Update &amp; Keys</b>	<i>Insert large datasets to the database using the wizard. Review alter and update, and map primary and foreign keys in our new schema</i>
3	<b>Replication, Backup &amp; Recovery</b>	<i>Learn about replication, and how and when we use it. Discuss methods for backup and recovery of data to ensure nothing is lost</i>

## [MID-COURSE PROJECT]

4	<b>Automating Database Activity</b>	<i>Use automation like triggers, stored procedures and scheduled events, and learn how website data is written to a DB in real-time</i>
5	<b>Creating Views &amp; EER Diagrams</b>	<i>Learn how Views &amp; EER Diagrams are used as tools for reporting and understanding data in your database</i>
6	<b>Database Security</b>	<i>Discuss best practices in database security, and learn what can happen when security isn't taken seriously</i>

## [FINAL PROJECT]

# INTRODUCING THE COURSE PROJECT

## THE SITUATION

You've just been hired as a **Database Administrator** for **Maven Bear Builders**, an online retailer which has just launched their first product.

## THE BRIEF

As a member of the startup team, you will work with the *CEO*, the *Head of Marketing*, and the *Website Manager* to help build their data infrastructure so they can grow their business.

You will create schemas and tables, load large datasets to the database, use automation to keep things running smoothly, and tackle serious issues like security, backup, and data recovery.

## THE OBJECTIVE

**Use SQL to:**

- *Build and maintain the Maven Bear Builders database*
- *Become the data expert for the company, and the go-to person for mission critical data enhancements*
- *Leverage automation, and make decisions about security and backup to plan for long-term success*

# SETTING EXPECTATIONS

1

## You'll learn MySQL Database Administration using MySQL Workbench

- *In your career, you may use other “flavors” of SQL (T-SQL, PL/SQL, PostgreSQL, etc.)*
- *Each flavor is very similar. The principles you learn here will apply universally, but syntax will vary slightly*

2

## We will focus on the creation and maintenance of databases

- *We will design and create schemas and tables from scratch, and alter existing tables to improve them*

3

## This course is meant for people who already have basic SQL + DBA skills

- *This course skips over the SQL basics, and jumps straight into a simulated real-world experience where students can practice building and maintaining a real-world database*
- *Prior completion of MySQL Database Administration for Beginners and MySQL for Data Analysis is recommended*

4

## We will NOT go deep on data analysis or into engineering

- *This course will focus on Database Administration concepts, like creating and maintaining databases*
- *We offer separate courses that focus on using SQL for data analysis and business intelligence (beginner + advanced-level)*
- *We will not be covering languages like PHP, Python, or Ruby*

# PREREQUISITE SKILLS REVIEW



# PREREQUISITE DBA SKILLS

---

- You should be familiar with **the basic statements** used in Data Definition and Data Manipulation: **CREATE, ALTER, INSERT, UPDATE, & DELETE**
  
- You should have some understanding of automation using database tools like **TRIGGERS**, and **STORED PROCEDURES**
  
- You should be able to create **EER DIAGRAMS**, and should be familiar with Database Design concepts like **NORMALIZATION, & CARDINALITY**

# COMMON MySQL DATA TYPES

Data Type	Specifications
<b>TINYINT</b>	<i>Integer (-128 to 127)</i>
<b>SMALLINT</b>	<i>Integer (-32768 to 32767)</i>
<b>MEDIUMINT</b>	<i>Integer (-8388608 to 8388607)</i>
<b>INT</b>	<i>Integer (-2147483648 to 2147483647)</i>
<b>BIGINT</b>	<i>Integer (-9223372036854775808 to 9223372036854775807)</i>
<b>FLOAT</b>	<i>Decimal (precise to 23 digits)</i>
<b>DOUBLE</b>	<i>Decimal (23 to 53 digits)</i>
<b>DECIMAL</b>	<i>Decimal (to 65 digits – most precise)</i>

Data Type	Specifications
<b>CHAR</b>	<i>String (0 – 255)</i>
<b>VARCHAR</b>	<i>String (0 – 255)</i>
<b>TINYTEXT</b>	<i>String (0 – 255)</i>
<b>TEXT</b>	<i>String (0 – 65535)</i>
<b>DATE</b>	<i>YYYY-MM-DD</i>
<b>DATETIME</b>	<i>YYYY-MM-DD HH:MM:SS</i>
<b>TIMESTAMP</b>	<i>YYYYMMDDHHMMSS</i>
<b>ENUM</b>	<i>One of a number of preset options</i>

# THE “BIG 6” ELEMENTS OF A SQL SELECT STATEMENT

START OF  
STATEMENT

**SELECT**

*Identifies the column(s) you want your query to select for your results*

**FROM**

*Identifies the table(s) your query will pull data from*

**WHERE**

*(Optional) Specifies record-filtering criteria for filtering your results*

**GROUP BY**

*(Optional) Specifies how to group the data in your results*

**HAVING**

*(Optional) Specifies group-filtering criteria for filtering your results*

**ORDER BY**

*(Optional) Specifies the order in which your query results are displayed*



END OF  
STATEMENT

**SELECT columnName**

**FROM tableName**

**WHERE logicalCondition**

**GROUP BY columnName**

**HAVING logicalCondition**

**ORDER BY columnName**

# COMMON JOIN TYPES

## INNER JOIN

*Returns records that exist in **BOTH** tables, and excludes unmatched records from either table*

**FROM** leftTableName  
**INNER JOIN** rightTableName

## LEFT JOIN

*Returns ALL records from the **LEFT** table, and any matching records from the **RIGHT** table*

**FROM** leftTableName  
**LEFT JOIN** rightTableName

## RIGHT JOIN

*Returns ALL records from the **RIGHT** table, and any matching records from the **LEFT** table*

**FROM** leftTableName  
**RIGHT JOIN** rightTableName

## FULL OUTER JOIN

*Returns ALL records from **BOTH** tables, including non-matching records*

**FROM** leftTableName  
**FULL JOIN** rightTableName

## UNION

*Returns all data from one table, with all data from another table **appended to the end***

**SELECT FROM** firstTableName  
**UNION**  
**SELECT FROM** secondTableName

# DOWNLOAD & SETUP



# MySQL DOWNLOAD & SETUP – OVERVIEW

Step 1

## Download Community Server

*This allows SQL to run on your machine*

Step 2

## Download MySQL Workbench

*This is the program you'll use to write and run SQL queries (it's intuitive, and works across operating systems)*

Step 3

## Connect Workbench to Server

*We'll get you connected to the server so you can use Workbench to start running your own SQL queries*

Step 4

## Review Workbench Interface

*We'll take a quick tour of the Workbench interface to get you familiar with the layout and key components*

Step 5

## Create the Database

*We'll run the SQL code to build the database which we'll be exploring throughout the course (this part is easy!)*



### HEY THIS IS IMPORTANT!

If you took one of my other courses, and have already installed Community Server and MySQL Workbench, then you can skip ahead to creating the database. No need to re-install. Whatever version you have is great.

# STEP 1: COMMUNITY SERVER (MAC)



## HEY THIS IS IMPORTANT!

If you took one of my other courses, and have already installed Community Server and MySQL Workbench, then you can skip ahead to creating the database. No need to re-install. Whatever version you have is great.

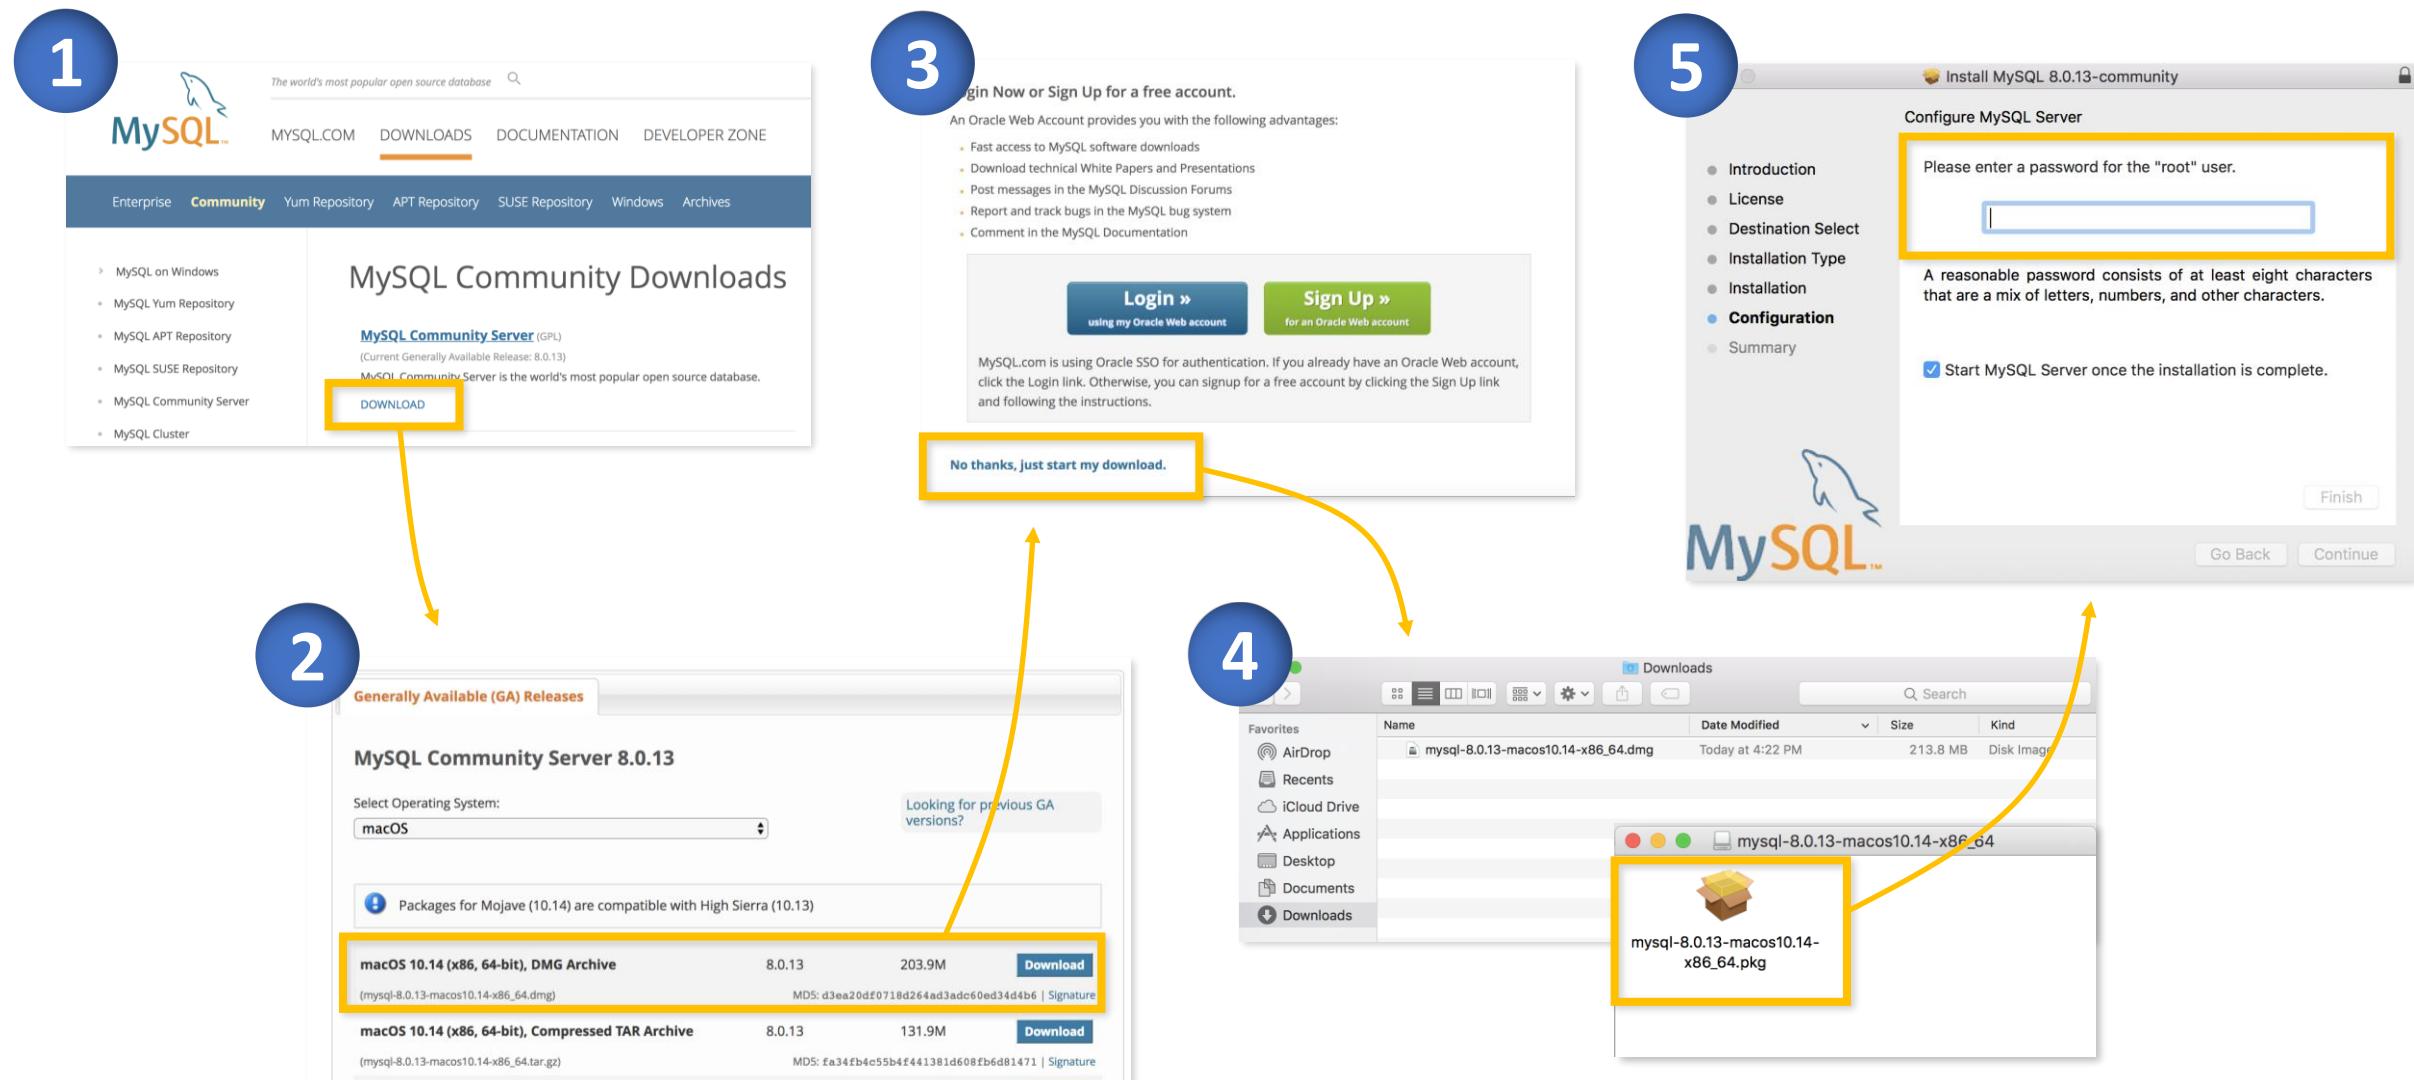
The screenshot shows the phpMyAdmin interface. On the left, there's a tree view of the database schema. In the center, a table named 'rental' is displayed with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Below it, a query results window shows a list of rows from the 'customer' table, ordered by the count of distinct rental IDs. On the right, a table list shows various tables like actor, address, city, country, customer, film, film\_actor, film\_category, inventory, and more, each with their respective details.

# MySQL COMMUNITY SERVER – MAC DOWNLOAD GUIDE



- 1 Go to <https://dev.mysql.com/downloads> and download **MySQL Community Server**
- 2 Select the **MacOS** operating system, and download the **DMG Archive** version
  - *Note: you'll likely see a later version than the one shown (just download the latest)*
- 3 No need to Login or Sign Up, just click "**No thanks, just start my download**"
- 4 Find the install file in your downloads, then double click to run the installer package
- 5 Click through each install step, leaving defaults unless you need customized settings
  - *Note: Make sure to store your **root password** somewhere, you'll need this later!*

# MySQL COMMUNITY SERVER – MAC DOWNLOAD GUIDE



# STEP 1: COMMUNITY SERVER (PC)



## HEY THIS IS IMPORTANT!

If you took one of my other courses, and have already installed Community Server and MySQL Workbench, then you can skip ahead to creating the database. No need to re-install. Whatever version you have is great.

The screenshot shows the phpMyAdmin interface with a blue header. On the left, there's a tree view of the database schema. The main area has four panes: 1) A table structure for 'rental' with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Data rows are listed from 1 to 10. 2) A query results pane showing rows 0-24 of a query that took 0.0212 seconds. 3) A SQL query editor with the following code:

```
1 SELECT
2     customer.first_name,
3     customer.last_name,
4     COUNT(DISTINCT rental_id) AS rental_count
5     FROM rental
6     LEFT JOIN customer
7     ON rental.customer_id = customer.customer_id
8     GROUP BY customer_id
9     ORDER BY
10    COUNT(DISTINCT rental_id) DESC
```

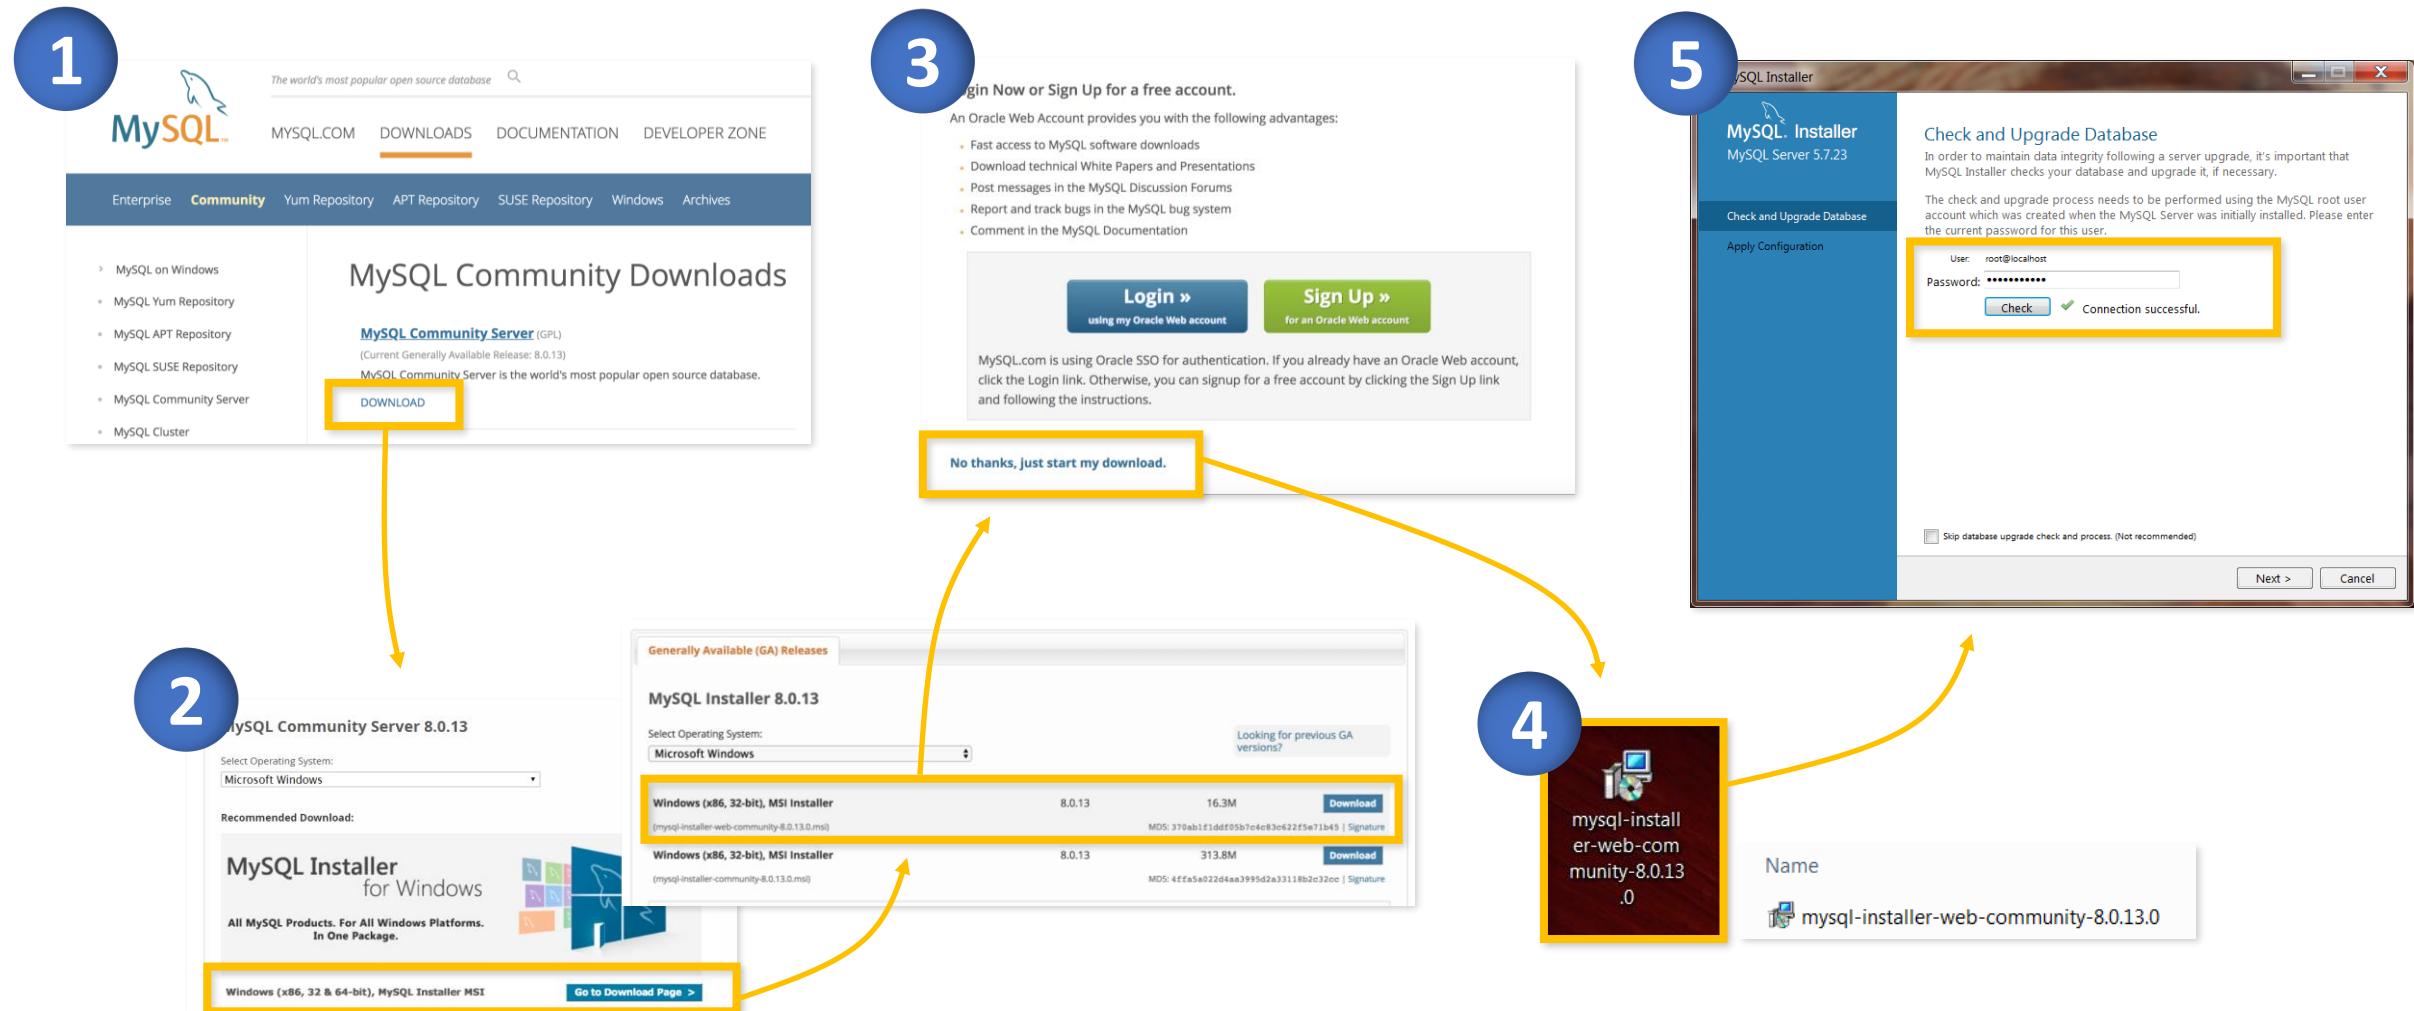
4) A privileges table showing various users and their permissions.

# MySQL COMMUNITY SERVER – PC DOWNLOAD GUIDE



- 1 Go to <https://dev.mysql.com/downloads> and download **MySQL Community Server**
- 2 Select the **Microsoft Windows** operating system, and the **Installer MSI** download
  - *Note: On the download page you may see two versions: select mysql-installer-web-community if you are connected to the internet, and keep in mind that you may see a later version than the one shown (just download the latest)*
- 3 No need to Login or Sign Up, just click “**No thanks, just start my download**”
- 4 Find the install file in your downloads, then double click to run the installer package
- 5 Click through each install step, leaving defaults unless you need customized settings
  - *Note: Make sure to store your **root password** somewhere, you'll need this later!*

# MySQL COMMUNITY SERVER – PC DOWNLOAD GUIDE



# STEP 2: MySQL WORKBENCH (MAC)



## HEY THIS IS IMPORTANT!

If you took one of my other courses, and have already installed Community Server and MySQL Workbench, then you can skip ahead to creating the database. No need to re-install. Whatever version you have is great.

The screenshot shows the MySQL Workbench interface. On the left, there's a tree view of a database schema with tables like actor, address, category, customer, film, film\_actor, film\_category, inventory, and rental. In the center, a query editor displays a SQL query and its results. The results show rows from the rental table, including rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. The total number of rows is 599. On the right, a table editor shows the structure of the rental table with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update.



# MySQL WORKBENCH – MAC DOWNLOAD GUIDE

- 1 Go to <https://dev.mysql.com/downloads/workbench>, scroll down to **Generally Available (GA) Releases**, and select the **MacOS** operating system
- 2 We'll be using version **8.0.16** for this course, so you can either click "**Looking for previous GA versions?**" to search for the same one, or simply download the latest available
- 3 No need to Login or Sign Up, just click "**No thanks, just start my download**"
- 4 Find the install file in your downloads, click the MySQL Workbench logo (*with the dolphin*) and drag it into your **Applications** folder
- 5 Look for MySQL workbench in your list of applications, double click to launch, then proceed to **Step 3: Connecting to the server**



# MySQL WORKBENCH – MAC DOWNLOAD GUIDE

1

The screenshot shows the MySQL Workbench 8.0.16 download page. A yellow box highlights the 'Select Operating System' dropdown menu where 'macOS' is chosen. Another yellow box highlights the 'Looking for previous GA versions?' link.

3

The screenshot shows the MySQL.com login or sign-up page. It includes a list of advantages for having an Oracle Web Account, two buttons ('Login' and 'Sign Up'), and a note about Oracle SSO. A yellow box highlights the 'No thanks, just start my download.' button.

5

The screenshot shows the Mac OS X Applications folder. A yellow arrow points from the 'MySQLWorkbench' icon in the list to the Applications folder icon in the sidebar.

2

*Look for version **8.0.16**, or download the latest*

The screenshot shows the MySQL Workbench 6.3.10 download page. A yellow box highlights the 'Select Version' dropdown menu where '6.3.10' is chosen. Another yellow box highlights the 'Looking for the latest GA version?' link.

4

The screenshot shows the MySQL Workbench 8.0 splash screen with instructions to drag the icon to the Applications folder. Two yellow boxes highlight the MySQLWorkbench icon and the Applications folder icon.

# STEP 2: MySQL WORKBENCH (PC)



## HEY THIS IS IMPORTANT!

If you took one of my other courses, and have already installed Community Server and MySQL Workbench, then you can skip ahead to creating the database. No need to re-install. Whatever version you have is great.

The screenshot shows the MySQL Workbench interface. On the left, there's a tree view of the database schema. In the center, a table named 'rental' is displayed with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Below the table is a preview of 24 rows of data. To the right, a query editor window shows a SELECT statement that retrieves the first name and last name of customers who have rented more than one item. The results of this query are also shown below the editor. At the bottom right, there's a list of tables and their details.

rental_id	rental_date	inventory_id	customer_id	return_date	staff_id	last_update
1	2005-05-24 22:53:30	367	130	2005-05-26 22:04:30	1	2006-02-15 21:30:53
2	2005-05-24 22:54:33	1525	459	2005-05-28 19:40:33	1	2006-02-15 21:30:53
3	2005-05-24 22:03:39	1711	408	2005-06-01 22:12:39	1	2006-02-15 21:30:53
4	2005-05-24 23:03:47	2455	333	2005-06-04 14:34:47	2	2006-02-15 21:30:53
5	2005-05-24 23:05:21	2079	222	2005-06-02 03:23:42	1	2006-02-15 21:30:53
6	2005-05-24 23:10:29	2792	549	2005-05-27 00:00:00	1	2006-02-15 21:30:53
7	2005-05-24 23:11:30	2079	222	2005-05-29 23:00:00	1	2006-02-15 21:30:53
8	2005-05-24 23:31:46	2346	239	2005-05-27 23:33:46	2	2006-02-15 21:30:53
9	2005-05-25 00:00:40	2580	128	2005-05-28 00:22:40	1	2006-02-15 21:30:53
10	2005-05-25 00:02:21	1824	399	2005-05-31 22:44:21	2	2006-02-15 21:30:53

```
✓ Showing rows 0 - 24 (599 total, Query took 0.0212 seconds.)  
1. SELECT  
2.     customer.first_name,  
3.     customer.last_name,  
4.     COUNT(DISTINCT rental_id) AS rentals  
5. FROM  
6.     rental  
7.     JOIN  
8.     customer  
9.     ON customer.customer_id = rental.customer_id  
10.    GROUP BY  
11.    customer.first_name  
12.    customer.last_name  
13.    ORDER BY  
14.        COUNT(DISTINCT rental_id) DESC
```

Table	Action	Rows	Type	Collation	Size	Overflow
actor	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	33,012 K	-
actor_info	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
address	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	84,412 K	-
category	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
city	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
country	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
customer	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	128,012 K	-
customer_address	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
customer_list	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	170,412 K	-
film	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	170,412 K	-
film_actor	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	170,412 K	-
film_category	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	89,612 K	-
film_text	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-
inventory	Browse  > Structure  > Search  > Insert  > Empty  > Drop	100	InnoDB	utf8_general_ci	16,412 K	-

# MySQL WORKBENCH – PC DOWNLOAD GUIDE

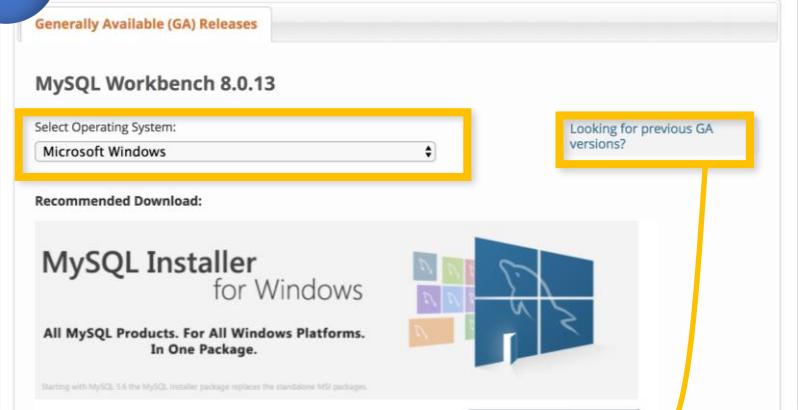


- 1 Go to <https://dev.mysql.com/downloads/workbench>, scroll down to **Generally Available (GA) Releases**, and select the **Microsoft Windows** operating system
- 2 We'll be using version **8.0.13** for this course, so you can either click "***Looking for previous GA versions?***" to search for the same one, or simply download the latest available
- 3 No need to Login or Sign Up, just click "***No thanks, just start my download***"
- 4 Find the install file in your downloads, double click to run the installation process, and stick with default settings unless you need a custom configuration
- 5 Look for MySQL workbench in your list of programs, double click to launch, then proceed to ***Step 3: Connecting to the server***
  - ***Note:*** You may see a warning if you aren't on **Windows 10+**, but most older systems (i.e. Windows 7) should be compatible

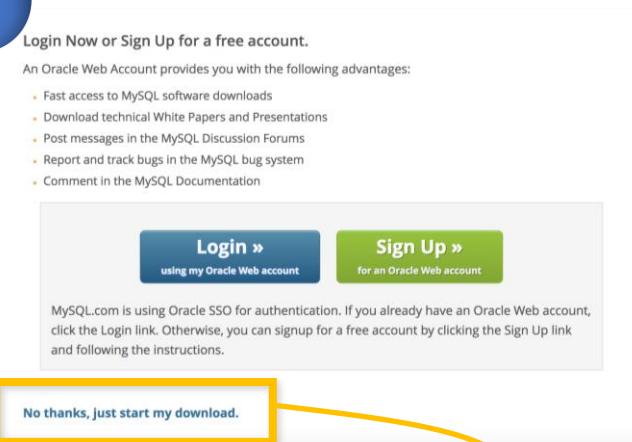


# MySQL WORKBENCH – PC DOWNLOAD GUIDE

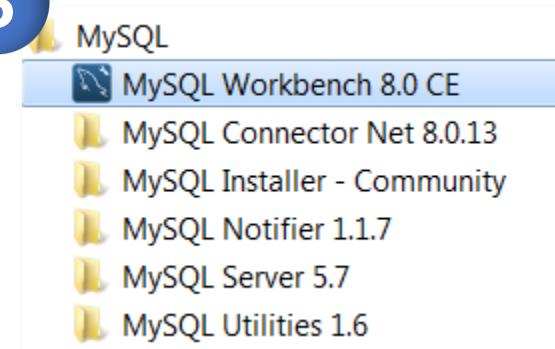
1



3

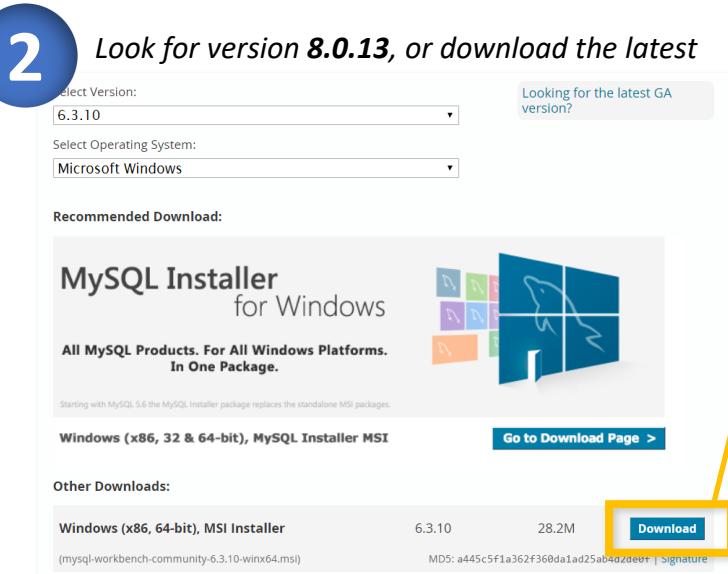


5

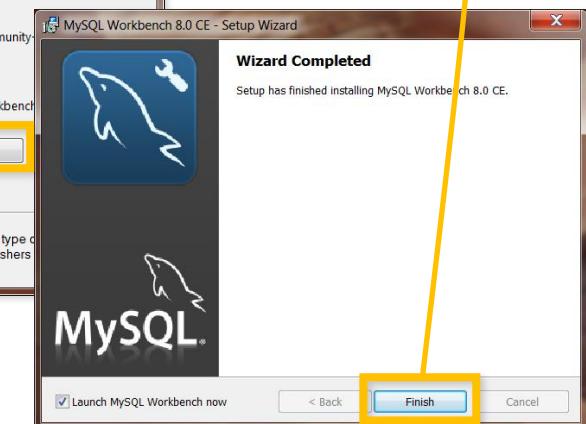
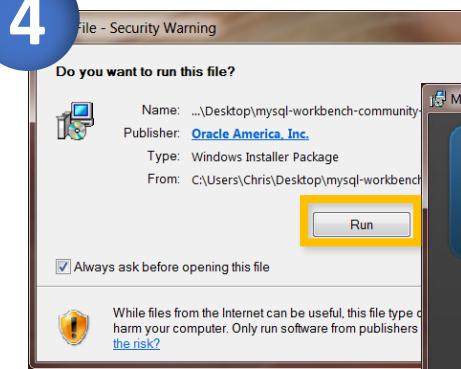


2

*Look for version 8.0.13, or download the latest*



4



# STEP 3: CONNECTING TO THE SERVER

The image shows the phpMyAdmin interface with a blue header. The main area contains four panels:

- Left Panel (Database Structure):** A tree view of the database schema.
- Top Center Panel (Table Data):** A grid showing data from the 'rental' table. The columns are: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. The data includes rows from May 24, 2005, to May 25, 2005.
- Middle Center Panel (Query Results):** A grid showing the results of a query. The columns are: customer.first\_name, customer.last\_name, COUNT(DISTINCT rental\_id) AS rentals. The results show the number of rentals for each customer.
- Right Panel (Table List):** A list of tables in the database, including actor, actor.info, address, category, customer, film, film\_actor, film\_category, inventory, and rental. Each entry has a 'Browse' button, a 'Structure' button, and a 'Search' button.

# CONNECTING TO THE SERVER

- 1 After launching Workbench, check the **MySQL Connections** section on the welcome page
  - *If you see a connection already, right-click to **Edit Connection**, otherwise click the **plus sign (+)** to add a new one*
- 2 Name the connection “**mavenmovies**”, confirm that the Username is “**root**”, and click **OK**
- 3 Once you see the **mavenmovies** connection on your welcome screen, simply click the tile and enter your **root password** to complete the connection

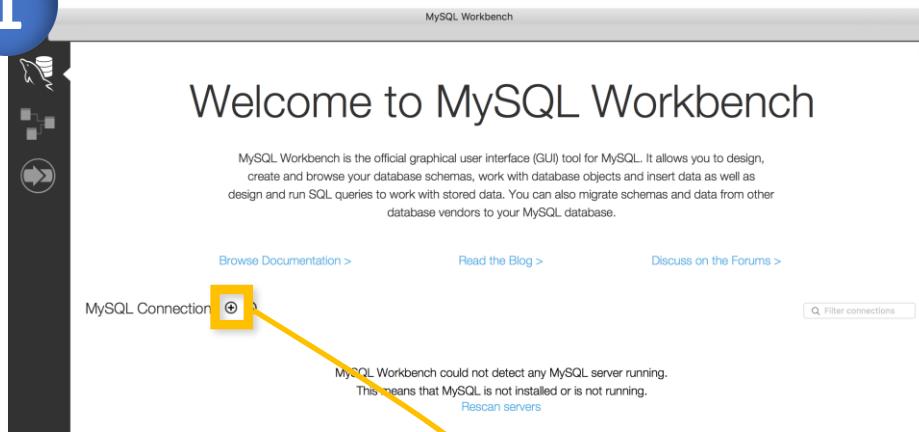


## Fun Fact!

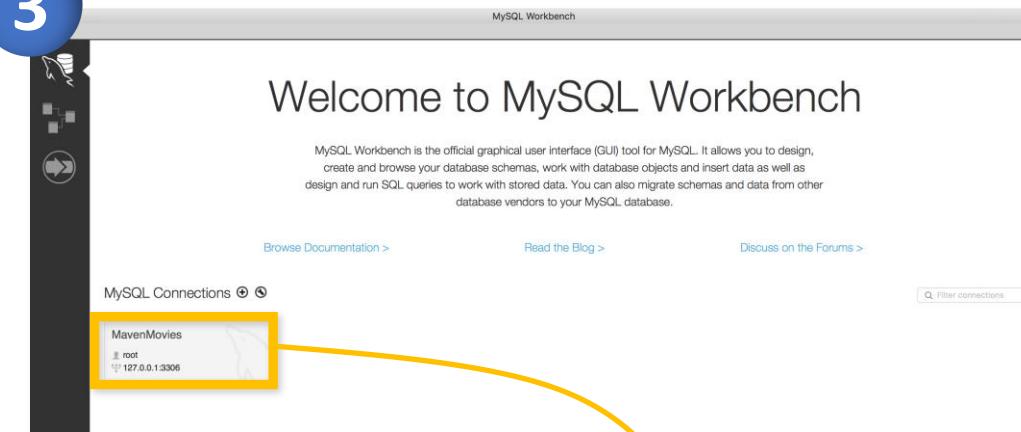
Maven Movies is the name of the database I used when I made my first course. I always name my connections ‘mavenmovies’ as tribute. **It does not matter what you name your connection. Name it anything you want!**

# CONNECTING TO THE SERVER

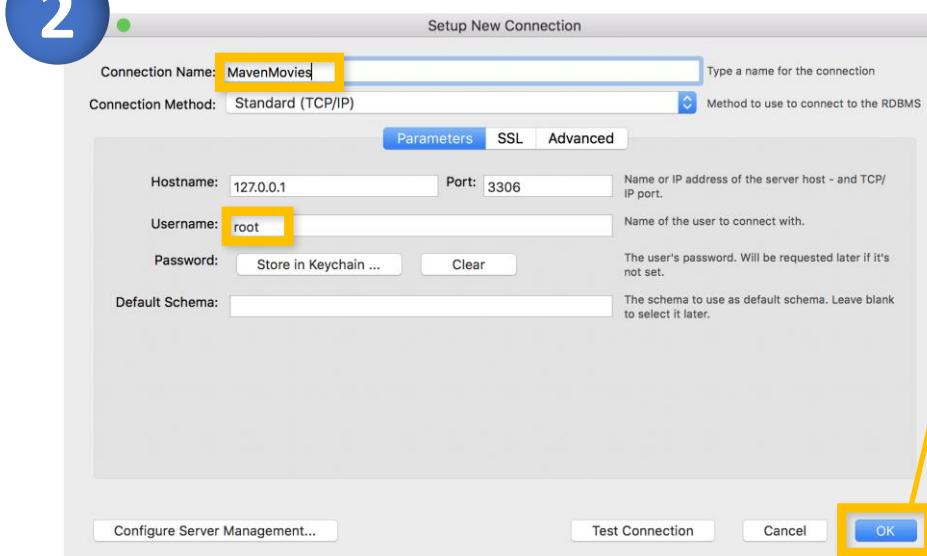
1



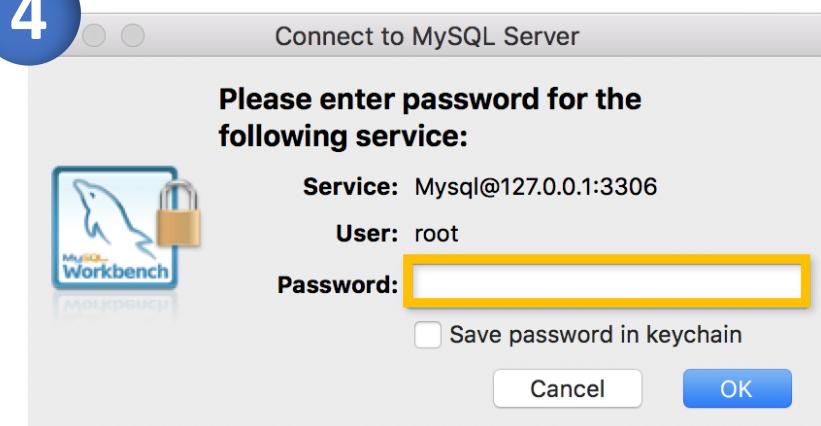
3



2



4



# STEP 4: MySQL WORKBENCH INTERFACE

The image shows a composite screenshot of the MySQL Workbench interface. On the left, there is a large, semi-transparent watermark of the MySQL Workbench interface itself, showing various tabs like 'Connections', 'Schemas', 'Tables', 'Data', 'Structure', 'SQL', and 'Logs'. Overlaid on this watermark are three main windows:

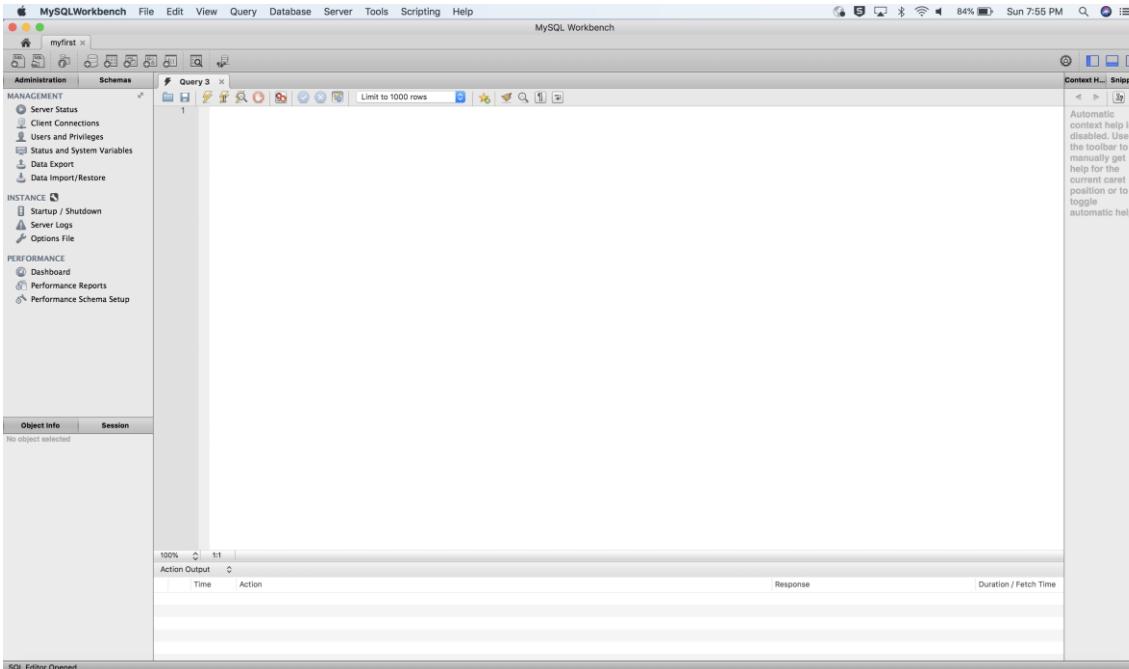
- phpMyAdmin Session:** A window titled 'phpMyAdmin' showing a database schema diagram with tables like actor, address, category, film, film\_actor, film\_text, and inventory. Below the diagram is a table named 'rental' with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Data rows are listed from 1 to 10.
- Query Editor:** A window titled 'Showing rows 0 - 24 (599 total, Query took 0.0212 seconds.)' containing the following SQL query:

```
1 SELECT
2   customer.first_name,
3   customer.last_name
4   COUNT(DISTINCT rental_id) AS rental_count
5   FROM inventory
6   LEFT JOIN rental
7   ON inventory.inventory_id = rental.inventory_id
8   GROUP BY
9   customer.first_name,
10  customer.last_name
11 ORDER BY
12  COUNT(DISTINCT rental_id) DESC
```
- Table Browser:** A window titled 'Tables' showing a list of tables with their respective structures and actions (Browse, Structure, Search, Insert, Drop, View). Tables listed include actor, actor\_info, address, category, film, film\_actor, film\_text, and inventory.

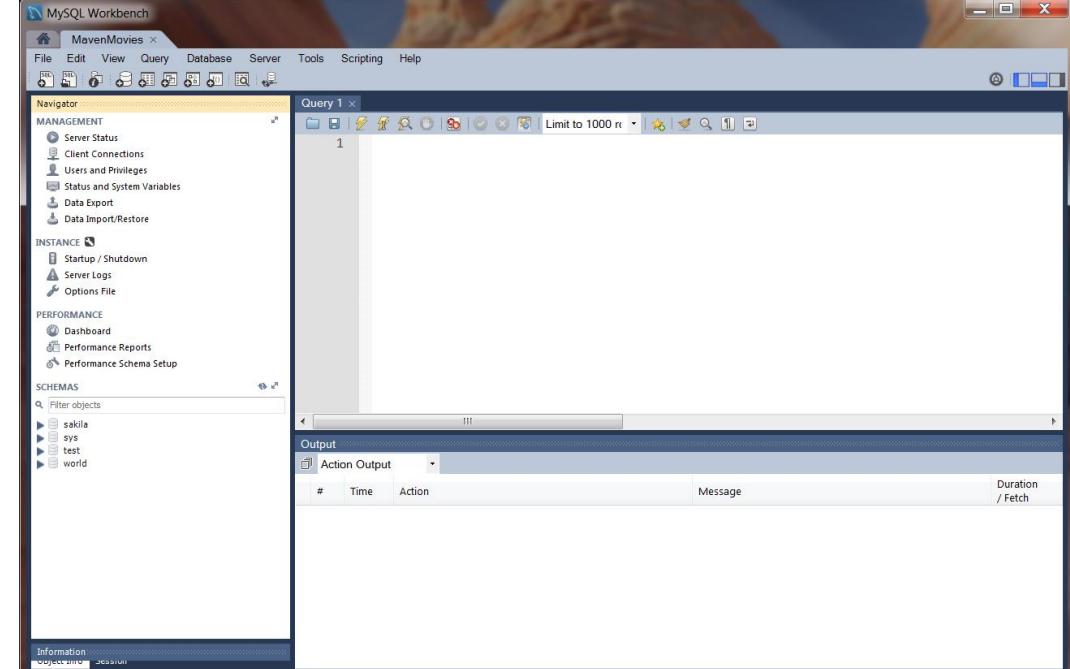
# MySQL WORKBENCH INTERFACE (MAC VS. PC)



*Mac interface*



*PC interface*



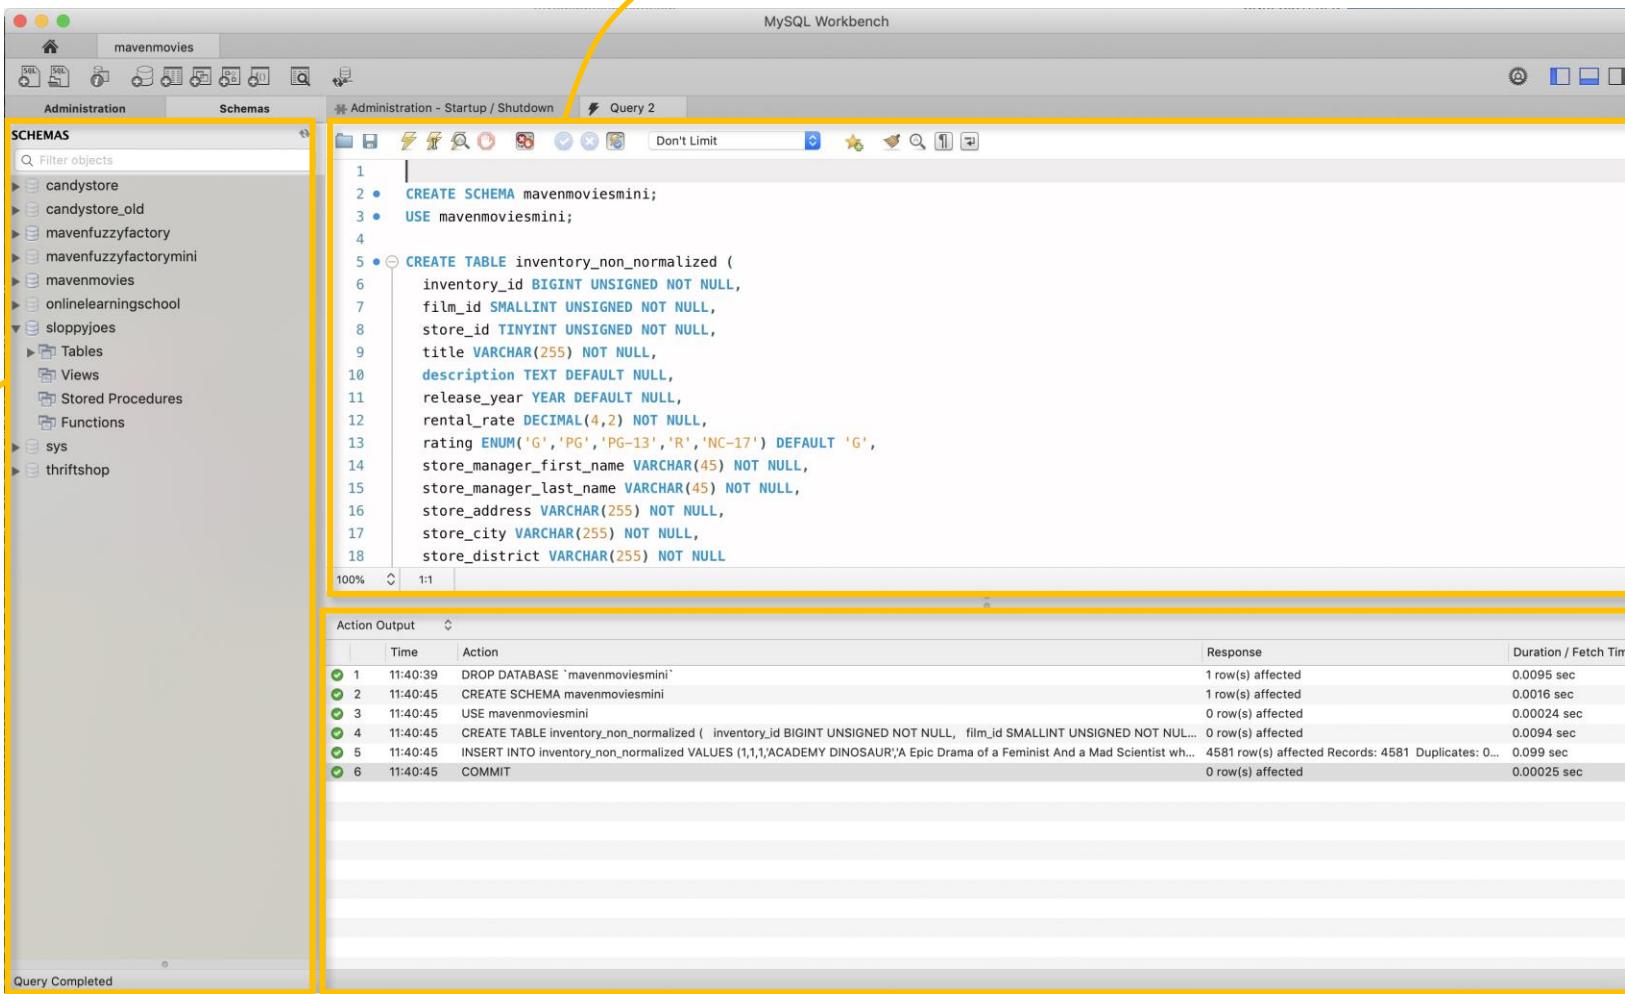
## HEY THIS IS IMPORTANT!

Workbench looks slightly different on **Mac** vs. **PC**, but everything you need is found in the same place. While the course is recorded on a Mac, but you should have no problem keeping up on a PC

# QUICK TOUR: THE WORKBENCH INTERFACE

## Query Editor Window

*This is where you write and run your code*



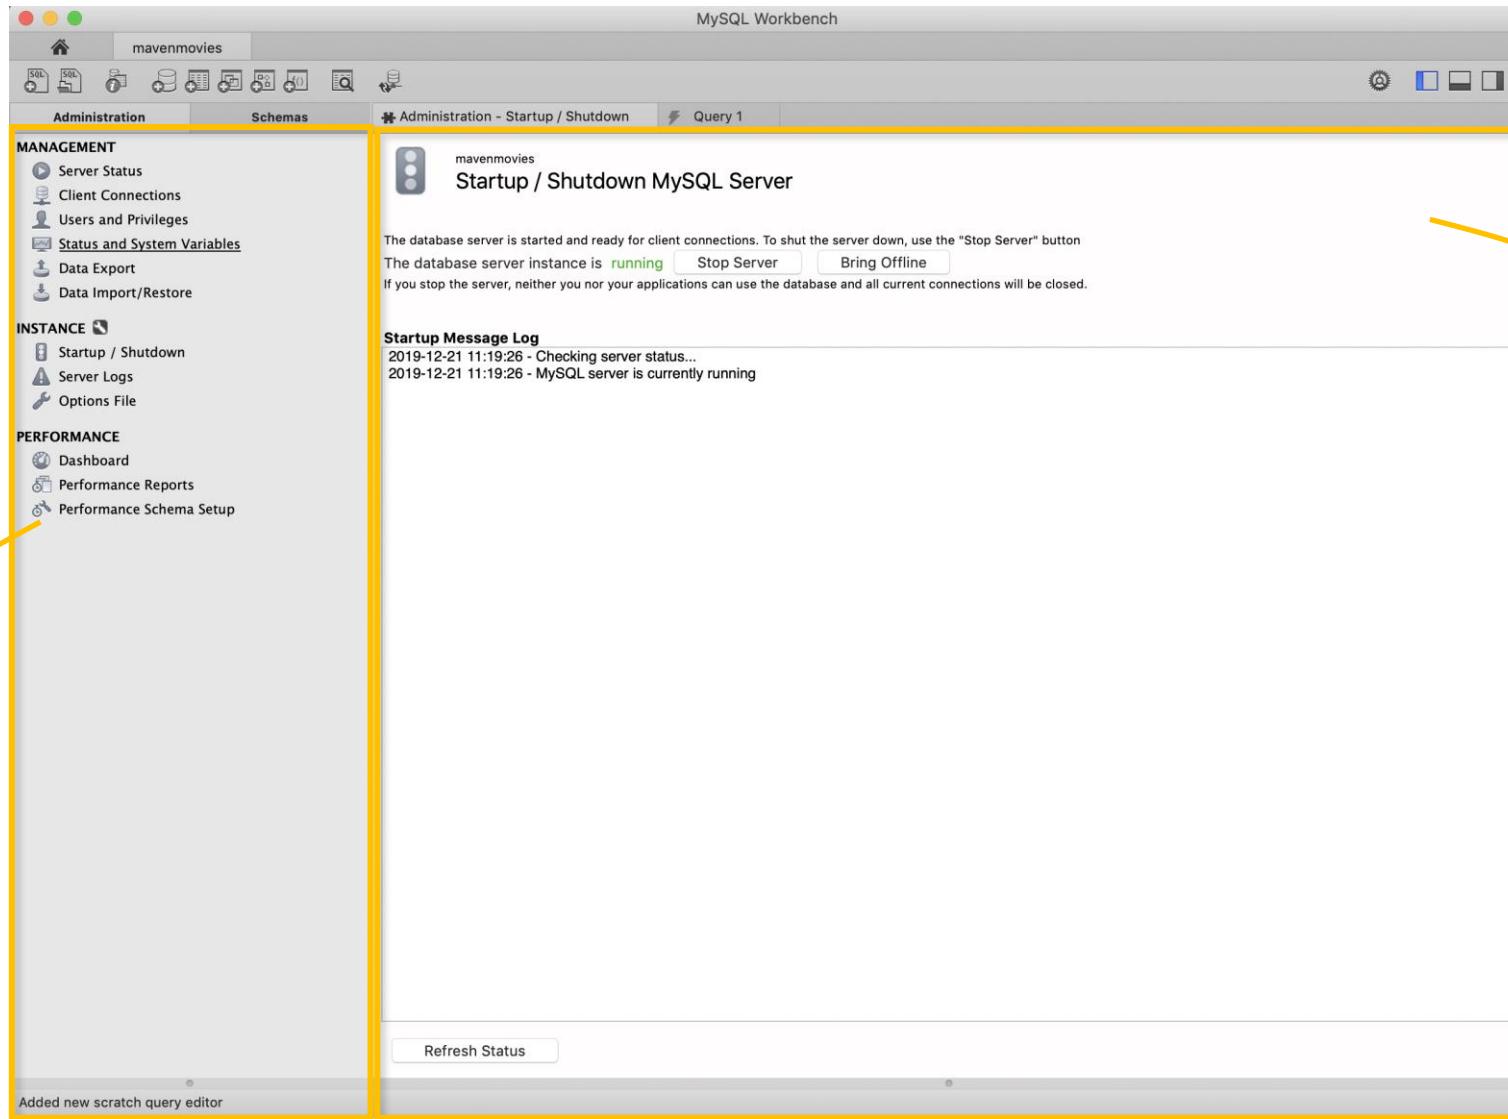
## Schemas Tab

*Here you can view tables and views in your schemas*

## Action Output

*This is a summary of actions taken by the server (TIP: the Response column is great for troubleshooting!)*

# QUICK TOUR: THE WORKBENCH INTERFACE



## Administration Tab

Here you can select the Management, Instance, and Performance tools

## Tool Tabs

This is where the various Management, Instance, and Performance tools show up after they are selected in the Administration Tab.

Note: these show up as additional tabs alongside tabs you have open for SQL query editing. They co-exist in the same section of Workbench, even though their functions are different

# IMPORTANT: PREPARE WORKBENCH SETTINGS



## THIS IS IMPORTANT!

Before we get started, we will quickly configure a few settings in Workbench which will dramatically cut down on problems you might otherwise face later in the course. Please do not skip this step.

- 1 First, we'll adjust a date setting, so we won't run into issues with our dates
- 2 Next, we'll update the GROUP BY setting to only allow FULL GROUP BY
- 3 Then, we'll adjust our MAX ALLOWED PACKET size, so any larger files we use will work
- 4 Last, we will shut down Workbench and restart, which will launch our new settings

# INSERTING DATA INTO THE DATABASE

The screenshot displays the phpMyAdmin interface with several panels:

- Left Panel:** Shows the database schema with tables like actor, address, category, customer, film, film\_actor, film\_category, inventory, and staff.
- Top Center:** A search bar with the placeholder "Search" and a dropdown menu for "Tables".
- Middle Left:** A table titled "rental" with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. It contains 10 rows of data.
- Middle Right:** A query results panel showing the output of a SELECT query. The results are identical to the data in the rental table.
- Bottom Right:** A table showing the structure of the "rental" table with columns: id, type, collation, size, and overhead.

# REVIEW OF CREATE, INSERT, & DELETE

- 1 First, we'll review CREATE SCHEMA to get started
- 2 Then, we'll review using CREATE TABLE to add a table to the schema
- 3 Next, we'll review the INSERT statement, used to add records to tables
- 4 Finally, we'll review the DELETE statement, used to remove records from tables



## THIS IS REVIEW!

We covered the basics of CREATE, INSERT, & DELETE in our MySQL Database Administration for Beginners course. If you already have these concepts covered, feel free to skip ahead to the lecture on Importing Data From a File.

# REVIEW OF CREATE SCHEMA

The screenshot shows the phpMyAdmin interface with several panels:

- Left Panel:** A tree view of the database schema.
- Top Center:** A table named "rental" with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, last\_update. It contains 10 rows of data.
- Middle Center:** A query results grid showing the output of a query: "SELECT customer.first\_name, customer.last\_name, COUNT(DISTINCT rental\_id) AS rental FROM customer JOIN rental ON customer.customer\_id = rental.staff\_id GROUP BY customer.first\_name, customer.last\_name ORDER BY COUNT(DISTINCT rental\_id) DESC". Rows 0-24 are displayed.
- Right Panel:** A table named "privileges" listing various database objects and their privileges.



## THIS IS REVIEW!

We covered the basics of creating schemas in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.

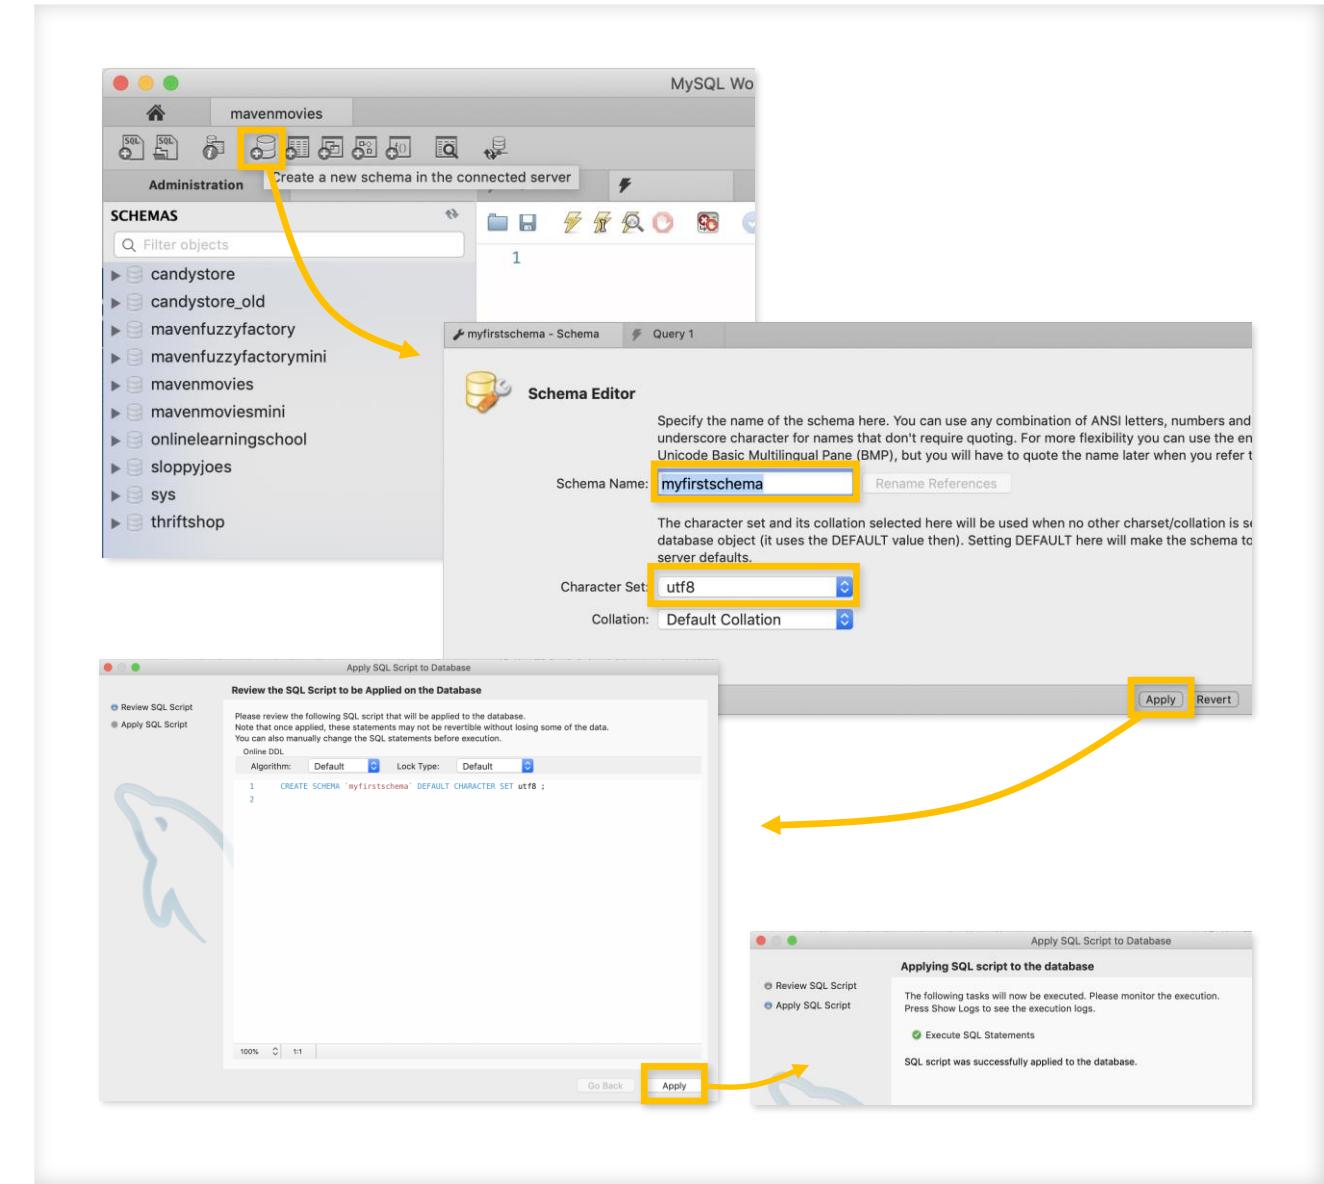
# CREATE SCHEMA W/ UI

- In MySQL, the terms ‘Schema’ and ‘Database’ will be used interchangeably to discuss a collection of data tables which relate to one another

We can create a new schema using Workbench’s user interface tools, with just a few quick clicks

- Whether we create schemas using the UI tools or using code is a matter of preference

## MySQL WORKBENCH TOOLS IN ACTION:



# REVIEW OF CREATE TABLE

The screenshot shows the phpMyAdmin interface with a blue header. On the left, there's a database schema diagram. The main area has three tabs: 'Structure' (selected), 'Query', and 'Privileges'. In the 'Structure' tab, there's a table named 'rental' with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Below it is a query results table with 24 rows of data. In the 'Query' tab, a SQL query is shown:

```
SELECT customer.first_name, customer.last_name, COUNT(DISTINCT rental_id) DESC
FROM rental
JOIN customer
ON customer.customer_id = rental.customer_id
GROUP BY customer.first_name, customer.last_name
ORDER BY COUNT(DISTINCT rental_id) DESC
```

The 'Privileges' tab lists various database objects with their actions (Browse, Structure, Search, Insert, Update, Delete, Drop, View) and details like Type (view), Collation (utf8\_general\_ci), and Size (various sizes).



## THIS IS REVIEW!

We covered the basics of creating tables in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.

# CREATE TABLE W/ CODE

- After our schema is created, we can begin populating our schema with tables that will contain records of data

- Creating a new table with code typically involves:

- Selecting a schema to add the table to
- Giving the new table a name
- Specifying which columns to include
- Setting data formats and types

## MySQL QUERY IN ACTION:

```
1
2 • USE myfirstcodeschema; -- specifying the schema to use
3
4 • CREATE TABLE myfirstcodetable (
5     myfirstcodetable_id BIGINT NOT NULL,
6     my_character_field VARCHAR(50),
7     my_text_field TEXT,
8     my_created_at TIMESTAMP
9 );
10
```

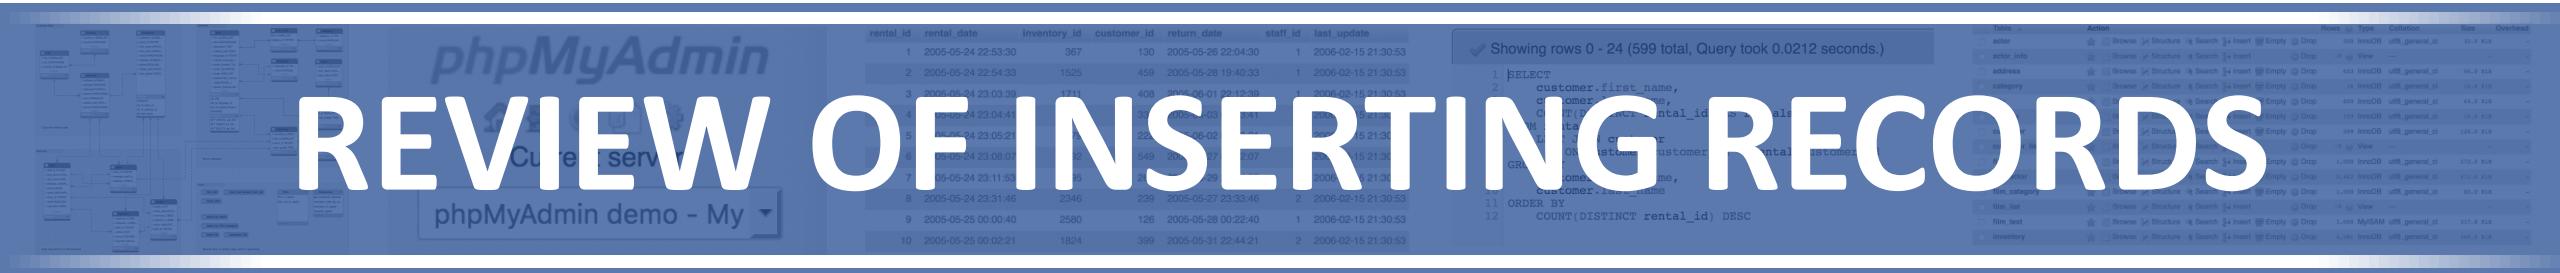
## ACTION OUTPUT:

Action Output			
	Time	Action	Response
1	14:59:51	USE myfirstcodeschema	0 row(s) affected
2	14:59:54	CREATE TABLE myfirstcodetable ( myfirstcodetable_i...	0 row(s) affected

# REVIEW OF INSERTING RECORDS

 THIS IS REVIEW!

We covered the basics of inserting records into tables in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.



# INSERTING NEW RECORDS

- MySQL makes it easy to insert additional records into a table using an **INSERT INTO** statement

- We can insert one row at a time, or we can batch insert many rows in one single SQL statement

- We may choose to specify values for every column in our **INSERT INTO** statement, or we may omit values for some columns. If we omit values, MySQL places the default value for that column

## MySQL QUERY IN ACTION:

```
USE thriftshop; -- sets the schema
```

```
SELECT * FROM inventory; -- just to see what's already there
```

```
INSERT INTO inventory (inventory_id, item_name, number_in_stock) VALUES  
(10, 'wolf skin hat', 1);
```

## (Another) MySQL QUERY IN ACTION:

```
-- this time we will batch insert
```

```
-- we will also not name the columns (this is optional)
```

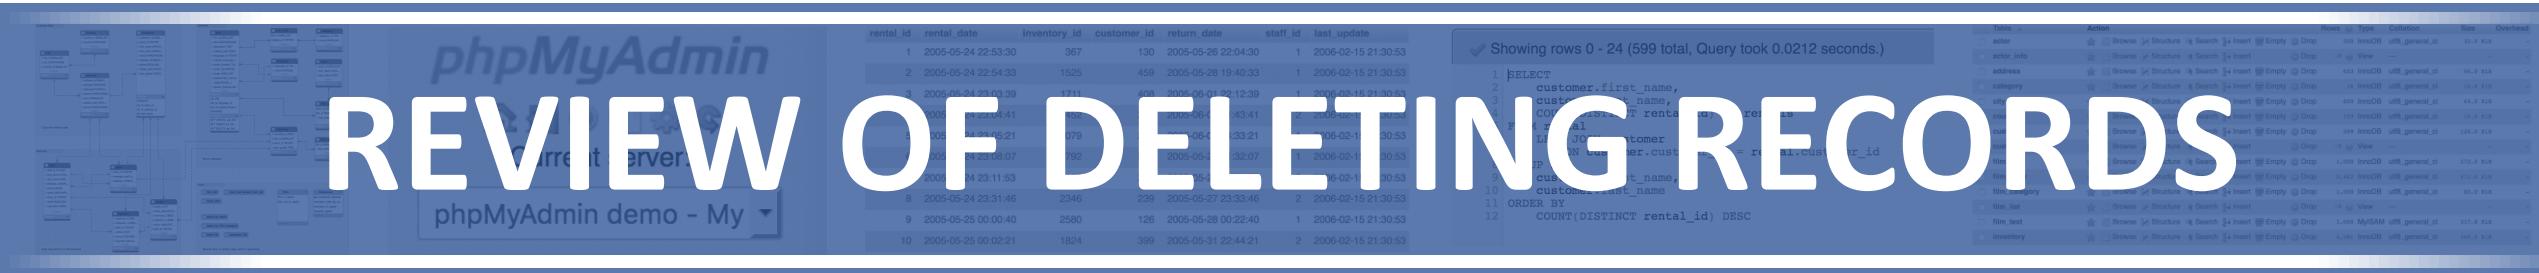
```
INSERT INTO inventory VALUES
```

```
(11, 'fur fox skin', 1),  
(12, 'plaid button up shirt', 8),  
(13, 'flannel zebra jammies', 6);
```

# REVIEW OF DELETING RECORDS

 THIS IS REVIEW!

We covered the basics of deleting records into tables in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.



# DELETE RECORDS

- MySQL allows us to target one or more records to **DELETE FROM** a given data table
- We will do this using the **DELETE FROM** statement, followed by the tablename
- As with **UPDATE** statements, we will usually want to specify which records to delete using a **WHERE** clause
- If we do not specify a **WHERE** clause, **DELETE FROM** will delete ALL records

## MySQL QUERY IN ACTION:

```
-- deleting ski blankets, no longer stocked  
DELETE FROM inventory  
WHERE inventory_id = 7
```

## ACTION OUTPUT:

inventory_id	item_name	number_in_stock
1	fur coat	0
2	moccassins	4
3	velour jumpsuit	12
4	house slippers	6
5	brown leather jacket	3
6	broken keyboard	6
7	ski blanket	1
8	kneeboard	2
9	pro wings sneakers	0
10	wolf skin hat	1
11	fur fox skin	1
12	plaid button up shirt	8
13	flannel zebra jammies	6

inventory_id	item_name	number_in_stock
1	fur coat	0
2	moccassins	4
3	velour jumpsuit	12
4	house slippers	6
5	brown leather jacket	3
6	broken keyboard	6
7	ski blanket	1
8	kneeboard	2
9	pro wings sneakers	0
10	wolf skin hat	1
11	fur fox skin	1
12	plaid button up shirt	8
13	flannel zebra jammies	6

# IMPORT DATA FROM A FILE

- Sometimes we will want to import a large amount of data at once
- Writing an `INSERT` statement and specifying every row of a large data set would be very tedious work

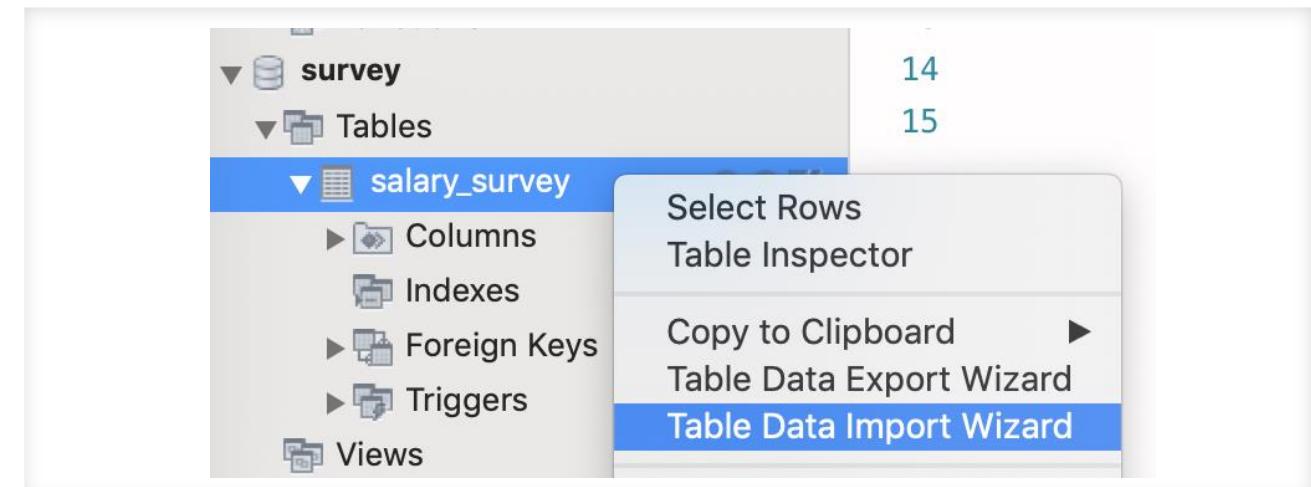
Luckily, MySQL Workbench provides a couple of ways to import data in bulk

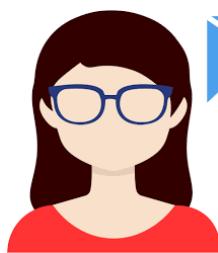
- The method we will primarily be using is the Table Data Import Wizard

## MySQL QUERY IN ACTION:

```
1 • CREATE SCHEMA survey;  
2  
3 • USE survey;  
4  
5 • CREATE TABLE salary_survey (  
6   country VARCHAR(50),  
7   years_experience INT,  
8   employment_status VARCHAR(50),  
9   job_title VARCHAR(50),  
10  is_manager VARCHAR(50),  
11  education_level VARCHAR(50)  
12 );
```

## Accessing the Table Data Import Wizard:





## NEW MESSAGE

April 2, 2012

From: **Sally Bleu (CEO)**

Subject: **Help get order data into the database**

Good morning,

I'm excited to get going on the data centralization project! As a first step, let's get the items customers are ordering stored in the database.

I've attached 2 files from our sales software. Can you create a table and import this data into our database?

Thanks!

-Sally



01.order\_items\_2012\_Mar  
02.order\_items\_2012\_Apr

Reply

Forward

## Helpful Hints

-- first, you'll want to create a new schema

```
CREATE SCHEMA mavenbearbuilders;
```

-- make sure to set it as the default schema too

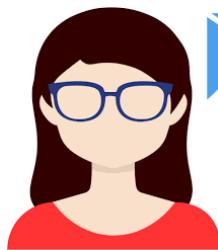
```
USE mavenbearbuilders;
```

-- then, create a new table, order\_items with columns...

- order\_item\_id
- created\_at
- order\_id
- price\_usd
- cogs\_usd
- website\_session\_id

-- if you get stuck on the syntax, check out the review videos

# TEST YOUR SKILLS: IMPORT ORDER ITEM DATA



## NEW MESSAGE

April 2, 2012

From: **Sally Bleu (CEO)**

Subject: **Help get order data into the database**

Good morning,

I'm excited to get going on the data centralization project! As a first step, let's get the items customers are ordering stored in the database.

I've attached 2 files from our sales software. Can you create a table and import this data into our database?

Thanks!

-Sally



01.order\_items\_2012\_Mar  
02.order\_items\_2012\_Apr

Reply

Forward

## Solution

```
29      -- if you've done this right, you'll have 159 records
30 •   SELECT COUNT(*) AS total_records FROM order_items;
```

100% 1:18

Result Grid



Filter Rows:

Search

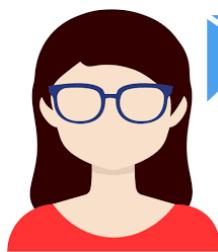
Export:



totalRecords
159

▶ 159

# TEST YOUR SKILLS: IMPORT ORDER ITEM DATA



## NEW MESSAGE

April 5, 2012

From: **Sally Bleu (CEO)**

Subject: **Can we aggregate refunds too?**

Hey there,

Now that we have **order\_items** built out, could you also import this attached April refund data (weren't any in March) in a new table called **order\_item\_refunds**?

It will be great to start tracking refund rates better so we can keep an eye on product quality.

Thanks!

-Sally



03.order\_item\_refunds\_2012\_Apr

Reply

Forward

## Helpful Hints

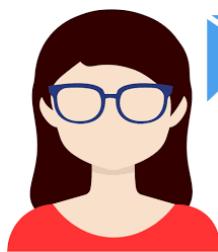
-- make sure to set it as the default schema

**USE mavenbearbuilders;**

-- then, create a new table, **order\_item\_refunds** with columns...

- **order\_item\_refund\_id**
- **created\_at**
- **order\_item\_id**
- **order\_id**
- **refund\_amount\_usd**

# TEST YOUR SKILLS: IMPORT REFUND DATA



## NEW MESSAGE

April 5, 2012

From: **Sally Bleu (CEO)**

Subject: **Can we aggregate refunds too?**

Hey there,

Now that we have **order\_items** built out, could you also import this attached April refund data (weren't any in March) in a new table called **order\_item\_refunds**?

It will be great to start tracking refund rates better so we can keep an eye on product quality.

Thanks!

-Sally



03.order\_item\_refunds\_2012\_Apr

Reply

Forward

## Solution

```
22      -- if you've done this right, you'll have 10 records
23 •  SELECT COUNT(*) AS total_records FROM order_item_refunds;
```

100% 3:19

Result Grid



Filter Rows:

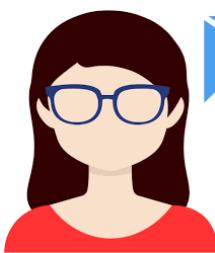
Search

Export:



total_records
10

# TEST YOUR SKILLS: IMPORT REFUND DATA



## NEW MESSAGE

April 8, 2012

From: **Sally Bleu (CEO)**

Subject: **Help Cleaning Up Some Bad Data**

Well...

Turns out the new guy messed up some of our data. He flagged order\_items **131, 132, 145, 151, and 153** as refunds, but they were actually customer inquiries.

Can you remove these from **order\_item\_refunds** to clean up our data?

Thanks!

-Sally

Reply

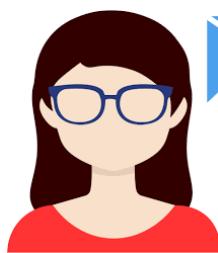
Forward

## Helpful Hints

```
-- I like to do a simple SELECT * to see what's in there  
SELECT * FROM order_item_refunds;
```

```
-- TIP: be careful deleting without using the PK
```

# TEST YOUR SKILLS: DELETING RECORDS



## NEW MESSAGE

April 8, 2012

From: **Sally Bleu (CEO)**

Subject: **Help Cleaning Up Some Bad Data**

Well...

Turns out the new guy messed up some of our data. He flagged order\_items **131, 132, 145, 151, and 153** as refunds, but they were actually customer inquiries.

Can you remove these from **order\_item\_refunds** to clean up our data?

Thanks!

-Sally

Reply

Forward

## Solution

```
31      -- if you've done this right, you'll have 5 records remaining
32 •  SELECT COUNT(*) AS total_records FROM order_item_refunds;
```

100% 23:27

Result Grid



Filter Rows:

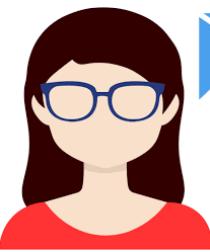
Search

Export:

total_records
5

▶ 5

# TEST YOUR SKILLS: DELETING RECORDS



## NEW MESSAGE

January 02, 2013

From: **Sally Bleu (CEO)**

Subject: **Time to pull in all 2012 data**

Hey,

The business finished 2012 strong, and I would like to get all of our **order\_items** and **order\_item\_refunds** data updated through the end of the year.

Can you help me import the 2 attached files into the correct tables?

Thanks!

-Sally



[04.order\\_items\\_2012\\_May-Dec](#)  
[05.order\\_item\\_refunds\\_May-Dec](#)

Reply

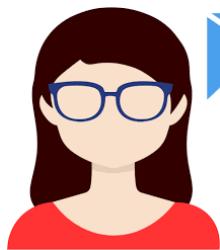
Forward

## Helpful Hint

-- first, import the `order_items` data

-- then, import the `order_item_refunds` data

# TEST YOUR SKILLS: IMPORT REMAINING 2012 DATA



## NEW MESSAGE

January 02, 2013

From: **Sally Bleu (CEO)**

Subject: **Time to pull in all 2012 data**

Hey,

The business finished 2012 strong, and I would like to get all of our **order\_items** and **order\_item\_refunds** data updated through the end of the year.

Can you help me import the 2 attached files into the correct tables?

Thanks!

-Sally



04.order\_items\_2012\_May-Dec  
05.order\_item\_refunds\_May-Dec

Reply

Forward

## Solution

```
-- if you did this right, you should see 2,586 TOTAL order_item records
-- note: this includes records from the previous assignment
```

```
SELECT COUNT(*) AS total_records FROM order_items;
```

```
-- if you did this right, you should see 169 TOTAL refund records
-- note: this includes records from the previous assignment
-- note: if you have more, you may have missed the DELETE assignment
```

```
SELECT COUNT(*) AS total_records FROM order_item_refunds;
```

# TEST YOUR SKILLS: IMPORT REMAINING 2012 DATA

# LOAD DATA

- Another method for importing large datasets is to use LOAD DATA
- One major benefit is that using LOAD DATA can be faster when your data sets get large

Another advantage is that Engineers can automate scripts that contain LOAD DATA scripts

## NOT FOR EVERYONE!

LOAD DATA is not enabled by default, and enabling it requires editing your configuration file using command line. This is too technical for many in our audience, so we are not going to teach it in detail.

## MySQL LOAD DATA SYNTAX

### 13.2.7 LOAD DATA Statement

```
1 LOAD DATA
2   [LOW_PRIORITY | CONCURRENT] [LOCAL]
3   INFILE 'file_name'
4   [REPLACE | IGNORE]
5   INTO TABLE tbl_name
6   [PARTITION (partition_name [, partition_name] ...)]
7   [CHARACTER SET charset_name]
8   [{FIELDS | COLUMNS}
9     [TERMINATED BY 'string']
10    [[OPTIONALLY] ENCLOSED BY 'char']
11    [ESCAPED BY 'char']
12  ]
13  [LINES
14    [STARTING BY 'string']
15    [TERMINATED BY 'string']
16  ]
17  [IGNORE number {LINES | ROWS}]
18  [(col_name_or_user_var
19    [, col_name_or_user_var] ...)]
20  [SET col_name={expr | DEFAULT}
21    [, col_name={expr | DEFAULT}] ...]
```

# ALTER, UPDATE, & KEYS

The screenshot displays the phpMyAdmin interface with a blue header and sidebar. The sidebar on the left shows the database schema with tables like actor, address, category, customer, film, film\_actor, film\_category, inventory, language, payment, rental, staff, and store. The main area has three panes:

- Query Results:** Shows a table of rental records with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. The data includes rows from May 24, 2005, to May 25, 2005.
- Query Editor:** Displays a SQL query that counts distinct rental IDs for each customer and orders them by count in descending order. The query is:

```
SELECT
    customer.customer_id,
    customer.last_name,
    COUNT(DISTINCT rental.rental_id) AS rental_count
FROM rental
LEFT JOIN customer
    ON rental.customer_id = customer.customer_id
GROUP BY
    customer.customer_id,
    customer.last_name
ORDER BY
    COUNT(DISTINCT rental_id) DESC
```
- Status Bar:** Shows the status "Showing rows 0 - 24 (599 total, Query took 0.0212 seconds.)" and a table of database statistics at the bottom right.

# REVIEW OF ALTER, UPDATE, AND KEY RELATIONSHIPS

- 1 First, we'll review the ALTER TABLE statement used for adding and removing columns
- 2 Then, we'll review using UPDATE to SET values of records already in the database
- 3 Next, we'll review Cardinality and Primary Key to Foreign Key relationships
- 4 After our review is complete, we'll use each of these concepts to update the database as the business rolls out a new product and data tracking needs become more complex



## THIS IS REVIEW!

We covered the basics of ALTER TABLE, UPDATE and Primary and Foreign Key relationships in our MySQL Database Administration for Beginners course. If you already have these concepts covered, feel free to skip ahead.

# REVIEW OF ALTER TABLE

The screenshot shows the phpMyAdmin interface with three main panels:

- Left Panel:** Database schema diagram showing tables like actor, address, category, customer, film, film\_actor, film\_category, inventory, and rental.
- Middle Panel:** A query results page titled "Showing rows 0 - 24 (599 total, Query took 0.0212 seconds.)". It displays a table with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. The data shows various rental records from May 2005 to June 2006.
- Right Panel:** A privileges page listing tables and their actions (Browse, Structure, Search, Insert, Update, Delete). It includes tables like actor, actor\_info, address, category, customer, film, film\_actor, film\_category, inventory, and rental.



## THIS IS REVIEW!

We covered the basics of modifying tables to add or delete columns using `ALTER TABLE` in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.

# ALTER TABLE W/ CODE

- We can use SQL code to alter tables by adding and dropping columns

- We specify **ALTER TABLE** tablename, and then we can either **ADD** or **DROP** a column

When we **ADD** a column, we must also specify its data type

- Optionally when we **ADD** a column, we may specify where it should appear, by using **FIRST** or **AFTER**

## MySQL QUERY IN ACTION:

```
SELECT * FROM customer_purchases; -- just to see what's here to start
```

```
ALTER TABLE customer_purchases  
DROP COLUMN customer_id; -- maybe this needs to get dropped for privacy
```

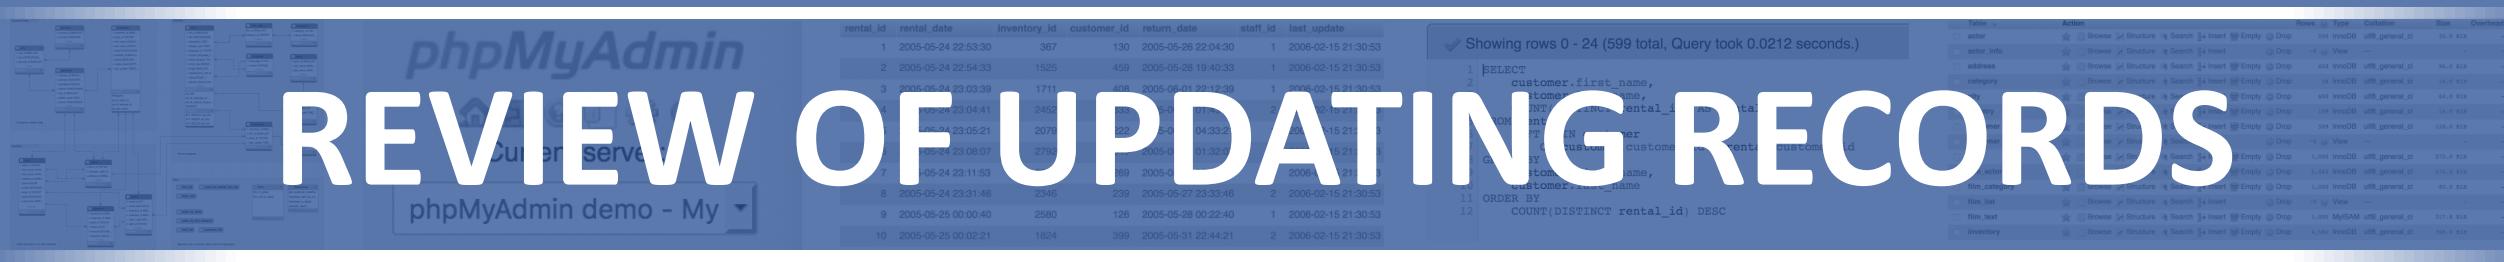
```
ALTER TABLE customer_purchases  
ADD COLUMN purchase_amount DECIMAL(6,2) AFTER customer_purchase_id;  
-- after is optional
```

```
SELECT * FROM customer_purchases; -- what's here at the end
```

## ACTION OUTPUT:

customer_purchase_id	purchase_amount	inventory_id	purchased_at
1	NULL	3	NULL
2	NULL	2	NULL
3	NULL	4	NULL
4	NULL	7	NULL
5	NULL	5	NULL
6	NULL	1	NULL
7	NULL	6	NULL
8	NULL	8	NULL
9	NULL	9	NULL
10	NULL	3	NULL
11	NULL	2	NULL
12	NULL	6	NULL

# REVIEW OF UPDATING RECORDS



# THIS IS REVIEW!

We covered the basics of updating records in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.

# UPDATING RECORDS

- We can update the values of one or more records using an **UPDATE** statement
- Our **UPDATE** statement will always require a **SET** clause, which tells the server what values to set
- In most cases, we will also use a **WHERE** clause (optional) to specify which records to update. If we do not use a **WHERE** clause, the **UPDATE** will process on all records

## MySQL QUERY IN ACTION:

```
UPDATE inventory
```

```
SET number_in_stock = 0 -- these are sold out
```

```
WHERE inventory_id = 9; -- only do this for id #9
```

## (ANOTHER) MySQL QUERY IN ACTION:

```
-- this time we'll do it for two records simultaneously!
```

```
UPDATE inventory
```

```
SET number_in_stock = 0 -- these are sold out
```

```
WHERE inventory_id IN(1,9); -- update 2 records at once
```

# REVIEW OF TABLE RELATIONSHIPS



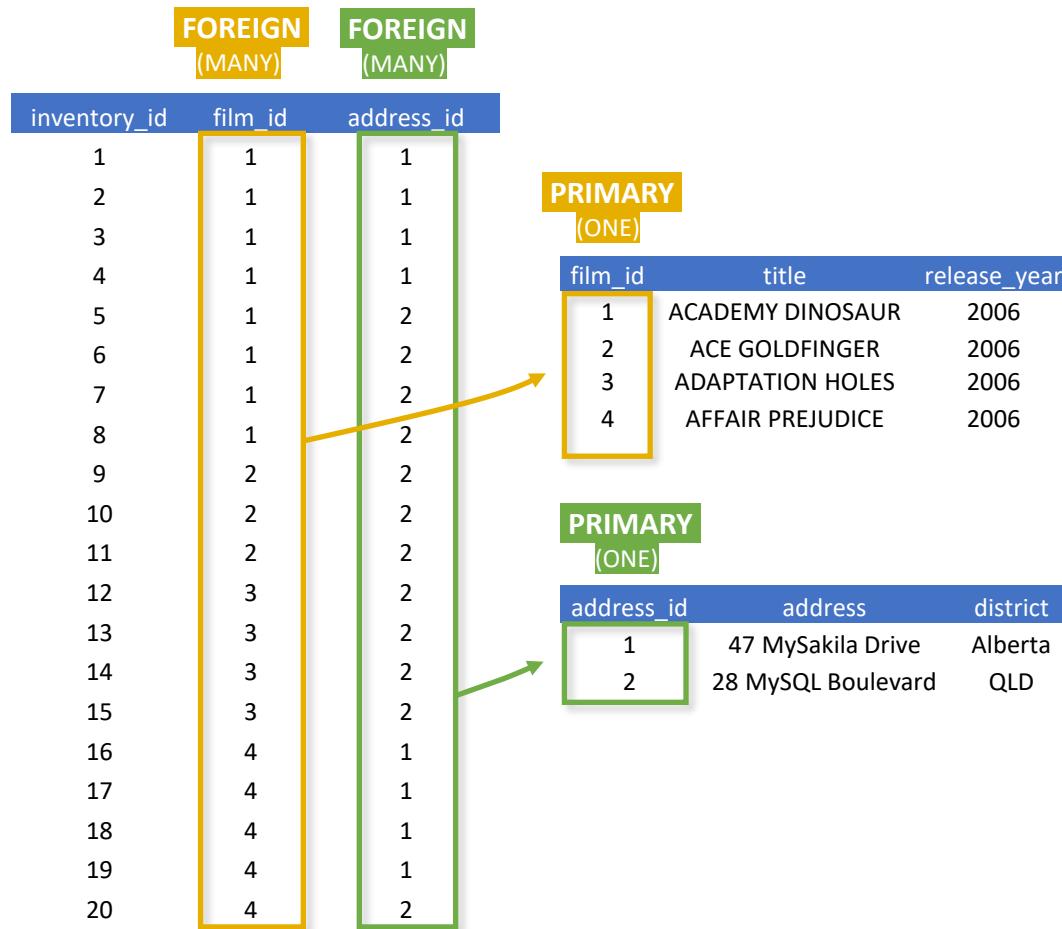
## THIS IS REVIEW!

We covered the basics of table relationships in our MySQL Database Administration for Beginners course. If you already have this concept covered, feel free to skip this lecture.

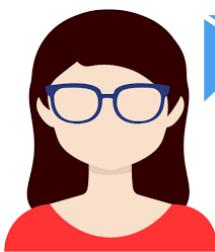
The screenshot shows the phpMyAdmin interface with a blue header. On the left, there's a tree view of database structures. The main area has three tabs: 'Tables' (selected), 'Actions', and 'Rows'. The 'Tables' tab lists tables like actor, address, category, customer, film, film\_actor, film\_text, and inventory. The 'Actions' tab shows various operations like 'Browse', 'Structure', 'Search', 'Insert', 'Drop', and 'View'. The 'Rows' tab shows statistics for each table. In the center, there are two panes: one for a table view showing rental data, and another for a query editor displaying a SELECT statement and its execution results.

# TABLE RELATIONSHIPS & CARDINALITY

Cardinality refers to the **uniqueness of values** in a column (or attribute) of a table and is commonly used to describe how two tables relate (**one-to-one**, **one-to-many**, or **many-to-many**). For now, here are the key points to grasp:



- Primary keys are unique
  - They **cannot** repeat, so there is **only one** instance of each primary key value in a column
- Foreign keys are non-unique
  - They **can** repeat, so there may be **many** instances of each foreign key value in a column
- We can create a **one-to-many** relationship by connecting a **foreign key** in one table to a **primary key** in another



## NEW MESSAGE

January 05, 2013

From: **Sally Bleu (CEO)**

Subject: **New table to track products**

Hey,

Tomorrow we're launching a new product called The Forever Love Bear to complement The Original Mr. Fuzzy.

Could you please create a **products** table in the database? Track when they launched (2012-03-09 at 9am and 2013-01-06 at 1pm, respectively), the product name, and assign an id so we can link to other tables later.

Thanks!

-Sally

Reply

Forward

## Result Preview

-- first, create a products table with columns...

-- product\_id

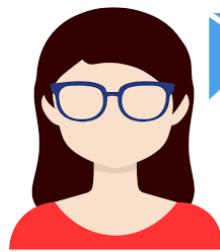
-- created\_at

-- product\_name

-- then, populate it with 2 records like this...

product_id	created_at	product_name
1	2012-03-19 09:00:00	The Original Mr. Fuzzy
2	2013-01-06 13:00:00	The Forever Love Bear

# TEST YOUR SKILLS: ADDING A NEW PRODUCT



## NEW MESSAGE

January 05, 2013

From: **Sally Bleu (CEO)**

Subject: **New table to track products**

Hey,

Tomorrow we're launching a new product called The Forever Love Bear to complement The Original Mr. Fuzzy.

Could you please create a **products** table in the database? Track when they launched (2012-03-09 at 9am and 2013-01-06 at 1pm, respectively), the product name, and assign an id so we can link to other tables later.

Thanks!

-Sally

Reply

Forward

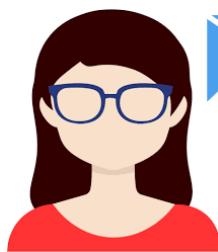
## Solution Queries

```
CREATE TABLE products (
    product_id BIGINT,
    created_at DATETIME,
    product_name VARCHAR(50),
    PRIMARY KEY (product_id)
);
```

```
INSERT INTO products VALUES
(1, '2012-03-19 09:00:00', 'The Original Mr. Fuzzy'),
(2, '2013-01-06 13:00:00', 'The Forever Love Bear');
```

product_id	created_at	product_name
1	2012-03-19 09:00:00	The Original Mr. Fuzzy
2	2013-01-06 13:00:00	The Forever Love Bear

# TEST YOUR SKILLS: ADDING A NEW PRODUCT



## NEW MESSAGE

January 06, 2013

From: **Sally Bleu (CEO)**

Subject: **Add product data to item sales**

Good morning,

Later today, we'll have multiple products selling, I would love to be able to tie our **order\_items** data to the product sold.

Can you please add **product\_id** to the **order\_items** table?

Thanks for the help!

-Sally

Reply

Forward

## Result Preview

-- don't worry about back-populating  
-- for this assignment, just add the column

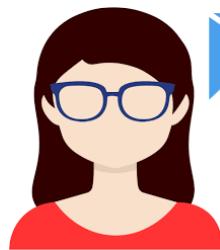
8     -- you should see a **product\_id** column in **order\_items**  
9 •   **SELECT \* FROM order\_items;**

100% 1:4

Result Grid Filter Rows: Search Edit: Export/Import

order_item_id	create_at	order_id	price_usd	cogs_usd	website_session_id	product_id
1	2012-03-19 10:42:46	1	49.99	19.49	20	NULL
2	2012-03-19 19:27:37	2	49.99	19.49	104	NULL
3	2012-03-20 06:44:45	3	49.99	19.49	147	NULL
4	2012-03-20 09:41:45	4	49.99	19.49	160	NULL
5	2012-03-20 11:28:15	5	49.99	19.49	177	NULL
6	2012-03-20 16:12:47	6	49.99	19.49	232	NULL
7	2012-03-20 17:03:41	7	49.99	19.49	241	NULL
8	2012-03-20 23:35:27	8	49.99	19.49	295	NULL
9	2012-03-21 02:35:01	9	49.99	19.49	304	NULL
10	2012-03-21 06:45:58	10	49.99	19.49	317	NULL

# TEST YOUR SKILLS: ADDING PRODUCT TO ORDER\_ITEMS



## NEW MESSAGE

January 06, 2013

From: **Sally Bleu (CEO)**

Subject: **Add product data to item sales**

Good morning,

Later today, we'll have multiple products selling, I would love to be able to tie our **order\_items** data to the product sold.

Can you please add **product\_id** to the **order\_items** table?

Thanks for the help!

-Sally

Reply

Forward

## Solution Query

```
ALTER TABLE order_items  
ADD product_id BIGINT;
```

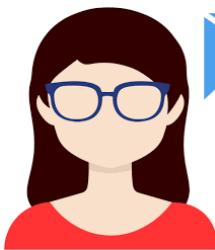
```
8 -- you should see a product_id column in order_items  
9 • SELECT * FROM order_items;
```

100% 1:4

Result Grid Filter Rows: Search Edit: Export/Import

order_item_id	create_at	order_id	price_usd	cogs_usd	website_session_id	product_id
1	2012-03-19 10:42:46	1	49.99	19.49	20	NULL
2	2012-03-19 19:27:37	2	49.99	19.49	104	NULL
3	2012-03-20 06:44:45	3	49.99	19.49	147	NULL
4	2012-03-20 09:41:45	4	49.99	19.49	160	NULL
5	2012-03-20 11:28:15	5	49.99	19.49	177	NULL
6	2012-03-20 16:12:47	6	49.99	19.49	232	NULL
7	2012-03-20 17:03:41	7	49.99	19.49	241	NULL
8	2012-03-20 23:35:27	8	49.99	19.49	295	NULL
9	2012-03-21 02:35:01	9	49.99	19.49	304	NULL
10	2012-03-21 06:45:58	10	49.99	19.49	317	NULL

# TEST YOUR SKILLS: ADDING PRODUCT TO ORDER\_ITEMS



## NEW MESSAGE

January 07, 2013

From: **Sally Bleu (CEO)**

Subject: **Back-populate sales with product\_id**

Hey there,

I noticed that all records in **order\_items** show NULL for **product\_id**.

All of the sales reflected in the database are for product 1, so could you update the records to reflect that? Then we'll have perfect data to use in the future.

Thanks for the help!

-Sally

Reply

Forward

## Result Preview

```
5      -- product_id should be 1 for all records
6 •   SELECT * FROM order_items;
```

100% 1:11

Result Grid



Filter Rows:



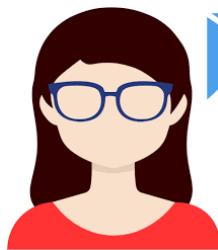
Search



Edit: Export/Import

order_item_id	create_at	order_id	price_usd	cogs_usd	website_session_id	product_id
1	2012-03-19 10:42:46	1	49.99	19.49	20	1
2	2012-03-19 19:27:37	2	49.99	19.49	104	1
3	2012-03-20 06:44:45	3	49.99	19.49	147	1
4	2012-03-20 09:41:45	4	49.99	19.49	160	1
5	2012-03-20 11:28:15	5	49.99	19.49	177	1
6	2012-03-20 16:12:47	6	49.99	19.49	232	1
7	2012-03-20 17:03:41	7	49.99	19.49	241	1
8	2012-03-20 23:35:27	8	49.99	19.49	295	1
9	2012-03-21 02:35:01	9	49.99	19.49	304	1
10	2012-03-21 06:45:58	10	49.99	19.49	317	1

# TEST YOUR SKILLS: UPDATING ORDER\_ITEMS



## NEW MESSAGE

January 07, 2013

From: **Sally Bleu (CEO)**

Subject: **Back-populate sales with product\_id**

Hey there,

I noticed that all records in **order\_items** show NULL for **product\_id**.

All of the sales reflected in the database are for product 1, so could you update the records to reflect that? Then we'll have perfect data to use in the future.

Thanks for the help!

-Sally

Reply

Forward

## Solution Query

```
1 • UPDATE order_items
2   SET product_id = 1
3   WHERE order_item_id > 0;
4
5   -- product_id should be 1 for all records
6 • SELECT * FROM order_items;
```

100% ◂ 1:11

Result Grid



Filter Rows:



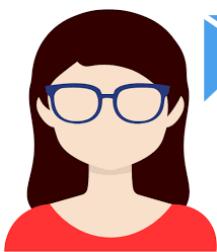
Search



Export/Import

order_item_id	create_at	order_id	price_usd	cogs_usd	website_session_id	product_id
1	2012-03-19 10:42:46	1	49.99	19.49	20	1
2	2012-03-19 19:27:37	2	49.99	19.49	104	1
3	2012-03-20 06:44:45	3	49.99	19.49	147	1
4	2012-03-20 09:41:45	4	49.99	19.49	160	1
5	2012-03-20 11:28:15	5	49.99	19.49	177	1
6	2012-03-20 16:12:47	6	49.99	19.49	232	1
7	2012-03-20 17:03:41	7	49.99	19.49	241	1
8	2012-03-20 23:35:27	8	49.99	19.49	295	1
9	2012-03-21 02:35:01	9	49.99	19.49	304	1
10	2012-03-21 06:45:58	10	49.99	19.49	317	1

# TEST YOUR SKILLS: UPDATING ORDER\_ITEMS



## NEW MESSAGE

January 07, 2013

From: **Sally Bleu (CEO)**

Subject: **Primary and foreign keys?**

Hey there,

I would like to make sure the newly-related **order\_items** and **products** tables have the right relationships specified in the database.

Can you set up the proper primary and foreign key relationships between those two tables?

Thank you!

-Sally

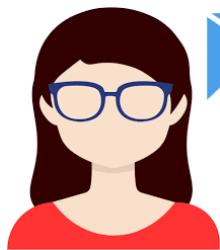
Reply

Forward

## Helpful Hint

-- TIP: you can do this with code or with the UI

# TEST YOUR SKILLS: PRIMARY AND FOREIGN KEYS



## NEW MESSAGE

January 07, 2013

From: **Sally Bleu (CEO)**

Subject: Primary and foreign keys?

Hey there,

I would like to make sure the newly-related **order\_items** and **products** tables have the right relationships specified in the database.

Can you set up the proper primary and foreign key relationships between those two tables?

Thank you!

-Sally

Reply

Forward

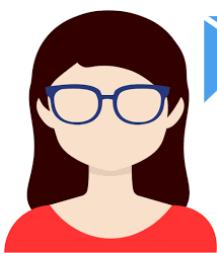
## Solution

The screenshot shows a database management interface for the 'mavenbearbuilders' schema. A foreign key named 'order\_items\_product\_id' is being configured for the 'order\_items' table, which references the 'products' table. The 'Foreign Keys' tab is selected. The configuration details are as follows:

- Foreign Key:** order\_items\_product\_id
- Referenced Table:** 'mavenbearbuilders'.products
- Column:** product\_id
- Referenced Column:** product\_id
- On Update:** NO ACTION
- On Delete:** NO ACTION

The 'product\_id' column is highlighted with a blue border. The interface also includes tabs for Columns, Indexes, Triggers, Partitioning, Options, and buttons for Apply and Revert.

# TEST YOUR SKILLS: PRIMARY AND FOREIGN KEYS



## NEW MESSAGE

April 1, 2013

From: **Sally Bleu (CEO)**

Subject: **Help get order data into the database**

Hey there,

Now that you've done all the work to get our products and `order_items` tables synced up, let's import the attached Q1 data into `order_items` and `order_item_refunds`.

Really curious to start digging into sales trends!

Thanks!

-Sally



[06.order\\_items\\_2013\\_Jan-Mar](#)  
[07.order\\_item\\_refunds\\_2013\\_Jan-Mar](#)

Reply

Forward

## Helpful Hint

-- TIP: remember order matters because of your keys

-- if you've done this correctly, there should be 3,859 records  
`SELECT COUNT(*) AS total_records FROM order_items ;`

-- if you've done this correctly, there should be 234 records  
`SELECT COUNT(*) AS total_records FROM order_item_refunds ;`

# TEST YOUR SKILLS: IMPORT 2013 Q1 DATA



## NEW MESSAGE

April 1, 2013

From: **Sally Bleu (CEO)**

Subject: **Help get order data into the database**

Hey there,

Now that you've done all the work to get our products and `order_items` tables synced up, let's import the attached Q1 data into `order_items` and `order_item_refunds`.

Really curious to start digging into sales trends!

Thanks!

-Sally



[06.order\\_items\\_2013\\_Jan-Mar](#)  
[07.order\\_item\\_refunds\\_2013\\_Jan-Mar](#)

[Reply](#)

[Forward](#)

## Solution

-- See solution video for more detail

-- if you've done this correctly, there should be 3,859 records  
`SELECT COUNT(*) AS total_records FROM order_items ;`

-- if you've done this correctly, there should be 234 records  
`SELECT COUNT(*) AS total_records FROM order_item_refunds ;`

# TEST YOUR SKILLS: IMPORT 2013 Q1 DATA

# DATABASE REPLICATION

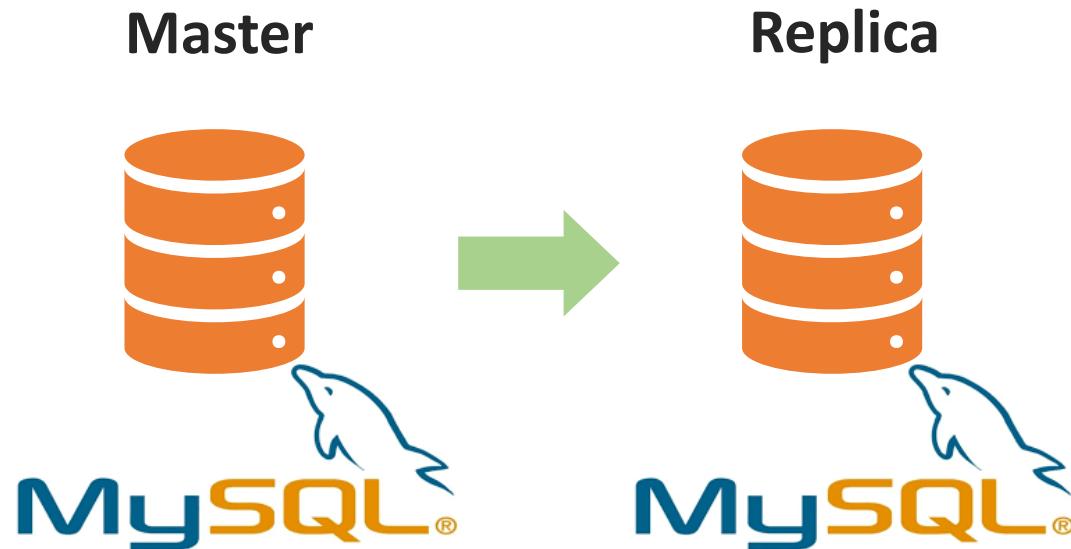
The screenshot displays the phpMyAdmin interface with three main panels:

- Left Panel:** Shows the database schema with tables like actor, address, category, customer, film, film\_actor, film\_category, inventory, and rental.
- Middle Panel:** A table titled "rental" showing data from rows 1 to 10. The columns are rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update.
- Right Panel:** A query results page showing rows 0 to 24 of a query. The query is:

```
SELECT customer.first_name, customer.last_name, COUNT(DISTINCT rental_id) AS rental_count
FROM customer
JOIN rental ON customer.customer_id = rental.customer_id
GROUP BY customer.first_name, customer.last_name
ORDER BY COUNT(DISTINCT rental_id) DESC
```

# INTRODUCTION TO REPLICATION

**Replication** enables us to store the same data on two or more servers by creating copies, known as replicas

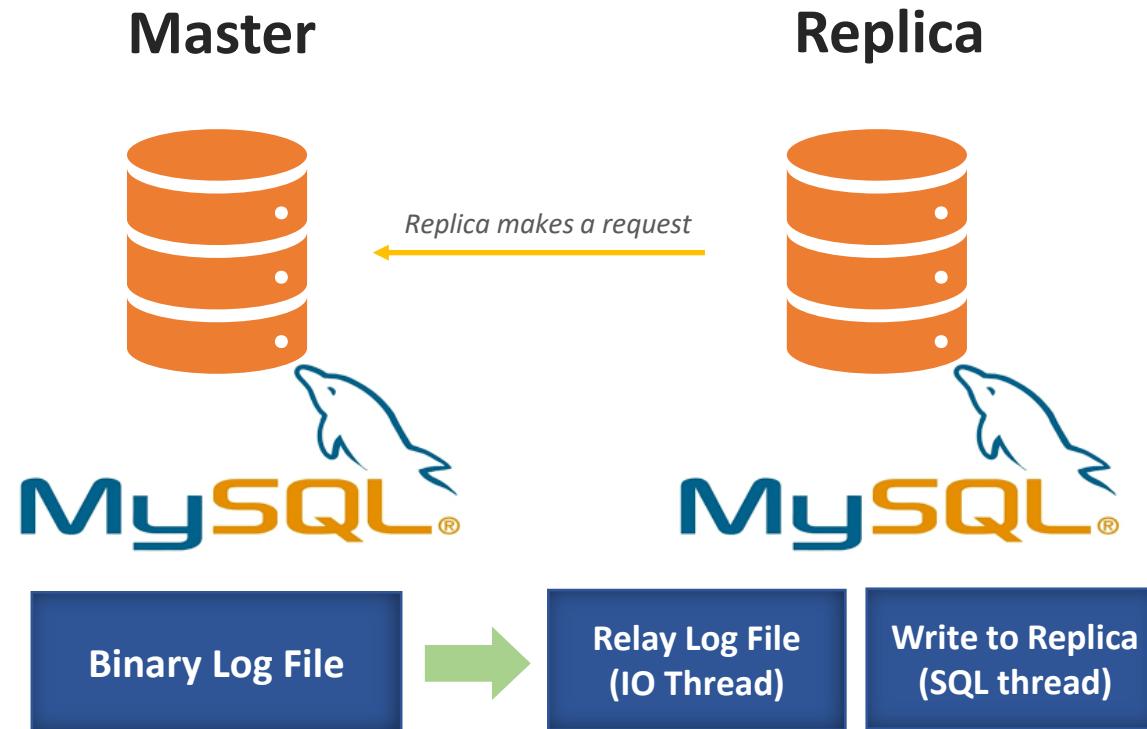


## BENEFITS & USES:

- Replication can serve as a form of backup
- Creating a read only replica for reporting purposes keeps analysis out of production
- In certain cases, splitting the load between two servers can improve application performance

# ROW-BASED BINARY LOG FILE REPLICATION

The first method of Replication we are going to focus on is **Row-based Replication using Binary Log files**

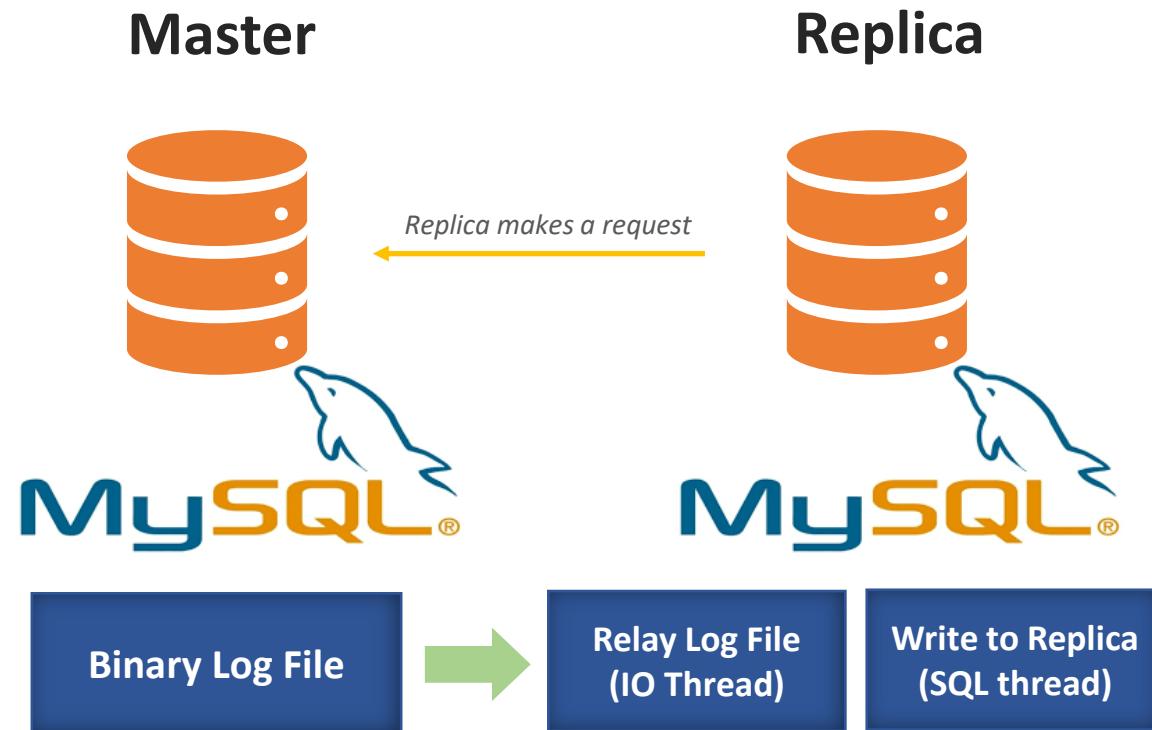


## IMPORTANT POINTS

- Binary Log files record changes made to records of data (think “before and after”)
- Replication typically (not always) happens asynchronously, so the Replica can lag behind
- Row-based replication is typically faster than statement-based replication, but can be harder to audit because it does not contain SQL statements

# STATEMENT-BASED BINARY LOG FILE REPLICATION

The next method of Replication we are going to discuss is **Statement-based Replication using Binary Log files**. Much of this will feel similar to row-based replication.



## IMPORTANT POINTS

- Binary Log files record the actual SQL statements that make changes to the database
- Statement-based replication is typically easier to read and audit compared to row-based
- Non-deterministic queries can be big trouble

# ROW-BASED vs STATEMENT-BASED REPLICATION

	ROW-BASED	STATEMENT-BASED
Changing Few Rows	✓ <b>PRO:</b> statements which update very few rows will execute very quickly	✗ <b>CON:</b> statements which update very few rows will usually perform more slowly than Row-Based
Changing Many Rows	✗ <b>CON:</b> statements which update a large number of rows will perform very slowly	✓ <b>PRO:</b> statements updating a large number of rows can still execute quickly
Consistent Data	✓ <b>PRO:</b> does not encounter issues with non-deterministic queries like statement-based	✗ <b>CON:</b> non-deterministic queries can spell trouble (example, INSERT w/ auto-incrementing PK)
Auditing Changes	✗ <b>CON:</b> harder to audit, because you only see changes, not the statements themselves	✓ <b>PRO:</b> easier to audit, because you see the statements themselves
Handling Triggers	✓ <b>PRO:</b> no problem handling stored routines and triggers	✗ <b>CON:</b> can create problems with stored routines and triggers.

# BACKUP & RECOVERY

The image shows a screenshot of the phpMyAdmin web application. On the left, there is a database schema diagram with various tables like actor, address, category, city, country, customer, customer\_list, film, film\_actor, film\_category, film\_text, and inventory. In the center, a query results page displays a table with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. The results show 24 rows of rental data. Above the results, it says "Showing rows 0 - 24 (599 total, Query took 0.0212 seconds.)". Below the results, the SQL query is shown:

```
1 SELECT
2   COUNT(DISTINCT rental_id) AS rental_count
3   FROM rental
4   JOIN customer
5   ON rental.customer_id = customer.customer_id
6   JOIN (
7     SELECT customer_id, COUNT(DISTINCT first_name) AS count_name
8     FROM customer
9     GROUP BY customer_id
10    ) AS customer_count_name
11   ON customer.customer_id = customer_count_name.customer_id
12   ORDER BY
13     COUNT(DISTINCT rental_id) DESC
```

On the right, there is a table structure page for the "Inventory" table. It lists columns: id, name, and last\_update. The table has 599 rows and a size of 317.4 KB.

# INTRODUCTION TO BACKUP & RECOVERY

---

We use Backups to **make sure our data is protected and recoverable** in the event of loss

## Physical vs Logical Backups

- **Physical** backups store the raw data in a file, where **Logical** backups store the SQL statements needed to recreate the database and populate it
  - With Logical backups, (ex: MySQL Dump), you'll be storing CREATE and INSERT statements
- 

## Online vs Offline Backups

- **Online** backups occur while the server is running. The advantage you don't have to take the server down, so it won't interfere with other clients using the server.
  - **Offline** backups happen when the server is stopped. Simpler, but other clients won't have access either.
- 

## Local vs Remote Backups

- **Local** backups happen on the same host that the MySQL server is running
- **Remote** backups are written somewhere else. This could be another host, a local machine, etc.

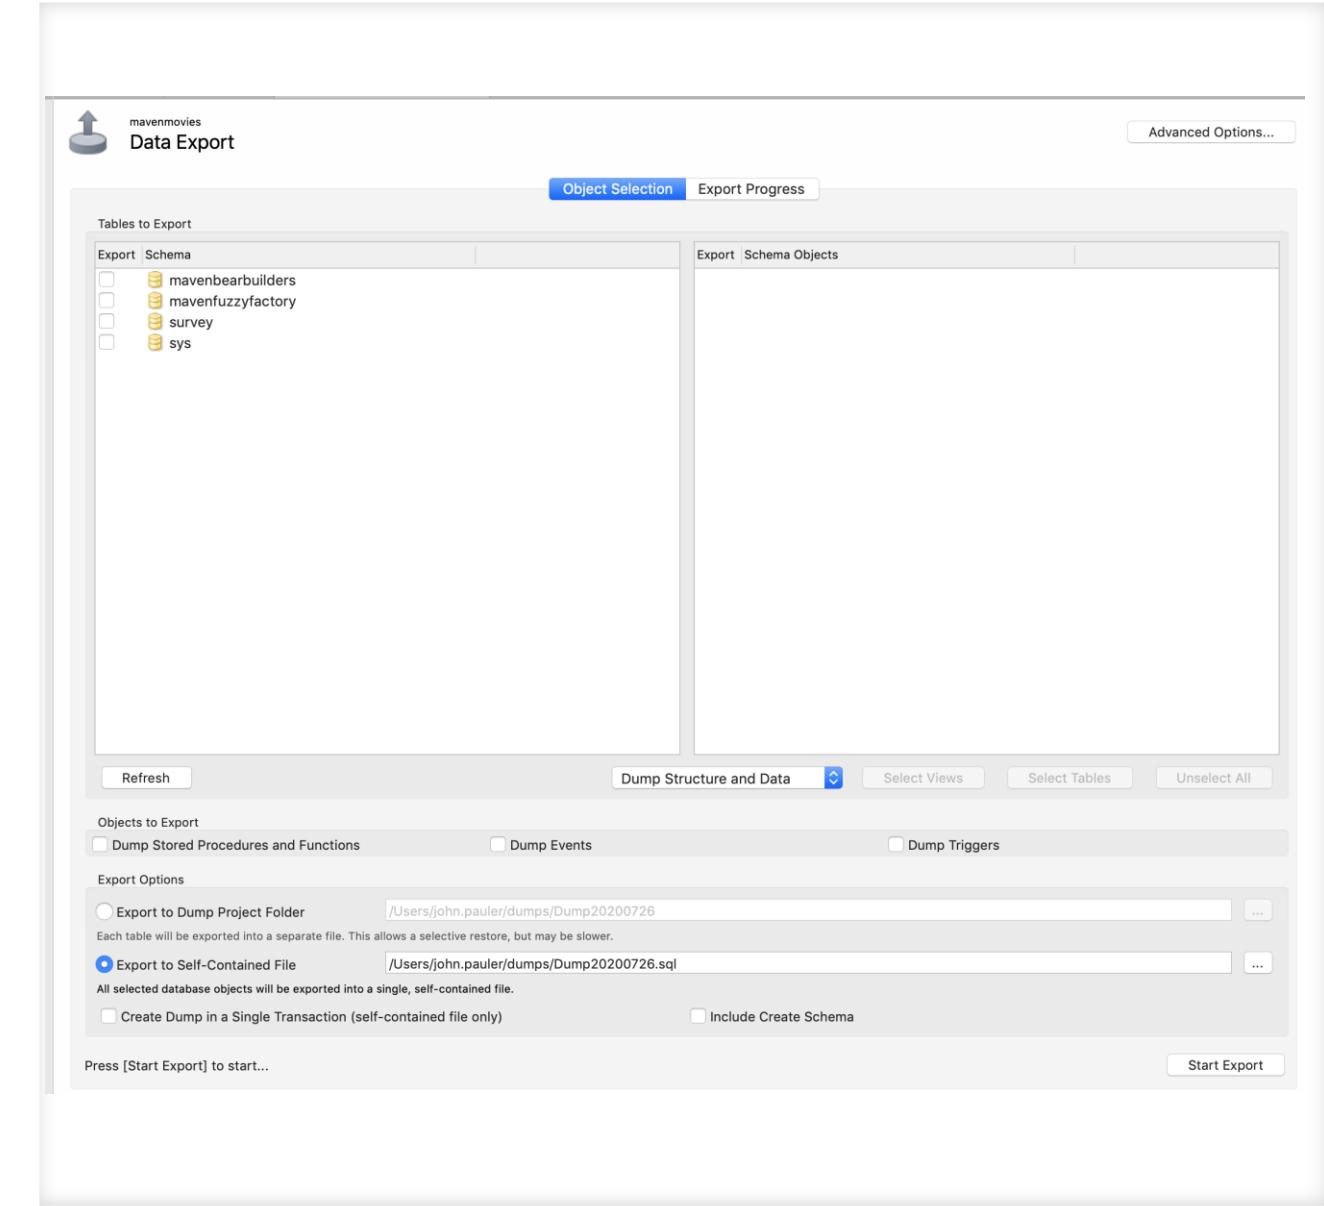
# mysqldump

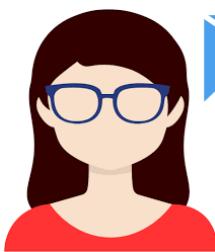
- mysqldump is a database backup method we can use to make sure we can restore data if we lose it

- You can initiate mysqldump from command line, or in Workbench using the wizard

When you use mysqldump, the output will be a .sql file that contains the create statements and insert statements for all of the tables you have dumped

## MySQL WORKBENCH IN ACTION:





## NEW MESSAGE

April 10, 2013

From: **Sally Bleu (CEO)**

Subject: **Need a Backup & Recover Plan**

Good morning,

One of our board members is upset because we haven't taken any steps to backup the data we're aggregating on the MySQL Server.

I would like you to put together a written proposal for everything we should be doing in terms of replication, backup, and recovery. And take a snapshot of the data while you're at it. Thanks!

-Sally

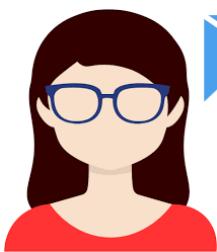
Reply

Forward

## Helpful Hints

- articulate the risks of data loss
- talk about how you will mitigate them
- for the snapshot, your output should be a .sql file
  - this file will recreate your database
  - if you don't know this part, review prior lectures

# TEST YOUR SKILLS: BACKUP & RECOVERY



## NEW MESSAGE

April 10, 2013

From: **Sally Bleu (CEO)**

Subject: **Need a Backup & Recover Plan**

Good morning,

One of our board members is upset because we haven't taken any steps to backup the data we're aggregating on the MySQL Server.

I would like you to put together a written proposal for everything we should be doing in terms of replication, backup, and recovery. And take a snapshot of the data while you're at it. Thanks!

-Sally

Reply

Forward

## *Solution Query*

-- Ideally, you would have addressed each of these...

-- data loss risks

-- statement-based vs row-based replication

-- creating a backup snapshot (`mysqldump`)

# TEST YOUR SKILLS: BACKUP & RECOVERY

# INTRODUCING THE MID COURSE PROJECT

## THE SITUATION

**Maven Bear Builders** has been up and running for a little over a year. You and your CEO have made some improvements to the database, but as the business continues to change, she needs more help tweaking the structure and importing additional data sets.

## THE OBJECTIVE

**Use SQL to:**

- Import additional data into the bearbuilders database
- Enhance the data structure to accommodate new tracking needs for the business

As a Database Administrator, part of your job is executing on specific tasks like altering tables. Another major focus area is staying on top of things like backup, recovery, and database security. Use any opportunities you see as chance to flex your muscle as a thought leader in these areas!

# MID COURSE PROJECT QUESTIONS

1

Import **Q2 orders and refunds** into the database using the files below:



08.order\_items\_2013\_Apr-June  
09.order\_item\_refunds\_2013\_Apr-Jun

~ 0:38

2

Next, help update the structure of the **order\_items** table:

~ 4:20

- The company is going to start cross-selling products and will want to track whether each item sold is the **primary** item (the first one put into the user's shopping cart) or a **cross-sold** item
- Add a binary column to the **order\_items** table called **is\_primary\_item**

3

Update all *previous* records in the **order\_items** table, setting **is\_primary\_item = 1** for all records

~ 6:15

- Up until now, all items sold were the primary item (since cross-selling is new)
- **Confirm this change has executed successfully**

# MID COURSE PROJECT QUESTIONS

4

**Add two new products to the `products` table, then import the remainder of 2013 orders and refunds, using the product details and files shown below:**

~ 9:00



`10.order_items_2013_Jul-Dec`  
`11.order_item_refunds_2013_Jul-Dec`

product_id	created_at	product_name
3	2013-12-12 09:00:00	The Birthday Sugar Panda
4	2014-02-05 10:00:00	The Hudson River Mini bear

5

Your CEO would like to make sure the database has a high degree of data integrity and avoid potential issues as more people start using the database. If you see any opportunities to **ensure data integrity by using constraints like `NON-NULL`**, add them to the relevant columns in the tables you have created.

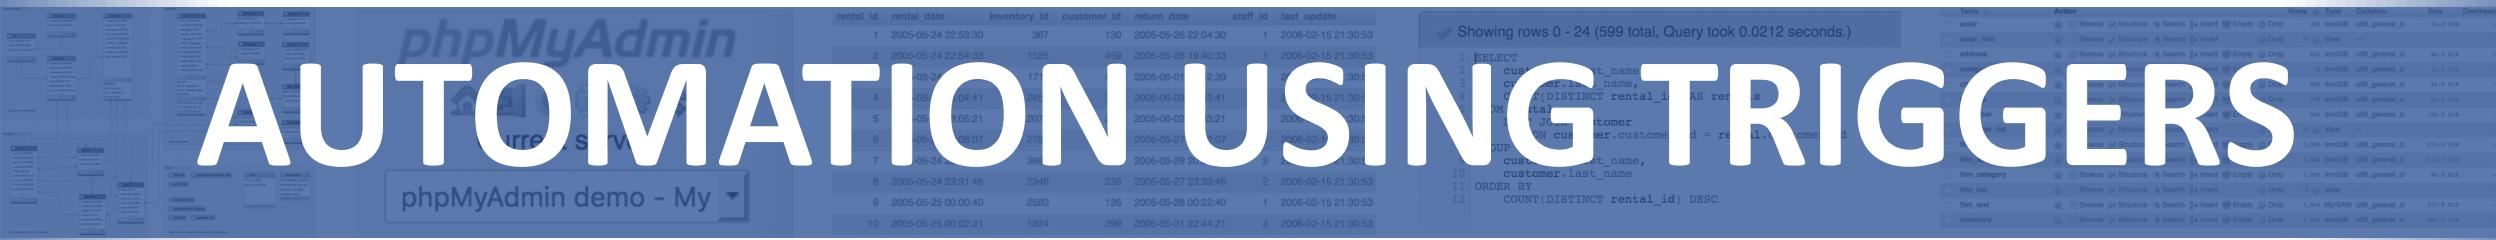
~ 13:38

6

One of the company's board advisors is pressuring your CEO on data risks and making sure she has a great backup and recovery plan. **Prepare a report on possible risks for data loss and steps the company can take to mitigate these concerns.**

~ 17:30

# AUTOMATION USING TRIGGERS



## THIS FIRST TRIGGERS VIDEO IS REVIEW!



We covered the basics of TRIGGERS in our MySQL Database Administration for Beginners course. If you already have these concepts covered, feel free to skip ahead. If you want review of this important concept, check it out!

# TRIGGERS REVIEW

- MySQL allows us to create Triggers, where we can prescribe certain actions on a table to trigger one or more other actions to occur

- We may prescribe that our triggered action occurs either BEFORE or AFTER an INSERT, UPDATE, or a DELETE

- Triggers are a very common way to make sure related tables remain in sync as they are updated over time

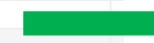
## MySQL Workbench in Action

```
CREATE TRIGGER purchaseUpdateInventory
AFTER INSERT ON customer_purchases
FOR EACH ROW
    UPDATE inventory
        -- subtracting an item for each purchase
        SET number_in_stock = number_in_stock - 1
        WHERE inventory_id = NEW.inventory_id;

INSERT INTO customer_purchases VALUES
(13,NULL,3,NULL), -- inventory_id = 3, velour jumpsuit
(14,NULL,4,NULL) -- inventory_id = 4, house slippers
;
```

```
SELECT * FROM inventory;
```

## Result Preview (SELECT \* FROM inventory)



Result Grid		
inventory_id	item_name	number_in_stock
1	fur coat	0
2	moccassins	4
3	velour jumpsuit	12
4	house slippers	6
5	brown leather jacket	3
6	broken keyboard	6
8	kneeboard	2
9	pro wings sneakers	0
10	wolf skin hat	1
11	fur fox skin	1
12	plaid button up shirt	8
13	flannel zebra jammies	6

Result Grid		
inventory_id	item_name	number_in_stock
1	fur coat	0
2	moccassins	4
3	velour jumpsuit	11
4	house slippers	5
5	brown leather jacket	3
6	broken keyboard	6
8	kneeboard	2
9	pro wings sneakers	0
10	wolf skin hat	1
11	fur fox skin	1
12	plaid button up shirt	8
13	flannel zebra jammies	6

# TRIGGERS DEEP DIVE

- MySQL supports only row-level Triggers, however in other syntax languages, statement-level Triggers are supported

- MySQL does not support ALTER TRIGGER, so to modify an existing Trigger, you'll use DROP TRIGGER and then CREATE TRIGGER

- In order to CREATE, DROP or view the Triggers on a table, you will need to have the Triggers privilege enabled for that table

## MySQL Workbench in Action

- 1 -- this removes the Trigger
- 2 -- you might do this to stop using the Trigger
- 3 -- or you could DROP it to recreate a modified Trigger
- 4 • **DROP TRIGGER** purchaseUpdateInventory;

## Trigger Privileges

The screenshot shows the MySQL Workbench interface with two main windows. The top window is titled 'Trigger Privileges' and shows the 'Administrative Roles' tab selected. It lists various global privileges such as CREATE ROUTINE, CREATE TABLESPACE, and TRIGGER. The bottom window shows the 'Schema Privileges' tab for the 'thrifishop' schema, where the 'TRIGGER' privilege is explicitly listed under the 'Privileges' column. Both windows have green circles highlighting the 'Administrative Roles' and 'Schema Privileges' tabs.

**Global**

**Schema-Specific**

# TRIGGERS DEEP DIVE

- In some cases, using Triggers with BEFORE and AFTER will produce different results
- Multiple Triggers can exist on the same table. By default, they will fire in creation order. This can be modified using **FOLLOWERS** or **PROCEDES**
- To see Triggers in a schema, use **SHOW TRIGGERS**. Or you can query the **INFORMATION\_SCHEMA** to see all triggers in your instance

## MySQL Workbench in Action

```
CREATE TRIGGER purchaseUpdatePurchaseSummary before
BEFORE INSERT ON customer_purchases
FOR EACH ROW
UPDATE purchase_summary
SET purchases_excluding_last = (
    SELECT COUNT(customer_purchase_id)
    FROM customer_purchases
    WHERE customer_purchases.customer_id = purchase_summary.customer_id
)
WHERE customer_id = NEW.customer_id
AND purchase_summary_id > 0; -- this is just to handle safe update mode;
```

```
CREATE TRIGGER purchaseUpdatePurchaseSummary after
AFTER INSERT ON customer_purchases
FOR EACH ROW
UPDATE purchase_summary
SET total_purchases = (
    SELECT COUNT(customer_purchase_id)
    FROM customer_purchases
    WHERE customer_purchases.customer_id = purchase_summary.customer_id
)
WHERE customer_id = NEW.customer_id
AND purchase_summary_id > 0; -- this is just to handle safe update mode;
```

## Check Out Active Triggers

12 • SHOW TRIGGERS;							
13							
Trigger	Event	Table	Statement	Timing	Created	sql_mode	
purchaseUpdateInventory	INSERT	customer_purchases	UPDATE inventory -- subtracting an item for...	AFTER	2020-08-15 12:02:13.19	ONLY_FULL_GROUP_BY	
14 • SELECT * FROM INFORMATION_SCHEMA.TRIGGERS;							
Trigger_Catalog	Trigger_Schema	Trigger_Name	Event_Manipulation	Event_Object_Catalog	Event_Object_Schema	Event	
def	sys	sys_config_insert_set_user	INSERT	def	sys	sys	
def	sys	sys_config_update_set_user	UPDATE	def	sys	sys	
def	thriftshop	purchaseUpdateInventory	INSERT	def	thriftshop	custo	

# TRIGGERS DEEP DIVE

- When creating a Trigger, BEFORE or AFTER is required, as is INSERT, UPDATE, or DELETE, naming the table, and including a valid SQL statement
- FOLLOWES or PRECEDES are optional
- Specifying OLD or NEW when referencing columns in your statement is an optional modifier
- Note that INSERT Triggers do not work with OLD (no OLD value exists) and DELETE Triggers won't work with NEW

## MySQL Workbench in Action

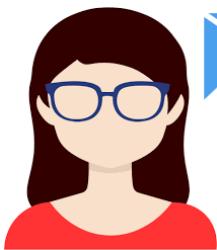
```
CREATE TRIGGER purchaseUpdatePurchaseSummary_before
BEFORE INSERT ON customer_purchases
FOR EACH ROW
UPDATE purchase_summary
SET purchases_excluding_last = (
    SELECT COUNT(customer_purchase_id)
    FROM customer_purchases
    WHERE customer_purchases.customer_id = purchase_summary.customer_id
)
WHERE customer_id = NEW.customer_id
AND purchase_summary_id > 0; -- this is just to handle safe update mode
```

Error Code: 1363.  
There is no OLD row  
in on INSERT trigger

```
CREATE TRIGGER purchaseUpdatePurchaseSummary_before
BEFORE INSERT ON customer_purchases
FOR EACH ROW
UPDATE purchase_summary
SET purchases_excluding_last = (
    SELECT COUNT(customer_purchase_id)
    FROM customer_purchases
    WHERE customer_purchases.customer_id = purchase_summary.customer_id
)
WHERE customer_id = OLD.customer_id
AND purchase_summary_id > 0; -- this is just to handle safe update mode;
```

## RECAP: Triggers Syntax w/ Optional Modifiers

```
CREATE TRIGGER trigger_name
{BEFORE|AFTER}{INSERT|UPDATE|DELETE}
ON table_name FOR EACH ROW
{{FOLLOWES|PRECEDES} existing_trigger_name}
[[some statement we want to execute by trigger]]
{WHERE column_name = {OLD|NEW}.column_name}
```



## NEW MESSAGE

January 03, 2014

From: **Sally Bleu (CEO)**

Subject: **Create an order summary table**

Morning! Now that we're selling multiple products, it would be great to have a table summarizing full orders.

Can you create a table to capture **order\_id**, a **created\_at** timestamp, **website\_session\_id**, primary **product\_id**, # of **items** purchased, **price** and **cogs** in USD?

Could you also back-populate the table using the records from our **order\_items** table?

Thank you!

-Sally

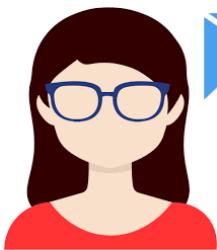
Reply

Forward

## Helpful Hint

```
-- back-populating the orders table
INSERT INTO orders
SELECT
    order_id,
    MIN(created_at) AS created_at,
    MIN(website_session_id) AS website_session_id,
    SUM(CASE
        WHEN is_primary_item = 1 THEN product_id
        ELSE NULL
    END) AS primary_product_id,
    COUNT(order_item_id) AS items_purchased,
    SUM(price_usd) AS price_usd,
    SUM(cogs_usd) AS cogs_usd
FROM order_items
GROUP BY 1
ORDER BY 1;
```

# TEST YOUR SKILLS: ORDER SUMMARY DATA



## NEW MESSAGE

January 03, 2014

From: **Sally Bleu (CEO)**

Subject: **Create an order summary table**

Morning! Now that we're selling multiple products, it would be great to have a table summarizing full orders.

Can you create a table to capture **order\_id**, a **created\_at** timestamp, **website\_session\_id**, primary **product\_id**, # of **items** purchased, **price** and **cogs** in USD?

Could you also back-populate the table using the records from our **order\_items** table?

Thank you!

-Sally

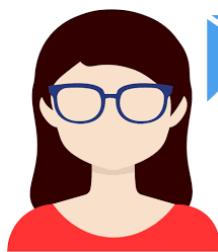
Reply

Forward

## Solution Query

```
-- back-populating the orders table
INSERT INTO orders
SELECT
    order_id,
    MIN(created_at) AS created_at,
    MIN(website_session_id) AS website_session_id,
    SUM(CASE
        WHEN is_primary_item = 1 THEN product_id
        ELSE NULL
    END) AS primary_product_id,
    COUNT(order_item_id) AS items_purchased,
    SUM(price_usd) AS price_usd,
    SUM(cogs_usd) AS cogs_usd
FROM order_items
GROUP BY 1
ORDER BY 1;
```

# TEST YOUR SKILLS: ORDER SUMMARY DATA



## NEW MESSAGE

January 05, 2014

From: **Sally Bleu (CEO)**

Subject: **Automation to update orders table**

Hey there,

The new **orders** table you created is so helpful!

Next, would you be able to set up some automation so that anytime **order\_items** records are inserted into the database, the **orders** table is updated as well?

Thank you!

-Sally

Reply

Forward

## Helpful Hints

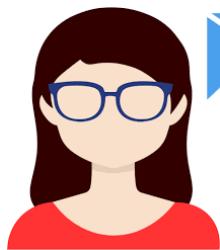
-- You're going to want your trigger to insert a record into orders  
-- You want this to happen whenever new order\_items records are added

**REPLACE INTO** orders

**SELECT**

```
order_id,  
MIN(created_at) AS created_at,  
MIN(website_session_id) AS website_session_id,  
SUM(CASE  
    WHEN is_primary_item = 1 THEN product_id  
    ELSE NULL  
    END) AS primary_product_id,  
COUNT(order_item_id) AS items_purchased,  
SUM(price_usd) AS price_usd,  
SUM(cogs_usd) AS cogs_usd  
FROM order_items  
GROUP BY 1  
ORDER BY 1;
```

# TEST YOUR SKILLS: AUTOMATED ORDERS TRIGGER



## NEW MESSAGE

January 05, 2014

From: **Sally Bleu (CEO)**

Subject: **Automation to update orders table**

Hey there,

The new **orders** table you created is so helpful!

Next, would you be able to set up some automation so that anytime **order\_items** records are inserted into the database, the **orders** table is updated as well?

Thank you!

-Sally

Reply

Forward

## Solution Query

```
CREATE TRIGGER insert_new_orders
AFTER INSERT ON order_items
FOR EACH ROW
REPLACE INTO orders
SELECT
    order_id,
    MIN(created_at) AS created_at,
    MIN(website_session_id) AS website_session_id,
    SUM(CASE
        WHEN is_primary_item = 1 THEN product_id
        ELSE NULL
    END) AS primary_product_id,
    COUNT(order_item_id) AS items_purchased,
    SUM(price_usd) AS price_usd,
    SUM(cogs_usd) AS cogs_usd
FROM order_items
WHERE order_id = new.order_id
GROUP BY 1
ORDER BY 1;
```

# TEST YOUR SKILLS: AUTOMATED ORDERS TRIGGER



## NEW MESSAGE

March 01, 2014

From: **Sally Bleu (CEO)**

Subject: **Putting your automation to the test**

Hey,

It's time to see if your trigger to sync the **order\_items** and **orders** tables is working correctly.

Why don't you go ahead and update the **order\_items** and **order\_item\_refunds** tables with the attached data and we'll see how everything is working.

Thank you!

-Sally



12.order\_items\_2014\_Jan-Feb  
13.order\_item\_refunds\_2014\_Jan-Feb

Reply

Forward

## Result Preview

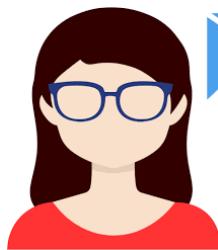
```
-- You should see 10033 records before trigger fires
SELECT COUNT(*) AS total_records FROM orders;

-- Then, test the trigger by inserting into order_items

-- TIP: after insert, find the max order_id
SELECT MAX(order_id) AS max_order_id
FROM order_items;

-- You should see a new row count that matches max_order_id
SELECT COUNT(*) AS total_records FROM orders;
```

# TEST YOUR SKILLS: TESTING YOUR TRIGGER



## NEW MESSAGE

March 01, 2014

From: **Sally Bleu (CEO)**

Subject: **Putting your automation to the test**

Hey,

It's time to see if your trigger to sync the **order\_items** and **orders** tables is working correctly.

Why don't you go ahead and update the **order\_items** and **order\_item\_refunds** tables with the attached data and we'll see how everything is working.

Thank you!

-Sally



12.order\_items\_2014\_Jan-Feb  
13.order\_item\_refunds\_2014\_Jan-Feb

Reply

Forward

## Solution

-- See solution video for more detail

-- You should see 10033 records before trigger fires  
**SELECT COUNT(\*) AS total\_records FROM orders;**

-- Then, test the trigger by inserting into **order\_items**

-- TIP: after insert, find the **max\_order\_id**  
**SELECT MAX(order\_id) AS max\_order\_id**  
**FROM order\_items;**

-- You should see a new row count that matches **max\_order\_id**  
**SELECT COUNT(\*) AS total\_records FROM orders;**

# TEST YOUR SKILLS: TESTING YOUR TRIGGER

# VIEWS

- A view is like a “virtual table”; contents are defined by a query, and data is not physically created

- One advantage of views is the ability to aggregate data from multiple tables into one view.

- Another advantage is security – you can grant someone access to certain views, exclusively

- Finally, altering views will not affect any downstream systems

## MySQL Workbench in Action

-- creating a view

```
CREATE VIEW country_averages AS
SELECT
    country,
    AVG(years_experience) AS yrs_experience,
    AVG(CASE WHEN is_manager = 'Yes' THEN 1 ELSE 0 END) AS pct_mgrs,
    AVG(CASE WHEN education_level = 'Masters' THEN 1 ELSE 0 END) AS pct_masters
FROM salary_survey
GROUP BY 1;
```

-- selecting from the view

```
SELECT * FROM country_averages;
```

## Result Preview

country	yrs_experience	pct_mgrs	pct_masters
United States	23.0379	0.2613	0.1731
United Kingdom	10.1005	0.2487	0.1005
Germany	10.6379	0.3276	0.3448
Netherlands	9.7955	0.2727	0.1364
Canada	11.3556	0.2593	0.0815
Poland	7.0000	0.3333	0.7083
Ukraine	7.1250	0.5000	0.7500
India	6.4795	0.3973	0.3699
Portugal	8.6667	0.3333	0.2222
South Africa	13.0645	0.2581	0.0000
Greece	10.5000	0.6667	0.4167
Israel	9.2727	0.2727	0.1818
Spain	7.4118	0.4118	0.4118
Norway	9.6667	0.0000	0.0000
France	9.7333	0.2000	0.4667
Jersey	11.6000	0.0000	0.2000
Argentina	6.4000	0.2000	0.0000
Belgium	9.7273	0.0909	0.0909
Switzerland	11.0909	0.2727	0.3182
Ireland	11.3478	0.4348	0.2174
Italy	10.8571	0.4286	0.1905



## NEW MESSAGE

March 07, 2014

From: **Brent Cheeseman (Marketing Lead)**

Subject: **Tying website activity to sales**

Hey data rockstar!

Sally tells me the great things you've done with our order data. I have this website session data that I would love to tie into that order data so we can better understand where sales are coming from.

Can you create a **website\_sessions** table and help me import the attached files?

-Brent



14.website\_sessions\_2014\_Jan  
15.website\_sessions\_2014\_Feb

Reply

Forward

## Helpful Hints

-- take a look at the datasets to determine table structure

-- then create the table and import the records

22 -- when you're done, there should be 31,110 records  
23 • `SELECT COUNT(*) AS total_records FROM website_sessions;`

Result Grid		Filter Rows:	Search	Export:
	total_records			
▶	31110			

# TEST YOUR SKILLS: TRACKING WEBSITE SESSIONS



## NEW MESSAGE

March 07, 2014

From: **Brent Cheeseman (Marketing Lead)**

Subject: **Tying website activity to sales**

Hey data rockstar!

Sally tells me the great things you've done with our order data. I have this website session data that I would love to tie into that order data so we can better understand where sales are coming from.

Can you create a **website\_sessions** table and help me import the attached files?

-Brent



14.website\_sessions\_2014\_Jan  
15.website\_sessions\_2014\_Feb

Reply

Forward

## Solution Query

```
CREATE TABLE website_sessions (
    website_session_id BIGINT,
    created_at DATETIME,
    user_id BIGINT,
    is_repeat_session INT,
    utm_source VARCHAR(50),
    utm_campaign VARCHAR(50),
    utm_content VARCHAR(50),
    device_type VARCHAR(50),
    http_referer VARCHAR(100),
    PRIMARY KEY (website_session_id)
);
```

# TEST YOUR SKILLS: TRACKING WEBSITE SESSIONS



## NEW MESSAGE

March 09, 2014

From: **Brent Cheeseman (Marketing Lead)**

Subject: **Reporting Views**

Hey, thanks for getting that **website\_session** data imported!

Next, would you be able to create a view summarizing performance for January and February? I would like to see the **number of sessions** sliced by **year, month, utm\_source, and utm\_campaign** if possible.

Thanks!

-Brent

Reply

Forward

## Helpful Hints

15 -- the view should look like this...

16 • **SELECT \* FROM monthly\_sessions;**

17 100% 18:14

Result Grid



Filter Rows:



Search

year	month	utm_source	utm_campaign	number_of_sessions
2014	1	gsearch	nonbrand	7500
2014	1	NULL	NULL	2724
2014	1	bsearch	nonbrand	1614
2014	1	gsearch	brand	1107
2014	1	bsearch	brand	262
2014	1	socialbook	pilot	1618
2014	2	bsearch	nonbrand	1646
2014	2	NULL	NULL	2819
2014	2	gsearch	nonbrand	8223
2014	2	gsearch	brand	1090
2014	2	socialbook	pilot	2237
2014	2	bsearch	brand	270

# TEST YOUR SKILLS: CREATING VIEWS FOR REPORTING



## NEW MESSAGE

March 09, 2014

From: Brent Cheeseman (*Marketing Lead*)

Subject: Reporting Views

Hey, thanks for getting that **website\_session** data imported!

Next, would you be able to create a view summarizing performance for January and February? I would like to see the **number of sessions** sliced by **month**, **utm\_source**, and **utm\_campaign** if possible.

Thanks!

-Brent

Reply

Forward

## Solution Query

```
6 • CREATE VIEW monthly_sessions AS
7   SELECT
8     YEAR(created_at) AS year,
9     MONTH(created_at) AS month,
10    utm_source,
11    utm_campaign,
12    COUNT(website_session_id) AS number_of_sessions
13   FROM website_sessions
14   GROUP BY 1,2,3,4;
15
16 • SELECT * FROM monthly_sessions;
```

100% ▲ 27:16

Result Grid

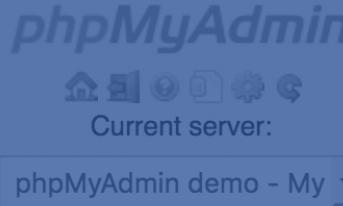


Filter Rows:

Search  Export:

year	month	utm_source	utm_campaign	number_of_sessions
2014	1	gsearch	nonbrand	7500
2014	1	HULL	HULL	2724
2014	1	bsearch	nonbrand	1614
2014	1	gsearch	brand	1107
2014	1	bsearch	brand	262
2014	1	socialbook	pilot	1618
2014	2	bsearch	nonbrand	1646
2014	2	HULL	HULL	2819
2014	2	gsearch	nonbrand	8223
2014	2	gsearch	brand	1090
2014	2	socialbook	pilot	2237
2014	2	bsearch	brand	270

# TEST YOUR SKILLS: CREATING VIEWS FOR REPORTING



# EER DIAGRAMS

Table	Action	Rows	Type	Collation	Size	Overflow
actor	Browse   Structure   Search   Insert   Empty   Drop	400	InnoDB	utf8_general_ci	33,04 Kib	-
actor.info	Browse   Structure   Search   Insert   Empty   Drop	400	InnoDB	utf8_general_ci	44,40 Kib	-
address	Browse   Structure   Search   Insert   Empty   Drop	400	InnoDB	utf8_general_ci	16,40 Kib	-
category	Browse   Structure   Search   Insert   Empty   Drop	400	InnoDB	utf8_general_ci	44,40 Kib	-
city	Browse   Structure   Search   Insert   Empty   Drop	400	InnoDB	utf8_general_ci	16,40 Kib	-
country	Browse   Structure   Search   Insert   Empty   Drop	100	InnoDB	utf8_general_ci	16,40 Kib	-
customer	Browse   Structure   Search   Insert   Empty   Drop	300	InnoDB	utf8_general_ci	128,00 Kib	-
customer.list	Browse   Structure   Search   Insert   Empty   Drop	300	InnoDB	utf8_general_ci	128,00 Kib	-
film	Browse   Structure   Search   Insert   Empty   Drop	1,000	InnoDB	utf8_general_ci	370,40 Kib	-
film.actor	Browse   Structure   Search   Insert   Empty   Drop	1,000	InnoDB	utf8_general_ci	370,40 Kib	-
film.category	Browse   Structure   Search   Insert   Empty   Drop	1,000	InnoDB	utf8_general_ci	400,00 Kib	-
film.list	Browse   Structure   Search   Insert   Empty   Drop	1,000	InnoDB	utf8_general_ci	400,00 Kib	-
film_text	Browse   Structure   Search   Insert   Empty   Drop	1,000	InnoDB	utf8_general_ci	400,00 Kib	-
Inventory	Browse   Structure   Search   Insert   Empty   Drop	4,000	InnoDB	utf8_general_ci	600,00 Kib	-

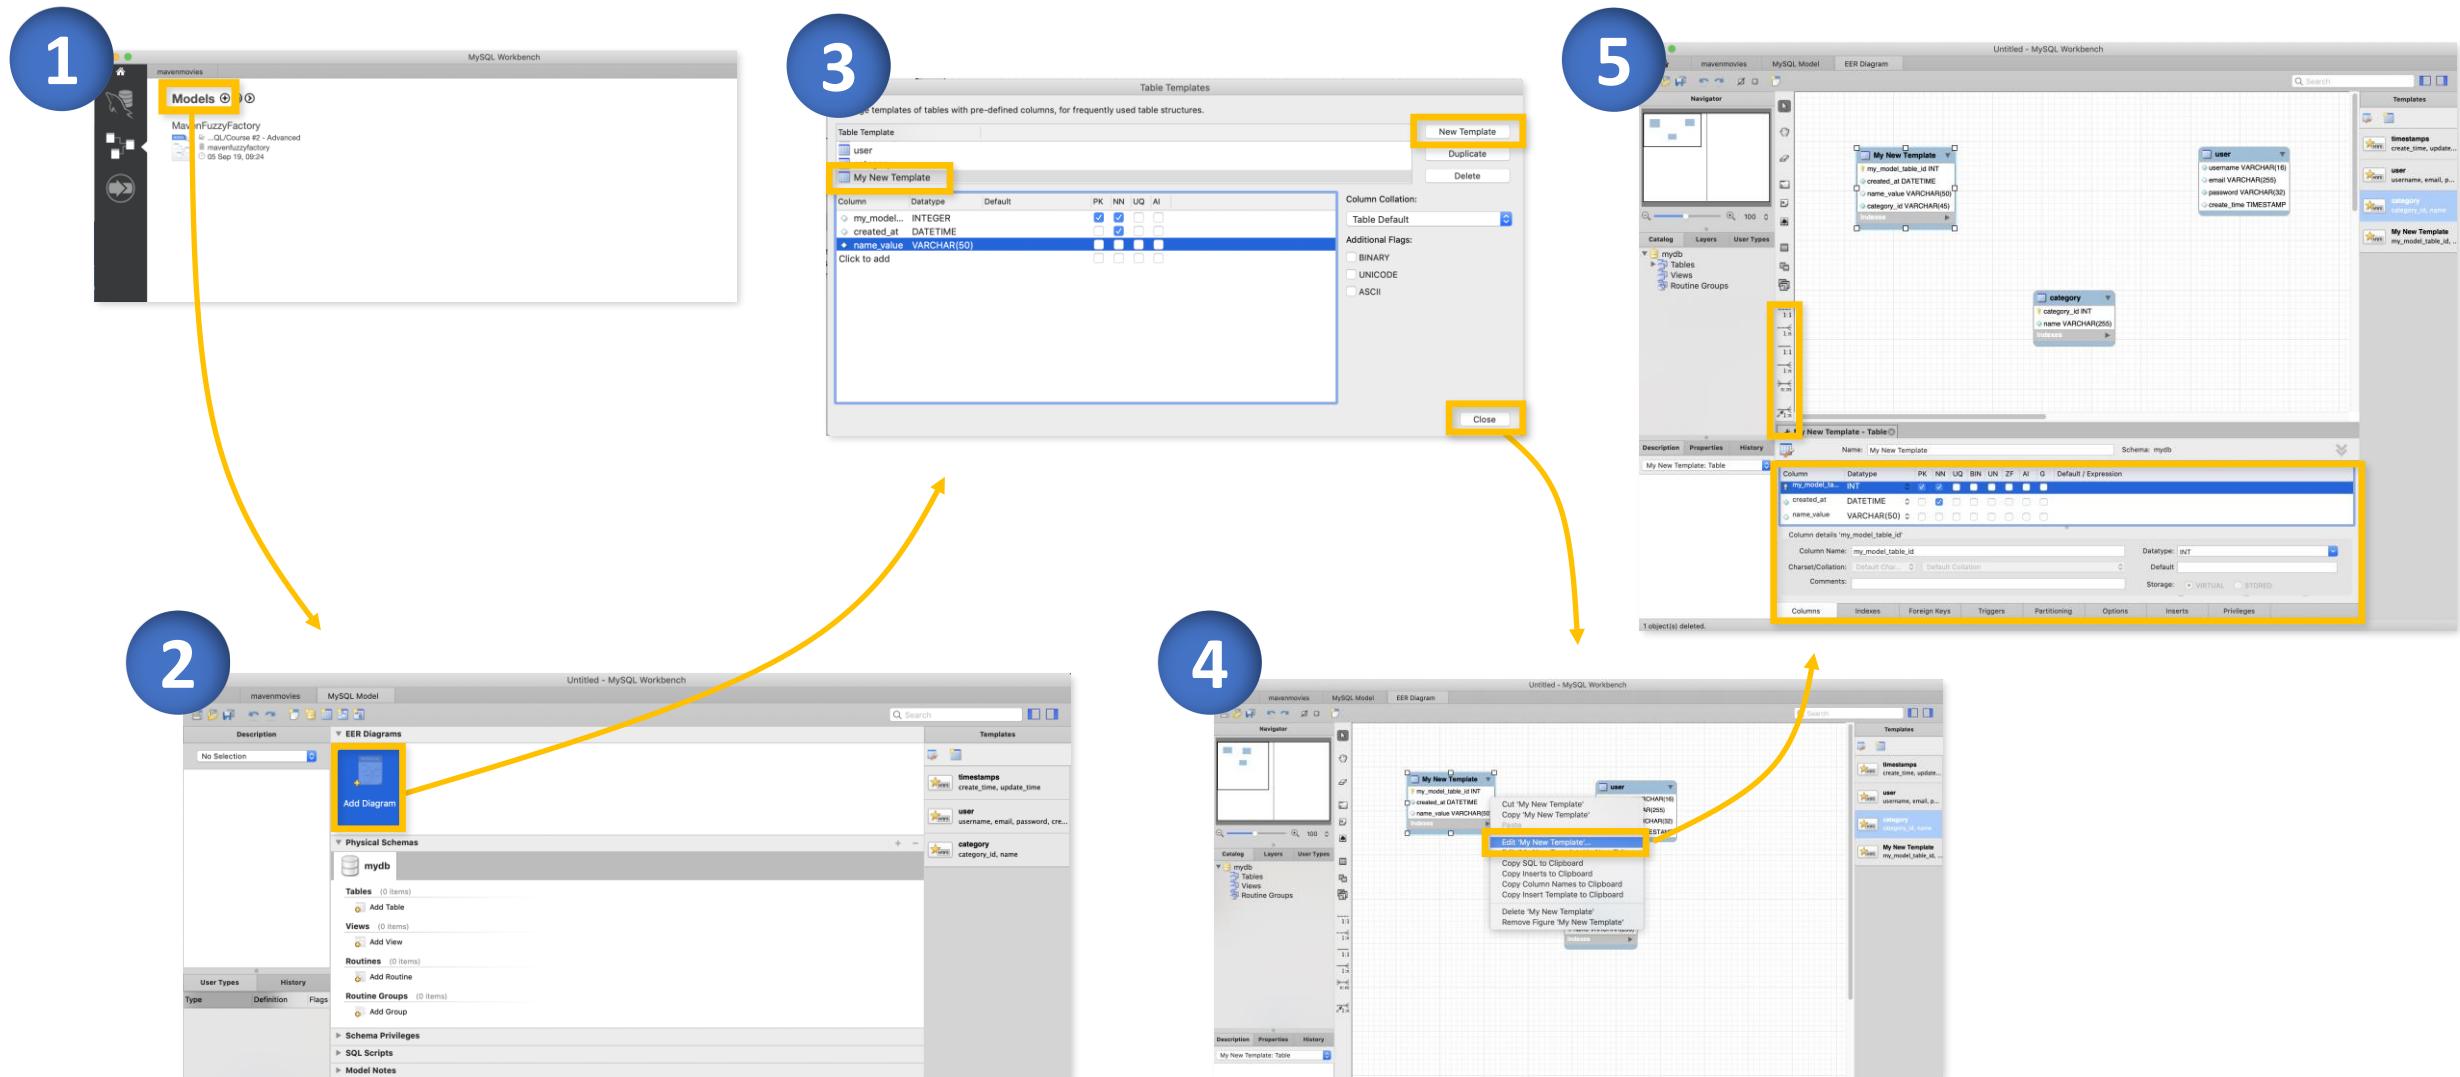
## THIS FIRST PART IS REVIEW!

We covered the basics of EER Diagrams in our MySQL Database Administration for Beginners course. If you already have these concepts covered, feel free to skip ahead. If you want to review this important concept, watch the video.



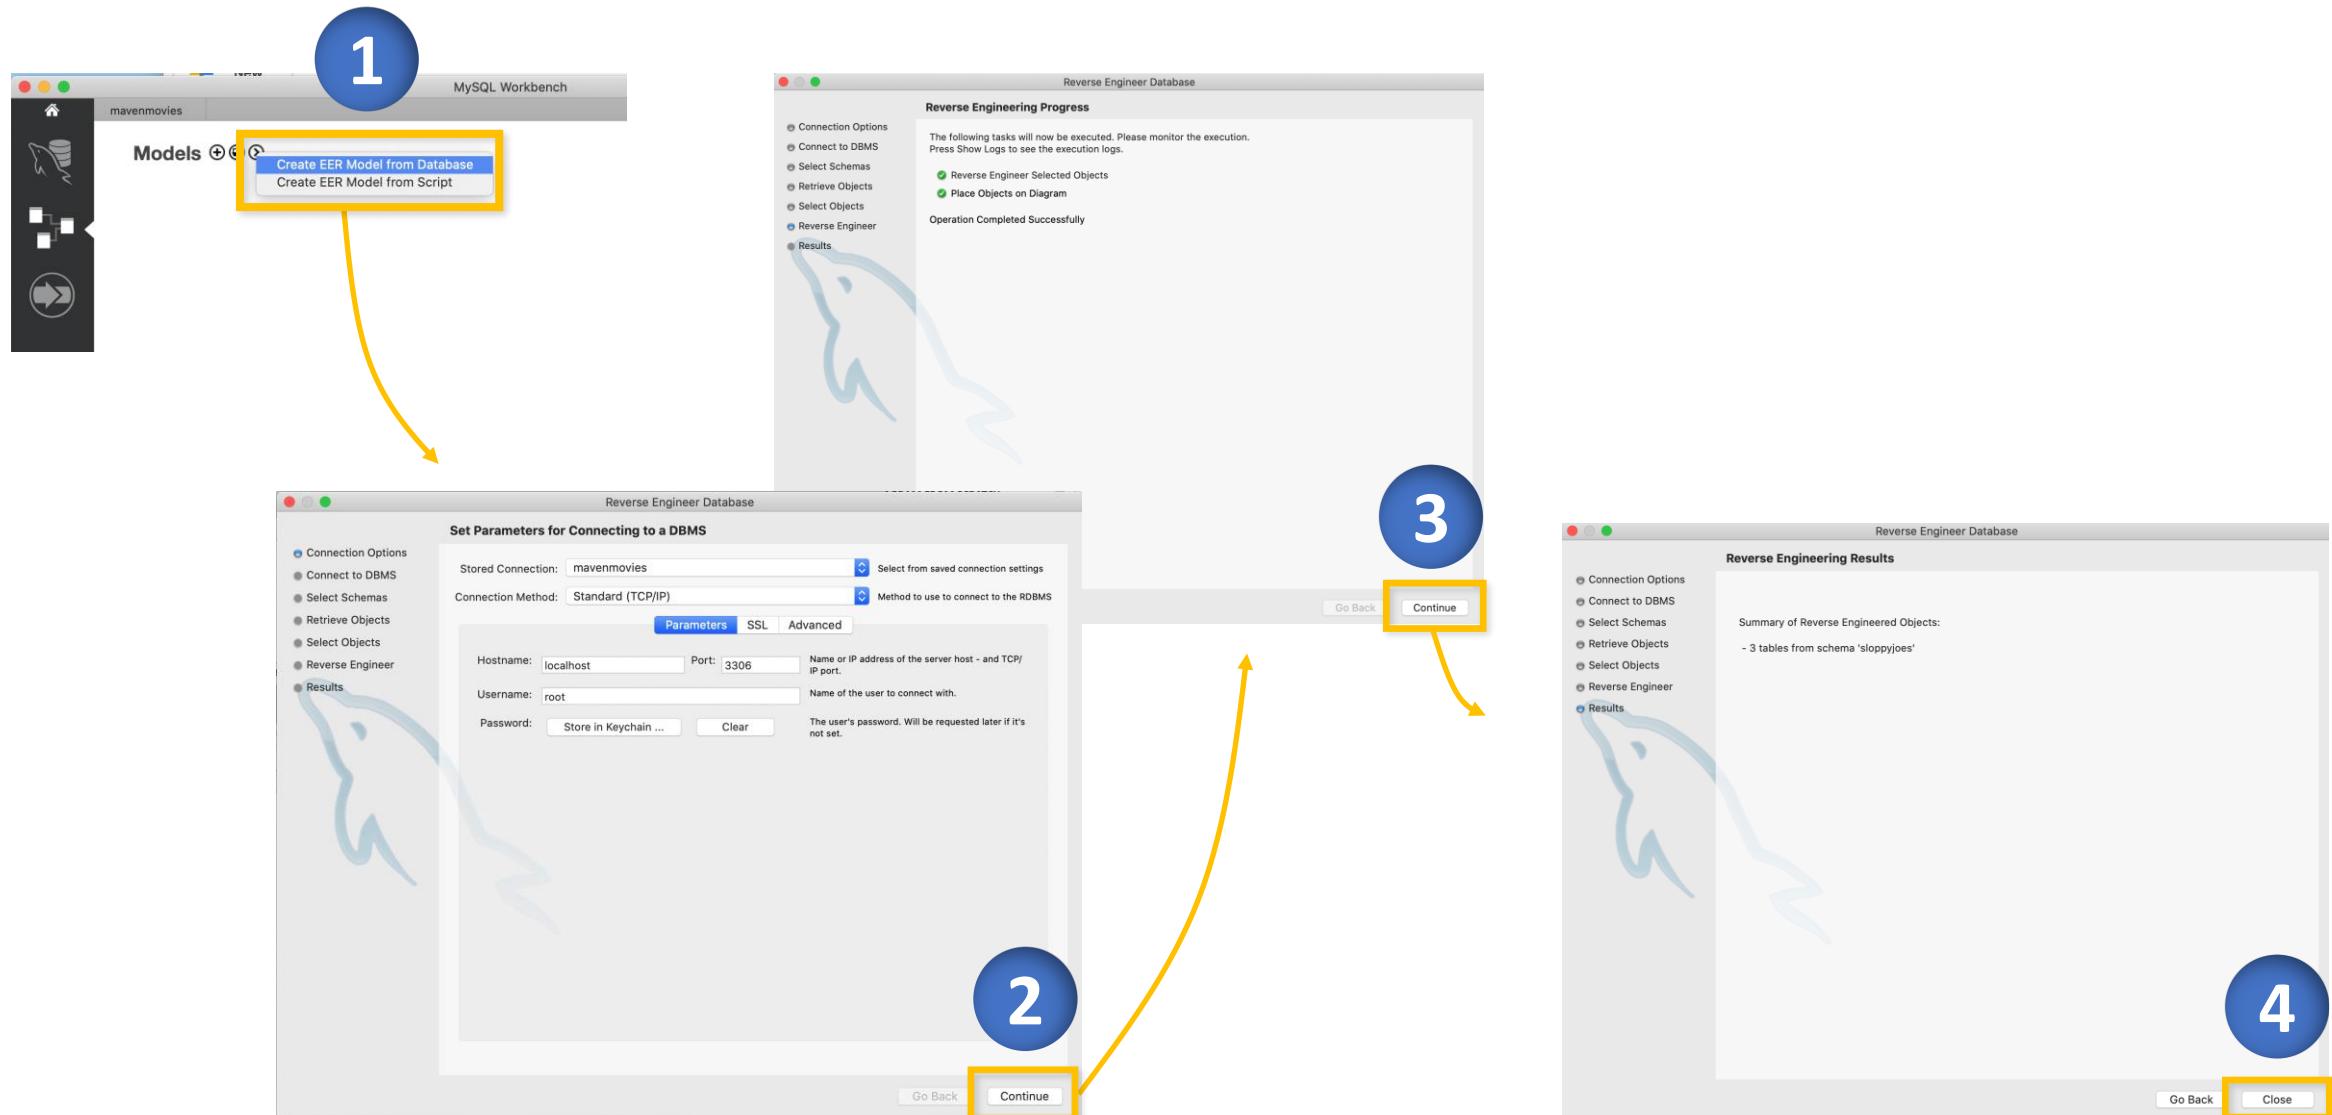


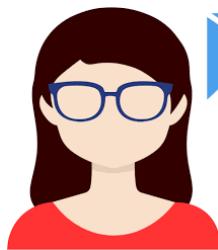
# CREATE AN EER DIAGRAM FROM SCRATCH





# CREATE AN EER DIAGRAM VIA REVERSE ENGINEERING





## NEW MESSAGE

March 15, 2014

From: **Sally Bleu (CEO)**

Subject: **Can we make an EER Diagram?**

Good morning,

Another CEO friend of mine just walked me through their database using an EER diagram. It was so easy to understand how everything connected. We should have one!

Can you put together an EER Diagram for us?

Thanks!

-Sally

Reply

Forward

## Helpful Hints

/\*

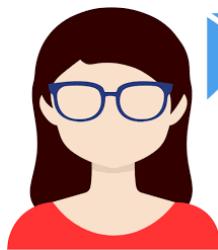
If you've seen a similar EER diagram in another course, don't just copy it!

There are some nuances to this database, so create one from scratch.

No shortcuts!

\*/

# TEST YOUR SKILLS: CREATING EER DIAGRAMS



## NEW MESSAGE

March 15, 2014

From: **Sally Bleu (CEO)**

Subject: **Can we make an EER Diagram?**

Good morning,

Another CEO friend of mine just walked me through their database using an EER diagram. It was so easy to understand how everything connected. We should have one!

Can you put together an EER Diagram for us?

Thanks!

-Sally

Reply

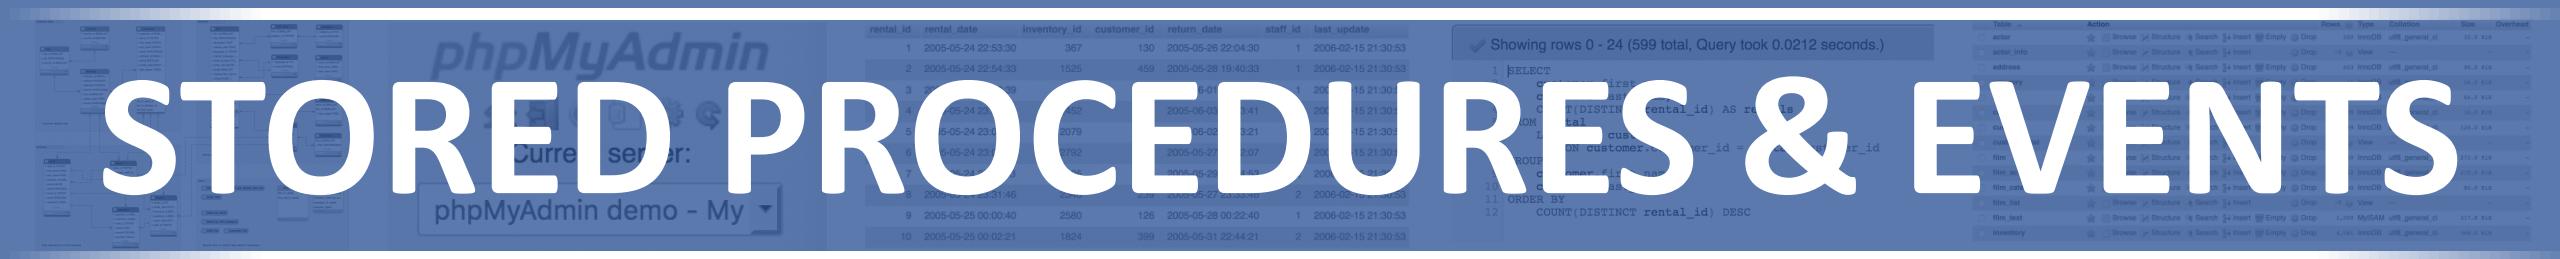
Forward

## *Solution Query*

--- see video for solution

# TEST YOUR SKILLS: CREATING EER DIAGRAMS

# STORED PROCEDURES & EVENTS



# STORED PROCEDURES

- MySQL gives us the ability to store and call frequently used queries on the server. These are then referred to as ‘Procedures’, or commonly ‘Stored Procedures’
- Benefits include more efficient query writing and performance, and the ability to share complex procedures more easily between Analysts and database users
- We invoke the stored procedure using a CALL statement

## MySQL Workbench in Action

```
USE thriftshop;  
  
-- changing the delimiter  
DELIMITER //  
-- creating the procedure  
CREATE PROCEDURE sp_selectAllInventory()  
BEGIN  
    SELECT * FROM inventory;  
END //  
-- changing the delimiter back to the default  
DELIMITER ;  
  
-- calling the procedure that we have created  
CALL sp_selectAllInventory();  
  
-- if we later want to DROP the procedure, we can use the this...  
-- DROP PROCEDURE IF EXISTS sp_selectAllInventory();
```

# STORED PROCEDURES

- Stored Procedures can be used with parameters, which can be IN, OUT, or INOUT
- IN parameters take in a value from the user or application. Values are not returned by the Procedure
- OUT parameters have their values returned by the Procedure
- INOUT is a combination of IN and OUT parameter modes. It can take in a value, which will also output

## MySQL Workbench in Action

```
USE mavenmovies;

-- changing the delimiter
DELIMITER //

-- creating the procedure
CREATE PROCEDURE customer_rentals
(IN custid BIGINT)
BEGIN
    SELECT * FROM rental WHERE customer_id = custid;
END //

-- next, changing the delimiter back
DELIMITER ;

-- calling the procedure, with 135 as the input value
CALL customer_rentals(135)
```

# STORED PROCEDURES

- OUT parameters have their values returned by the Procedure
- If you want, you can specify a variable to store the output returned in an OUT parameter
- INOUT is a combination of IN and OUT parameter modes. It can take in a value, which will also output

## MySQL Workbench in Action

```
-- next, we'll add an OUT...
-- changing the delimiter
DELIMITER //
-- creating the procedure
CREATE PROCEDURE customer_rentals_w_total
(IN custid BIGINT, OUT total_rentals BIGINT) -- this time, we've added an OUT parameter
BEGIN
    SELECT COUNT(rental_id) INTO total_rentals FROM rental WHERE customer_id = custid;
    SELECT * FROM rental WHERE customer_id = custid;
END //
-- next, changing the delimiter back
DELIMITER ;

CALL customer_rentals_w_total (135,@total_rentals);

SELECT @total_rentals;
```

# SCHEDULED EVENTS

- MySQL allows us to specify Events, often referred to as Scheduled Events

- Similar to Stored Procedures, Events are created with a user-defined SQL operation

- Unlike Procedures, Events will fire at a certain time, or at regular recurring intervals

## MySQL Workbench in Action

```
-- just making a new schema
CREATE SCHEMA schema_for_events;

USE schema_for_events;

-- we'll create timestamp values for this table using events
CREATE TABLE sillytable (
    timestamps_via_event DATETIME
);

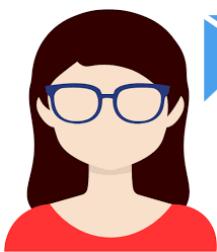
CREATE EVENT myfirstevent
    ON SCHEDULE AT CURRENT_TIMESTAMP + INTERVAL 1 MINUTE
    DO
        INSERT INTO sillytable VALUES (NOW());

-- run this quick, wait a minute, try again. Magic!
SELECT * FROM sillytable;

-- now we'll create a recurring event
CREATE EVENT mysecondevent
    ON SCHEDULE
        EVERY 20 SECOND
    DO
        INSERT INTO sillytable VALUES (NOW());

SELECT * FROM sillytable;

DROP EVENT mysecondevent; -- to keep it from running forever
```



## NEW MESSAGE

April 1, 2014

From: **Sally Bleu (CEO)**

Subject: **Help me pull reports on my own?**

Hey there!

With all this great data, I would love to be able to quickly pull together the total orders and revenue for a given time period. I'm not a SQL guru though.

Is there a way I could specify a **startDate** and **endDate** and see total orders and revenue during that period?

Thanks for the help!

-Sally

Reply

Forward

## Result Preview

-- you'll want her to be able to type something like this...

```
CALL order_performance('2013-11-01','2013-12-31');
```

5 -- the results should look like this...

6

100% 1:2

Result Grid



Filter Rows:

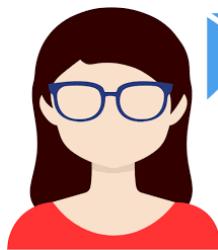


Search

Export

total_orders	total_revenue
1908	104893.62

# TEST YOUR SKILLS: STORED PROCEDURES



**NEW MESSAGE**  
April 1, 2014

From: **Sally Bleu (CEO)**

Subject: **Help me pull reports on my own?**

Hey there!

With all this great data, I would love to be able to quickly pull together the total orders and revenue for a given time period. I'm not a SQL guru though.

Is there a way I could specify a **startDate** and **endDate** and see total orders and revenue during that period?

Thanks for the help!

-Sally

Reply

Forward

## *Solution Query*

DELIMITER //

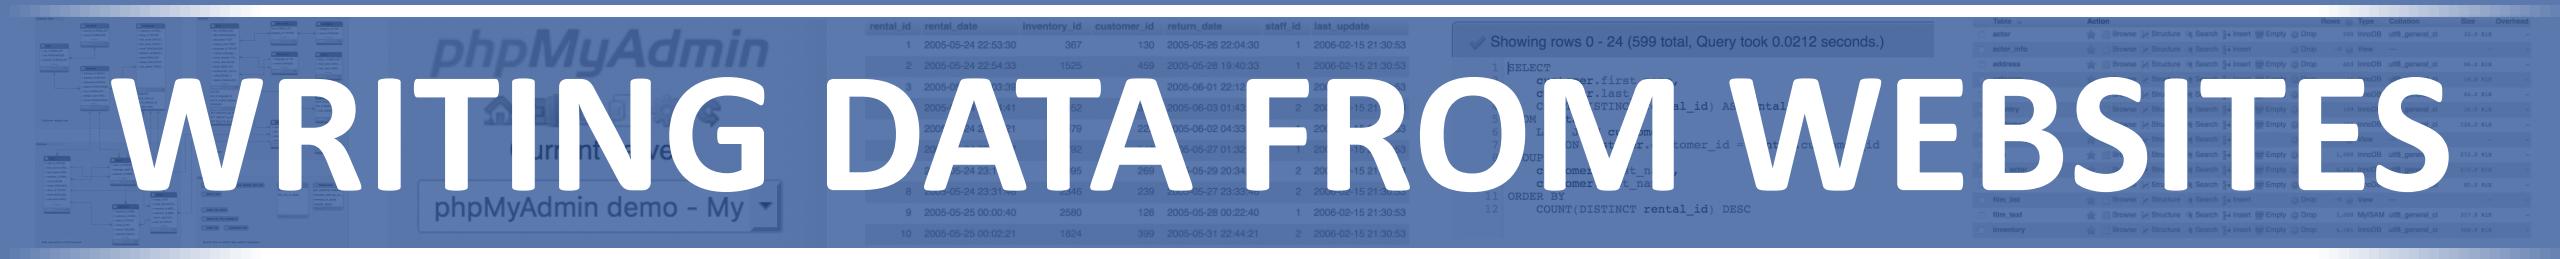
```
CREATE PROCEDURE order_performance
(IN startdate DATE, IN enddate DATE)
BEGIN
SELECT
COUNT(order_id) AS total_orders,
SUM(price_usd) AS total_revenue
FROM orders
WHERE DATE(created_at) BETWEEN startdate AND enddate;
END //
```

DELIMITER ;

```
CALL order_performance('2013-11-01','2013-12-31');
```

# TEST YOUR SKILLS: STORED PROCEDURES

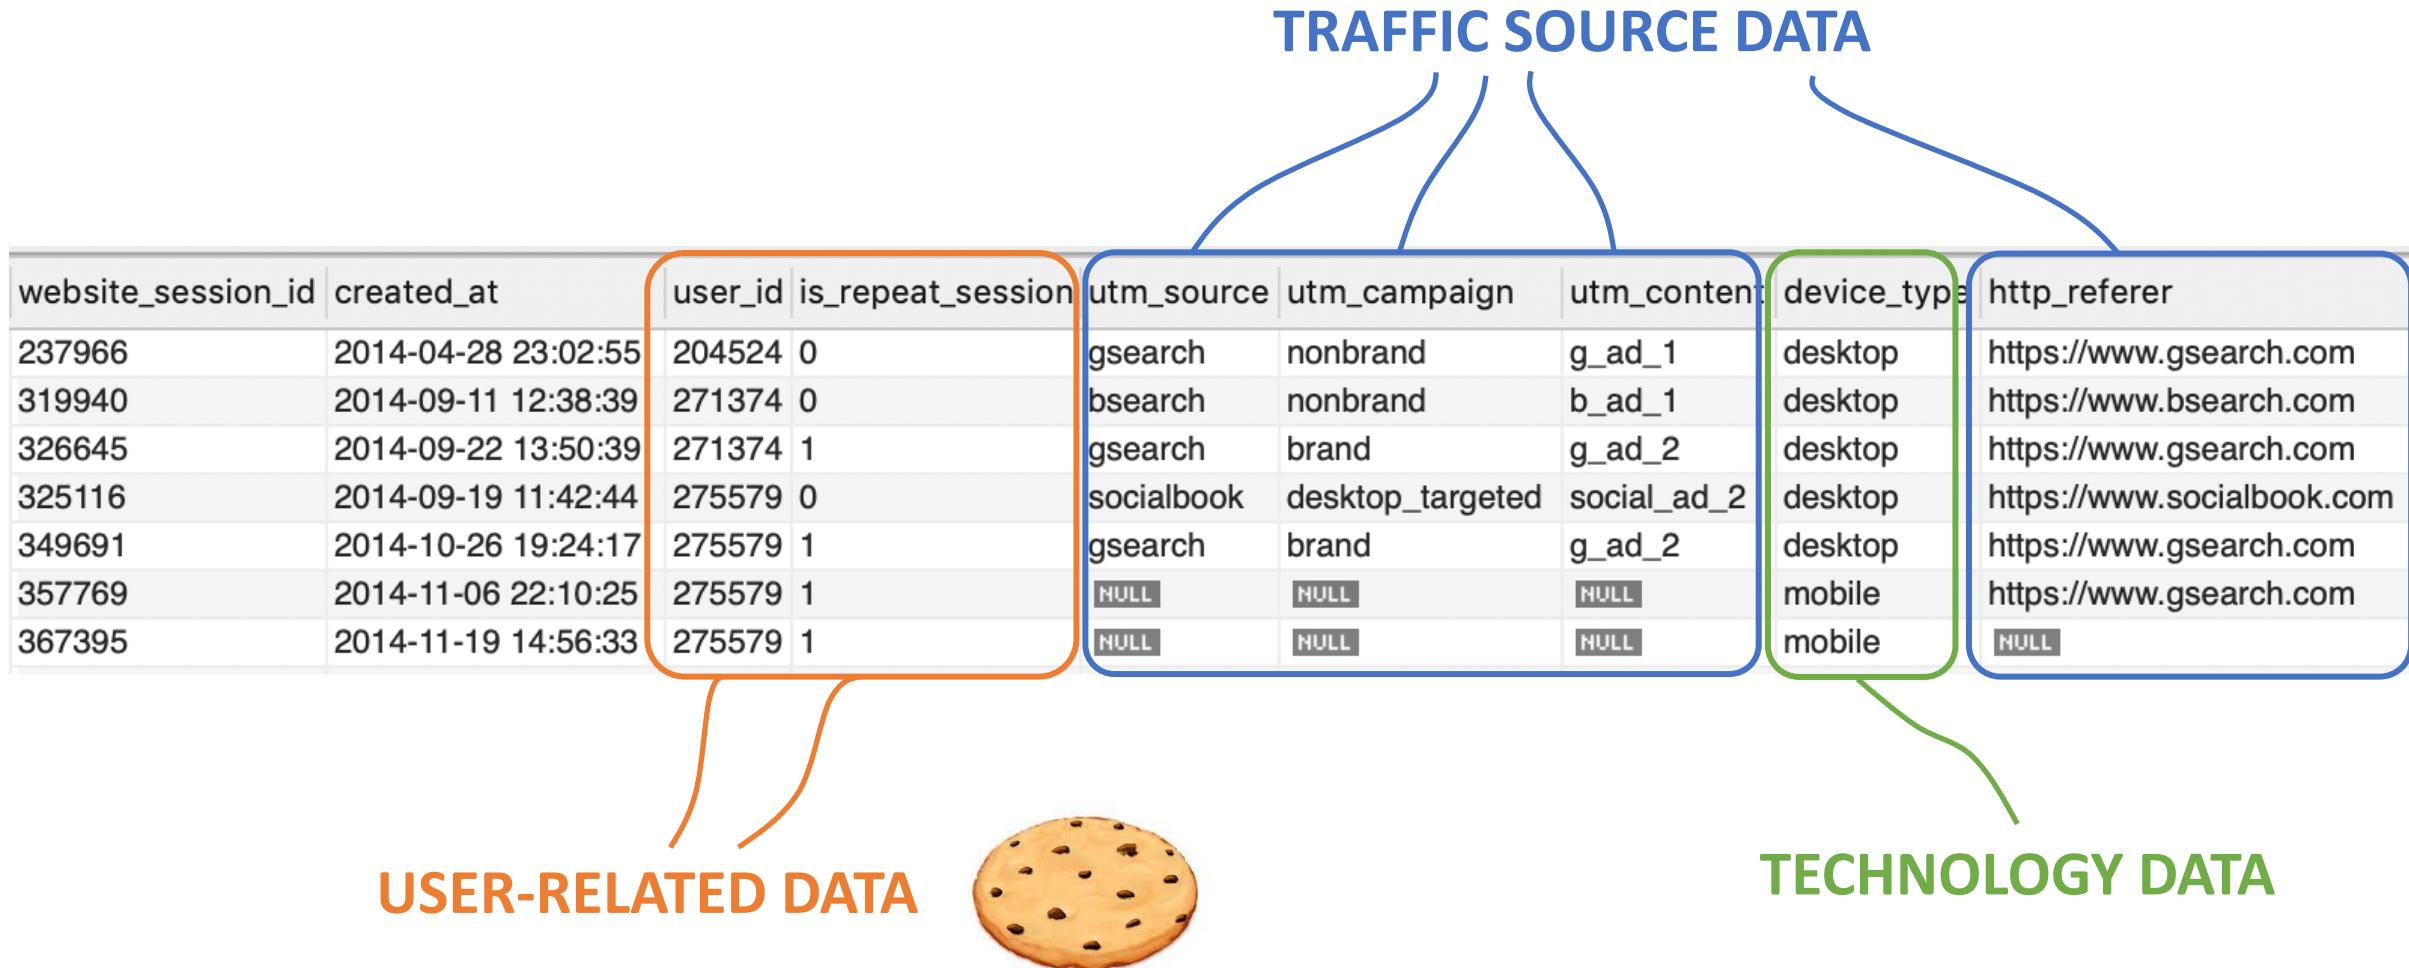
# WRITING DATA FROM WEBSITES



# WRITING DATA FROM WEBSITES



# TYPES OF DATA WEBSITES CAPTURE AT SESSION START



# TYPES OF DATA WEBSITES CAPTURE ALONG THE WAY

website_pageview_id	created_at	website_session_id	pageview_url
500157	2014-03-09 20:50:49	210749	/cart
500158	2014-03-09 20:50:49	210751	/lander-4
500159	2014-03-09 20:51:48	210747	/shipping
500160	2014-03-09 20:53:49	210747	/billing-2
500161	2014-03-09 20:54:21	210752	/lander-3
500162	2014-03-09 20:54:47	210753	/home
500163	2014-03-09 20:54:51	210754	/lander-3
500164	2014-03-09 20:55:50	210753	/products
500165	2014-03-09 20:56:34	210753	/the-original-mr-fuzzy
500166	2014-03-09 20:56:45	210754	/products
500167	2014-03-09 20:58:27	210755	/home
500168	2014-03-09 20:59:08	210747	/thank-you-for-your-order
500169	2014-03-09 21:04:39	210756	/home
500170	2014-03-09 21:06:14	210757	/lander-3
500171	2014-03-09 21:07:14	210757	/products

## Purchases + Key Business Activities

- Business will always want to capture data related to purchases, subscriptions, lead signups, etc.
- Many times this data is triggered by website activities

order_id	created_at	website_session_id	user_id	primary_product_id	items_purchased	price_usd	cogs_usd
20648	2014-09-23 10:07:43	327430	277430	1	2	19.98	20.98
20649	2014-09-23 16:17:08	327437	277442	1	1	49.99	19.49
20650	2014-09-23 16:20:46	327448	277450	2	1	59.99	22.49
20651	2014-09-23 16:41:23	327474	277472	1	2	79.98	28.98
20652	2014-09-23 16:42:05	327472	277471	1	1	49.99	19.49
20653	2014-09-23 16:43:13	327475	277473	3	2	95.98	33.98
20654	2014-09-23 17:01:26	327486	277482	1	1	49.99	19.49
20655	2014-09-23 17:17:51	327499	277492	3	2	75.98	23.98
20656	2014-09-23 17:21:42	327503	277495	1	2	79.98	28.98
20657	2014-09-23 17:48:29	327531	277518	1	1	49.99	19.49
20658	2014-09-23 17:58:38	327536	277522	1	2	95.98	33.98
20659	2014-09-23 18:00:05	327554	277539	1	1	49.99	19.49
20660	2014-09-23 18:56:46	327583	277560	1	2	79.98	28.98
20661	2014-09-23 18:57:52	327582	277559	1	1	49.99	19.49
20662	2014-09-23 18:58:09	327578	231588	1	2	79.98	28.98
20663	2014-09-23 19:12:32	327596	277571	2	1	59.99	22.49

## Pageview and Other Event / Activity Data:

- As activities happen on the website, data records are created to track and measure behavior
- This can happen on a page load, the download of a white paper, etc.

# RECAP OF CAPTURING WEBSITE DATA



- User-related data (cookies)
  - Traffic source data
  - Technology data
  - Pageview data
- Pageviews
  - Other important events
- Purchases
  - Subscriptions
  - Lead signups
  - Other business activity



## NEW MESSAGE

April 6, 2014

From: **Molly Monterey (Website Manager)**  
Subject: Any recommended enhancements?

Good morning!

Given your expertise and your knowledge of the website data we are already tracking in our SQL database, is there any additional data you recommend we add?

Thanks!

-Molly

Reply

Forward

## Helpful Hints

- think about the website data topics we covered
- is anything missing that you would want?

# TEST YOUR SKILLS: WRITING WEBSITE DATA



## NEW MESSAGE

April 6, 2014

From: **Molly Monterey (Website Manager)**

Subject: Any recommended enhancements?

Good morning!

Given your expertise and your knowledge of the website data we are already tracking in our SQL database, is there any additional data you recommend we add?

Thanks!

-Molly

Reply

Forward

## Solution

-- see video for solution

# TEST YOUR SKILLS: WRITING WEBSITE DATA



## NEW MESSAGE

April 15, 2014

From: **Molly Monterey (Website Manager)**

Subject: **Loading Website Pageview Data**

Hey!

I was able to get this February pageview data out of our web analytics tool. Would you be able to help me load it into the database so we can tie it to all of your other great data?

Thanks!

-Molly



16.website\_pageviews\_2014\_Feb

Reply

Forward

## Helpful Hints

- take a look at the pageview data structure
- create a website\_pageviews table
- then upload the dataset

# TEST YOUR SKILLS: IMPORTING PAGEVIEW DATA



## NEW MESSAGE

April 15, 2014

From: **Molly Monterey (Website Manager)**

Subject: **Loading Website Pageview Data**

Hey!

I was able to get this February pageview data out of our web analytics tool. Would you be able to help me load it into the database so we can tie it to all of your other great data?

Thanks!

-Molly



16.website\_pageviews\_2014\_Feb

Reply

Forward

## Solution Query

```
CREATE TABLE website_pageviews (
    website_pageview_id BIGINT,
    create_at DATETIME,
    website_session_id BIGINT,
    pageview_url VARCHAR(50),
    PRIMARY KEY (website_pageview_id)
)
;
```

```
22 -- you should see 38,391 records if you've done this right
23 • SELECT COUNT(*) AS total_records FROM website_pageviews;
```

Result Grid		Filter Rows:	Search	Export:
total_records				
▶	38391			

# TEST YOUR SKILLS: IMPORTING PAGEVIEW DATA

**SECURITY**

The image shows a screenshot of the phpMyAdmin web application. On the left, there is a database schema diagram with various tables like actor, address, category, city, country, customer, film, film\_actor, film\_category, film\_text, and inventory. In the center, a query results page displays a table of rental records with columns: rental\_id, rental\_date, inventory\_id, customer\_id, return\_date, staff\_id, and last\_update. Below this, a SQL query is shown:

```

SELECT customer.first_name,
       customer.last_name,
       COUNT(DISTINCT rental_id) AS rentals
FROM rental
LEFT JOIN customer
ON customer.customer_id = rental.customer_id
GROUP BY customer.first_name,
         customer.last_name
ORDER BY COUNT(DISTINCT rental_id) DESC

```

On the right, a privilege list table shows rows for various database objects with actions like Browse, Structure, Search, Insert, Update, Delete, and View.

Table	Action	Rows	Type	Collation	Size	Overflow
actor	Browse	200	Empty	utf8_general_ci	33.0 K	-
actor.info	Browse	200	Empty	utf8_general_ci	44.0 K	-
address	Browse	400	Empty	utf8_general_ci	16.0 K	-
category	Browse	400	Empty	utf8_general_ci	44.0 K	-
city	Browse	400	Empty	utf8_general_ci	44.0 K	-
country	Browse	100	Empty	utf8_general_ci	16.0 K	-
customer	Browse	300	Empty	utf8_general_ci	120.0 K	-
customer.list	Browse	300	Empty	utf8_general_ci	120.0 K	-
film	Browse	1000	Empty	utf8_general_ci	370.0 K	-
film.actor	Browse	1000	Empty	utf8_general_ci	370.0 K	-
film.category	Browse	1000	Empty	utf8_general_ci	49.0 K	-
film.list	Browse	1000	Empty	utf8_general_ci	370.0 K	-
film.text	Browse	1000	Empty	utf8_general_ci	370.0 K	-
inventory	Browse	400	Empty	utf8_general_ci	60.0 K	-

# COMMON SECURITY THREATS & ATTACKS

1

## Weak Authentication

*Attackers can take advantage of limited or weak authentication methods in place to break through your perimeter*

2

## Denial of Service

*DoS attackers will attempt to flood or crash your system, rendering it useless for legitimate users*

3

## Privilege Escalation

*Attackers will use this strategy to gain access to higher levels of your systems and information you don't want them to see*

4

## SQL Injection

*Various endpoints of your system can create vulnerabilities for hackers who understand how your back-end works*

5

## Buffer Overflow

*Similar to SQL injection, attackers can take advantage of holes in your application and make it do things it shouldn't*

6

## Ransomware

*Malware attacks will render your data unusable unless you agree to pay a ransom*

# WEAK AUTHENTICATION

---

**Authentication** refers to the processes you employ for users to prove to your system that they are who they say they are.

Attackers can gain access to your system by tricking your system into thinking they are one of your trusted users.

**They may do this by taking advantage of:**

- Weak passwords
- Passwords that do not change frequently
- Lack of a two-factor or multi-factor authentication mechanism

# DENIAL OF SERVICE ATTACKS

With **Denial of Service (DoS) attacks**, the hackers will attempt to overwhelm your system with requests, which will render it useless for your legitimate users.

There are services you can employ that will filter out suspicious activity, and when a DoS attack is detected, there are a number of mechanisms for blocking the attack.

In one special case -- a **Distributed Denial of Service (DDoS)** -- attackers will coordinate the attack from multiple machines simultaneously, often distributed globally. This makes it harder to determine the attacking source and shut them off quickly.



## THIS IS IMPORTANT!

A DoS or DDoS can sometimes be used to distract from and cover up additional attacks on your systems. In some cases, the intent is to extract data elsewhere. When you encounter a DoS, be on high alert across all systems.

# PRIVILEGE ESCALATION ATTACKS

With **privilege escalation attacks**, hackers will attempt to gain access to increasingly secure levels of information by authenticating as a lower level user first, then climbing the ladder.

Once they have gained access to some of your systems, it can become easier for them to identify additional security holes and gain higher levels of access.

Some common methods for implementing these attacks:

- **Cookie-based** authentication
- **Form-based** authentication

*With cookie-based authentication, hackers may spoof this to impersonate another user*



`user_id = 666`

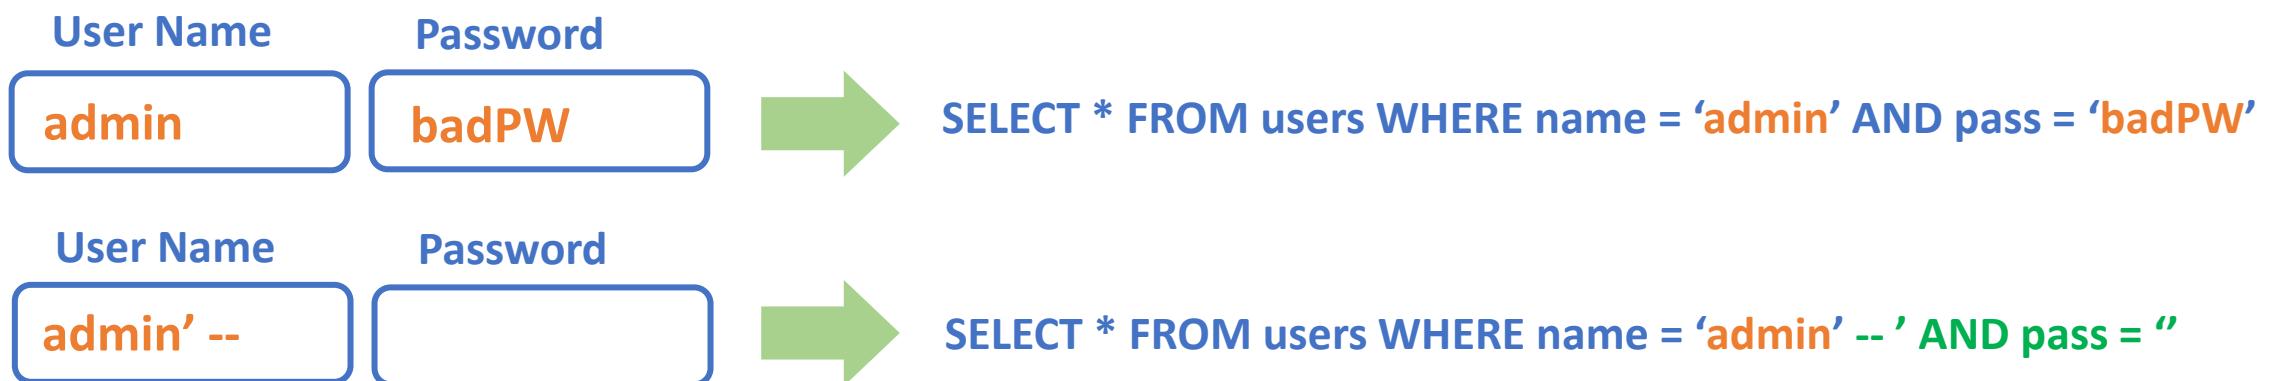


# SQL INJECTION ATTACKS

With **SQL injection attacks**, hackers will attempt to make your application do something it's not supposed to do by editing SQL statements utilized on the back-end.

Common ways attackers might use SQL injection maliciously:

- Gain information about your database (*structure, version, etc.*)
- Return results from your database that they shouldn't see



# BUFFER OVERFLOW VULNERABILITY

---

A **buffer** is a sequential section of memory which will contain anything from a string to an array of numbers.

A **buffer overflow** is when more data is put into the buffer than it can handle. This causes the data to overflow and get written elsewhere.

Attackers can deliberately overflow your application to crash it or to force it to do things you might not want it to do.

# RANSOMWARE

**Ransomware** has become more prevalent in recent years. It represents an attractive opportunity for hackers and an expensive problem for businesses.

When a business is attacked by ransomware, systems will be rendered unusable and the attacker will demand payment (*typically in the form of cryptocurrency*)

Malware most often breaches systems by getting employees to download a file from their email, but social media attacks have been reported as well.



## THIS IS IMPORTANT!

Preparing for and defending against ransomware attacks needs to be a team effort. The Database Administrator and CTO should prepare by having adequate and safe backups, and employees should be trained to avoid phishing.

# SECURITY BEST PRACTICES

1

## Practice Safe Data Storage

*Steer clear of storing personal information if you can. Anything sensitive that you do need to store should be encrypted.*

2

## Limit Access

*Practice “minimum viable access”. Don’t grant more permissions than are needed for the job.*

3

## Take Authentication Seriously

*Require strong passwords, require changing frequently, and enable multi-factor authentication.*

4

## Dedicate Resources to Security

*Whether it is a team, a person, or some percentage of a person’s time, someone in your organization should be focused on security.*

5

## Have a Data Security Policy

*Come up with a formal security policy. Write it down. Train your employees on best practices.*

6

## Backup, Log, Monitor & Audit

*Backup your data. Log changes & monitor access. Perform periodic audits & invite third parties to try and hack you.*

# PRACTICE SAFE DATA STORAGE

---

- Make every attempt to store as little **personally identifiable information (PII)** as possible
  - *If you don't absolutely need it for your application, consider NOT storing it*
- For sensitive data that you need to store (*credit cards, social security #, email address, phone number, passwords, etc.*), make sure to **encrypt** it whenever possible
- Make a list** of the applications and people in your organization that have access to sensitive data, and where you are vulnerable. Is the list longer than it should be?

# LIMIT ACCESS TO SYSTEMS AND DATA

---

- Make sure to practice “**minimum viable access**” in granting your permission levels
  - *Very few employees will need full access to systems and data. Grant only what they need.*
- When thinking about permission levels, consider whether employees need edit rights or if **read-only** access is sufficient for their job function
- Make sure that you have a plan for quickly **removing access** to systems when employees leave, or when you need to limit access on a case-by-case basis

# TAKE AUTHENTICATION SERIOUSLY

---



- Require **strong passwords** for employees, and don't allow them create weak passwords
  - Minimum 8 characters, containing a number, a special character, and both upper- and lower-case letters*



- Require employees to **change their password** once every 60 to 90 days
  - They may hate it, but it will mitigate potential attacks*



- Leverage **two-factor or multi-factor** authentication methods, especially for employees with access to sensitive information

# DEDICATE RESOURCES TO SECURITY

---

- Whether it is a whole team, a single person, or part of someone's time (*in very small companies*), **someone in your organization should be responsible for security planning**
- Make security a part of someone's **job description and performance review**
- Give the responsible employee the **authority** to work with the rest of your organization to ensure that things are set up properly

# HAVE A WRITTEN SECURITY POLICY

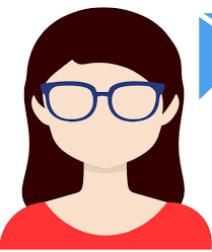
---

- Writing down your security policy forces you to create a **cohesive plan** and determine **specific action items** that you will be held accountable for completing
- Share your written security policy with all employees
- Make sure you are continually **reviewing and revising** your security policy as your business and your risks evolve over time

# BACKUP, LOG & MONITOR YOUR DATA. PERFORM AUDITS

---

- Back up** your data, especially the most important information; in the event of a breach, system malfunction, or malware, you will be happy to have another copy
- Create logs** on admin systems and database activity so you can monitor activity
- Conduct periodic reviews** and invite friendly third parties to try hacking your system to expose potential vulnerabilities



## NEW MESSAGE

April 20, 2014

From: **Sally Blue (CEO)**

Subject: **We need a security plan!**

Hey there!

I know you've been learning a lot about security, and the board is giving me a lot of pressure on that topic.

Can you put together a comprehensive security plan for us? Include anything you think makes sense for today, and also as we scale up.

-Sally

Reply

Forward

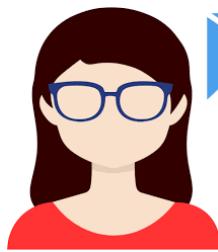
## Result Preview

-- recommend you write out a formal plan

-- it can be a bulleted list, nothing fancy

-- writing it down will force you to think about it

# TEST YOUR SKILLS: DATABASE SECURITY



## NEW MESSAGE

April 20, 2014

From: **Sally Blue (CEO)**

Subject: **We need a security plan!**

Hey there!

I know you've been learning a lot about security, and the board is giving me a lot of pressure on that topic.

Can you put together a comprehensive security plan for us? Include anything you think makes sense for today, and also as we scale up.

-Sally

Reply

Forward

## *Solution Query*

-- see video for solution

# TEST YOUR SKILLS: DATABASE SECURITY

# INTRODUCING THE FINAL COURSE PROJECT

## THE **SITUATION**

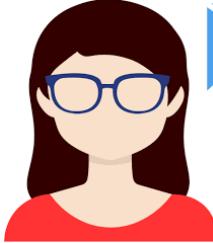
There have been some exciting developments for Maven Bear Builders. The company is going to start offering chat support on the website, and needs your help planning. The company has also been approached by potential acquirers, and you'll be asked to help with due diligence.

## THE **OBJECTIVE**

**Use SQL to:**

- Create a plan for handling chat support, including the database infrastructure, EER diagrams explaining your plan, and reports to help management understand performance
- Provide support for questions relating to the potential acquisition, to help your CEO keep them interested and hopefully close the deal

# INTRODUCING THE FINAL PROJECT



## NEW MESSAGE

May 1, 2014

From: **Sally Bleu (CEO)**  
Subject: **Exciting Developments!**

Good morning!

Two big updates for you...

- 1.** We're adding chat support to the website, and need your help planning for the data structure.
  
- 2.** We were approached by a larger company who is interested in acquiring us. Will need some help here too.

Let's do this!

-Sally

Reply

Forward

## YOUR OBJECTIVES:

- Update the database with the most recent data
  
- Create a plan for the company's service expansion to include chat support
  
- Help Sally with some support for asks related to the potential acquisition

# FINAL COURSE PROJECT QUESTIONS

1

**Import the latest order\_items and order\_item\_refunds data below into the database, and verify the order summary trigger you created previously still works (*if not, recreate it*)**

~ 0:37



17.order\_items\_2014\_Mar  
18.order\_items\_2014\_Apr  
19.order\_item\_refunds\_2014\_Mar  
20.order\_item\_refunds\_2014\_Apr

2

**Import the website\_sessions and website\_pageviews data for March and April, provided below:**

~ 5:30



21.website\_sessions\_2014\_Mar  
22.website\_sessions\_2014\_Apr  
23.website\_pageviews\_2014\_Mar  
24.website\_pageviews\_2014\_Apr

3

The company is adding chat support to the website. You'll need to **design a database plan** to track which customers and sessions utilize chat, and which chat representatives serve each customer

~ 9:14

4

Based on your tracking plan for chat support, **create an EER diagram** that incorporates your new tables into the existing database schema (*including table relationships*)

~ 14:00

# FINAL COURSE PROJECT QUESTIONS

---

5

**Create the tables** from your chat support tracking plan in the database, and include relationships to existing tables where applicable

~ 20:40

6

Using the new tables, **create a stored procedure** to allow the CEO to pull a count of chats handled by chat representative for a given time period, with a simple CALL statement which includes two dates

~ 26:37

7

**Create two Views** for the potential acquiring company; one detailing monthly order volume and revenue, the other showing monthly website traffic. Then **create a new User**, with access restricted to these Views

~ 31:18

8

The potential acquirer is commissioning a third-party security study, and your CEO wants to get in front of it. Provide her with a **list of your top data security threats** and recommendations for mitigating risk

~ 37:00